

Deep Image Classification (Medical Image Diagnosis)

[Spring 2020 CS-8395 Deep Learning in Medical Image Computing]

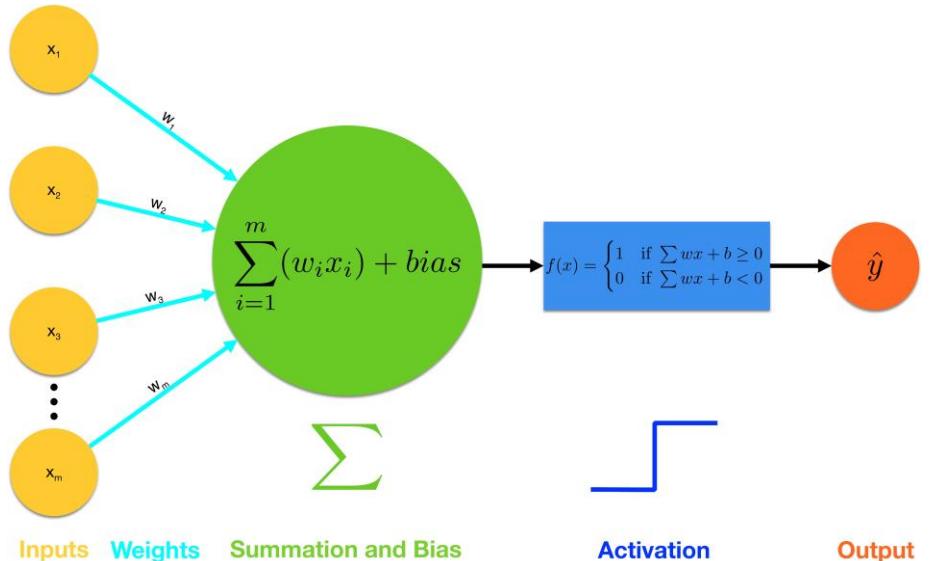
Instructor: Yuankai Huo, Ph.D.
Department of Electrical Engineering and Computer Science
Vanderbilt University

Topics



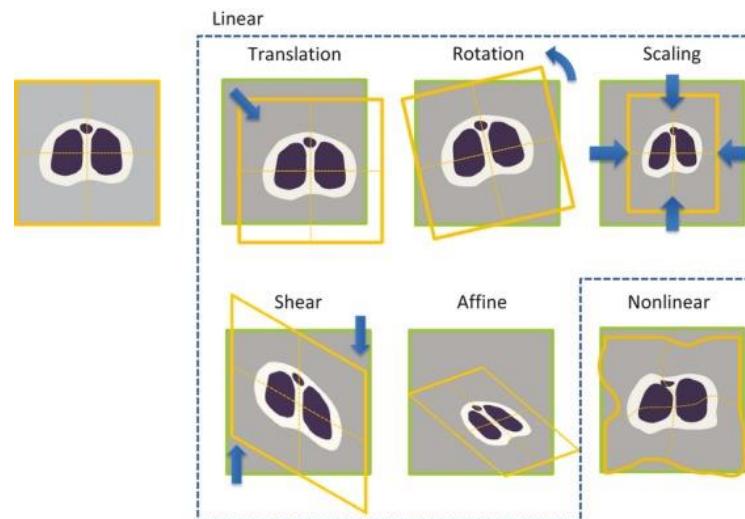
- Review
- Classification
- Discuss the Paper

Neural Network



$$\hat{y} = \sigma(w^T x + b)$$

$w^T x + b$
linear transformation
(e.g. affine registration)



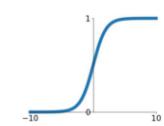
<http://nahuakang.com/page3/>

$$\sigma(\cdot)$$

Activation Functions

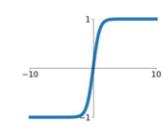
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



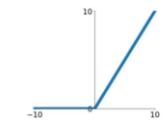
tanh

$$\tanh(x)$$



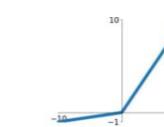
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

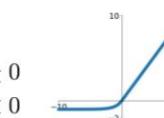


Maxout

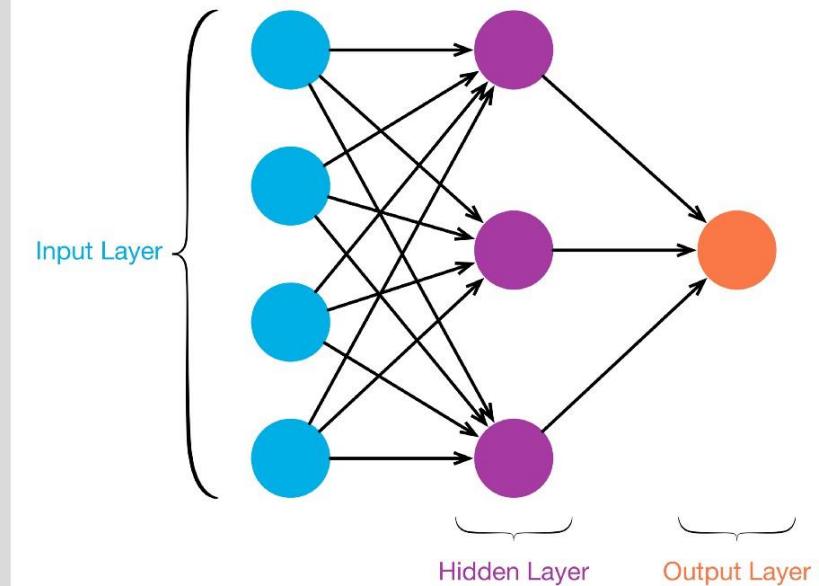
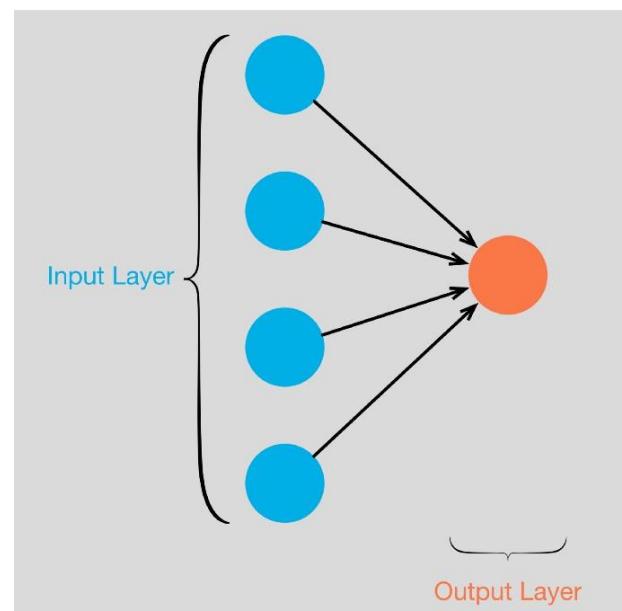
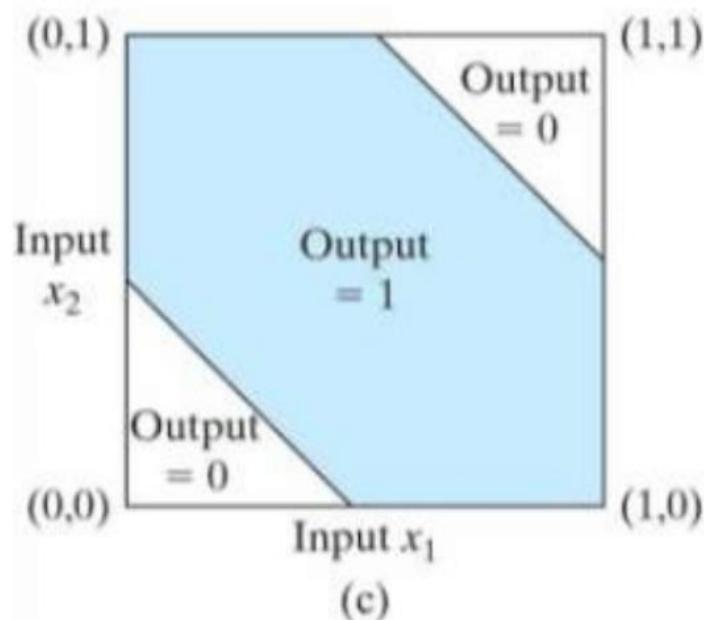
$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



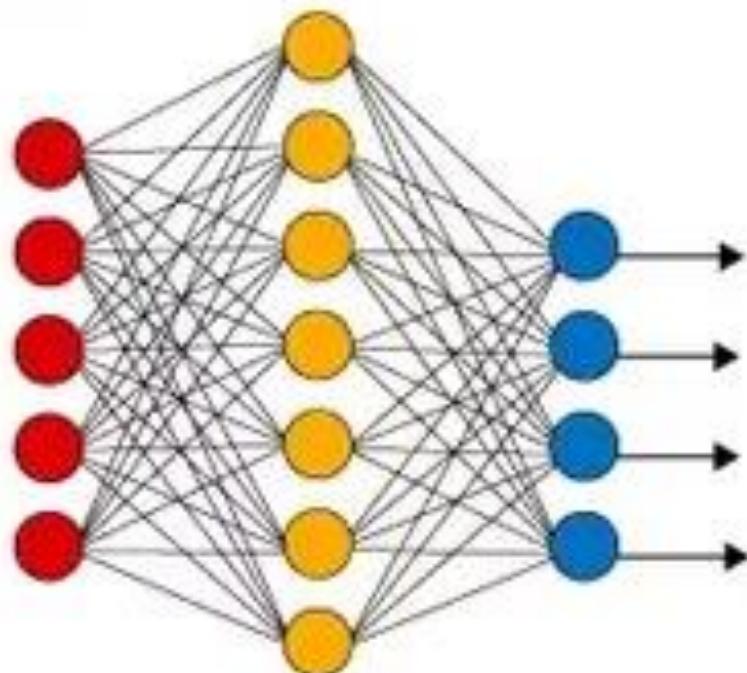
Multi-Layer Perceptrons



<https://towardsdatascience.com/multi-layer-neural-networks-with-sigmoid-function-deep-learning-for-rookies-2-bf464f9eb7f>

Deep Learn = Deep Neural Network

Simple Neural Network

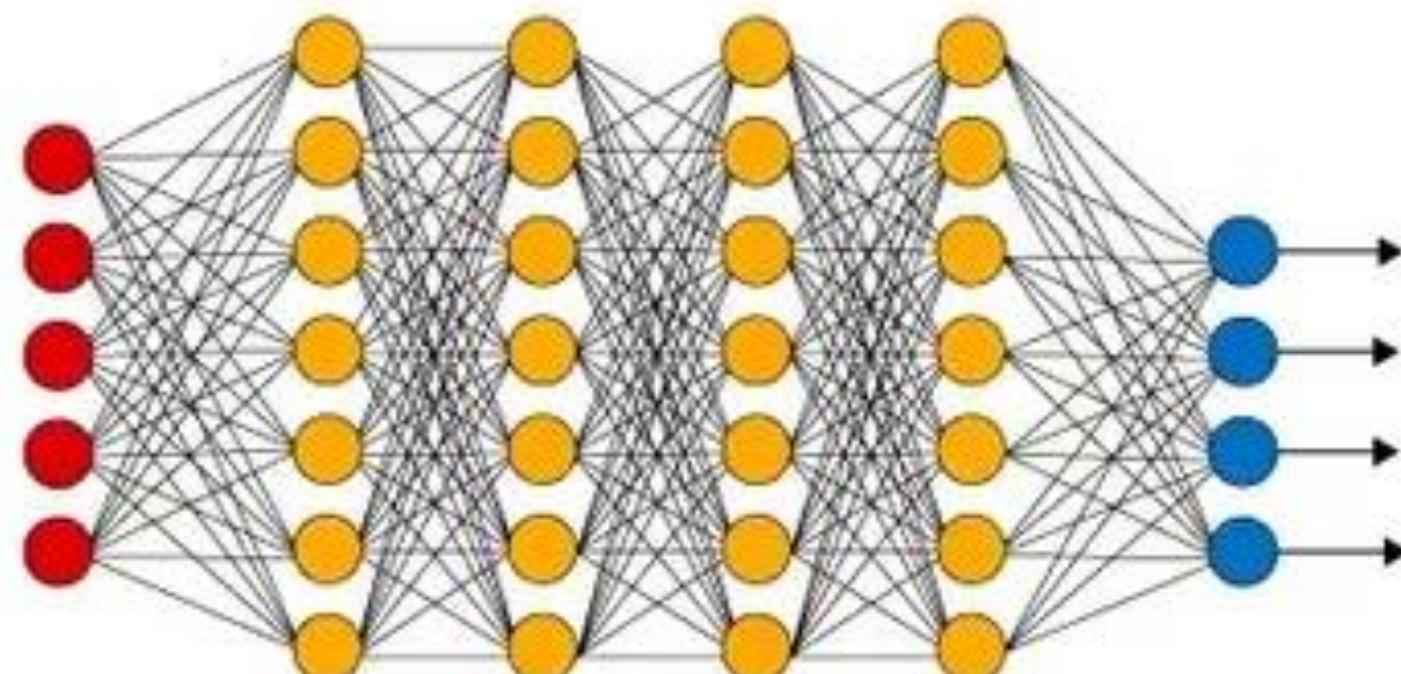


● Input Layer

○ Hidden Layer

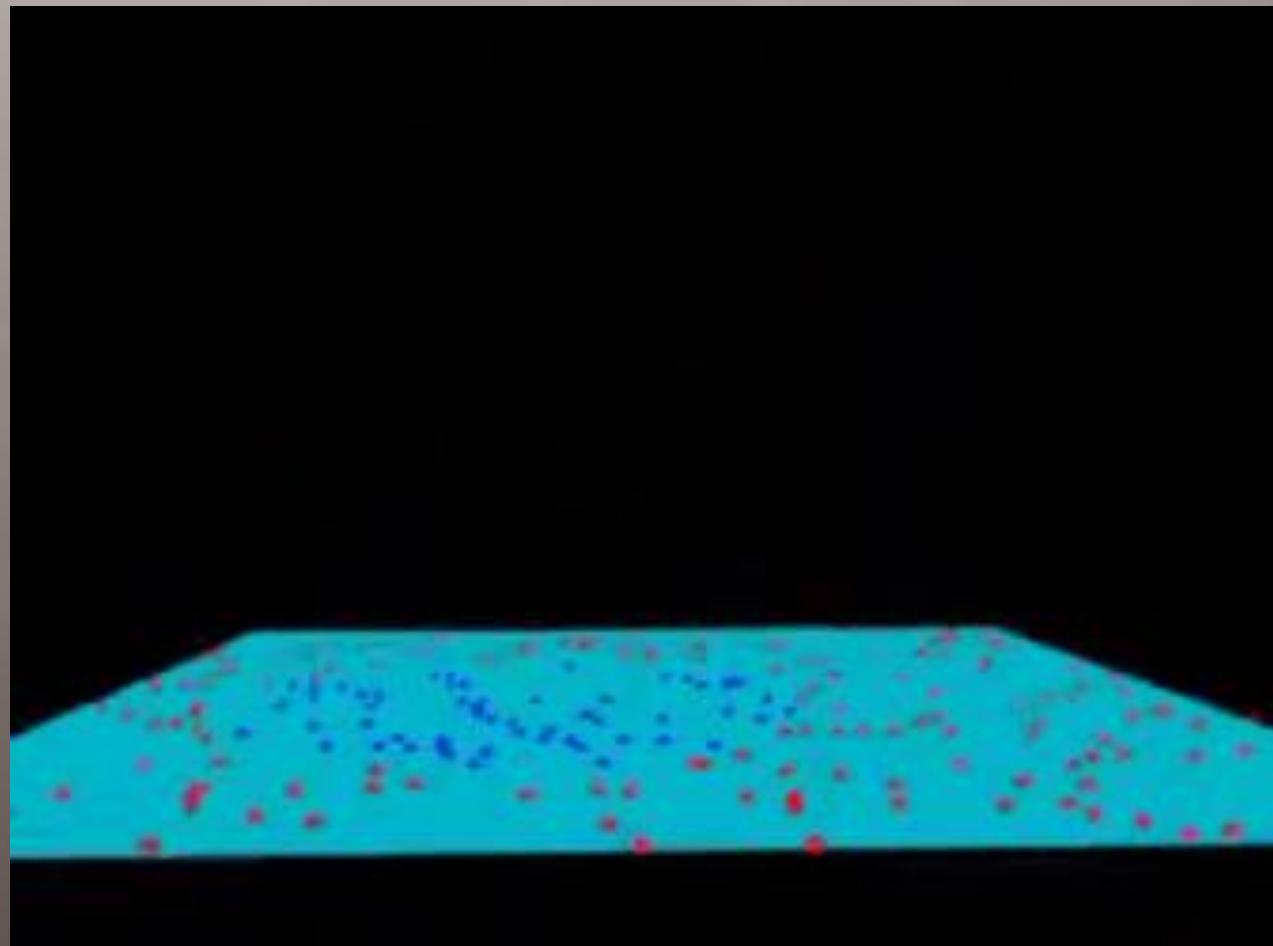
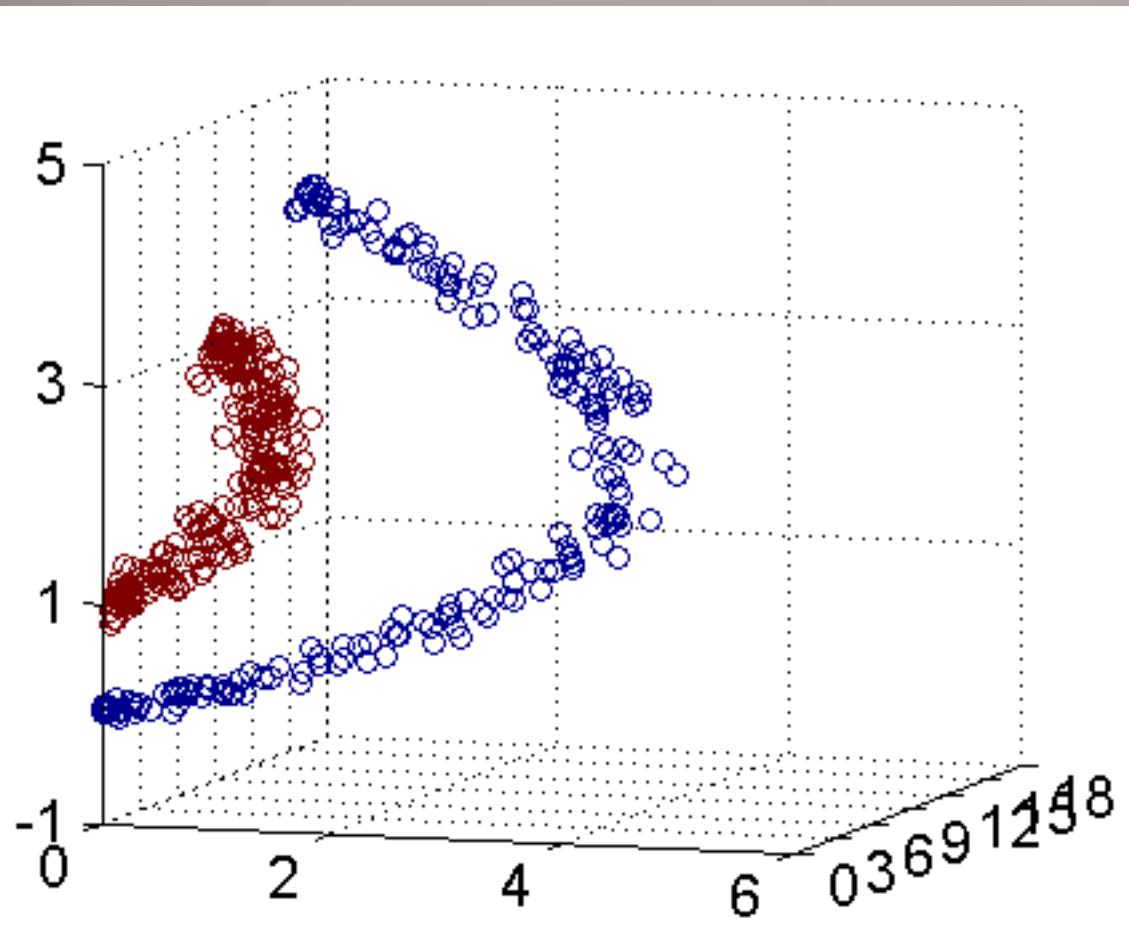
● Output Layer

Deep Learning Neural Network

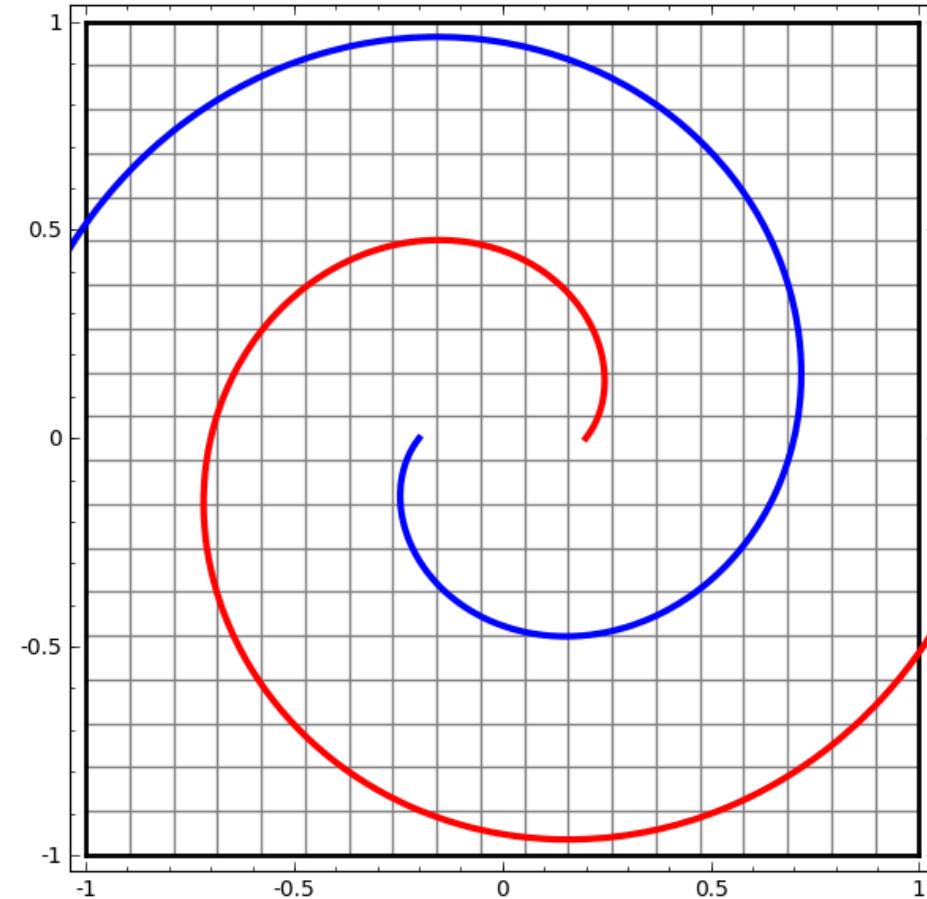
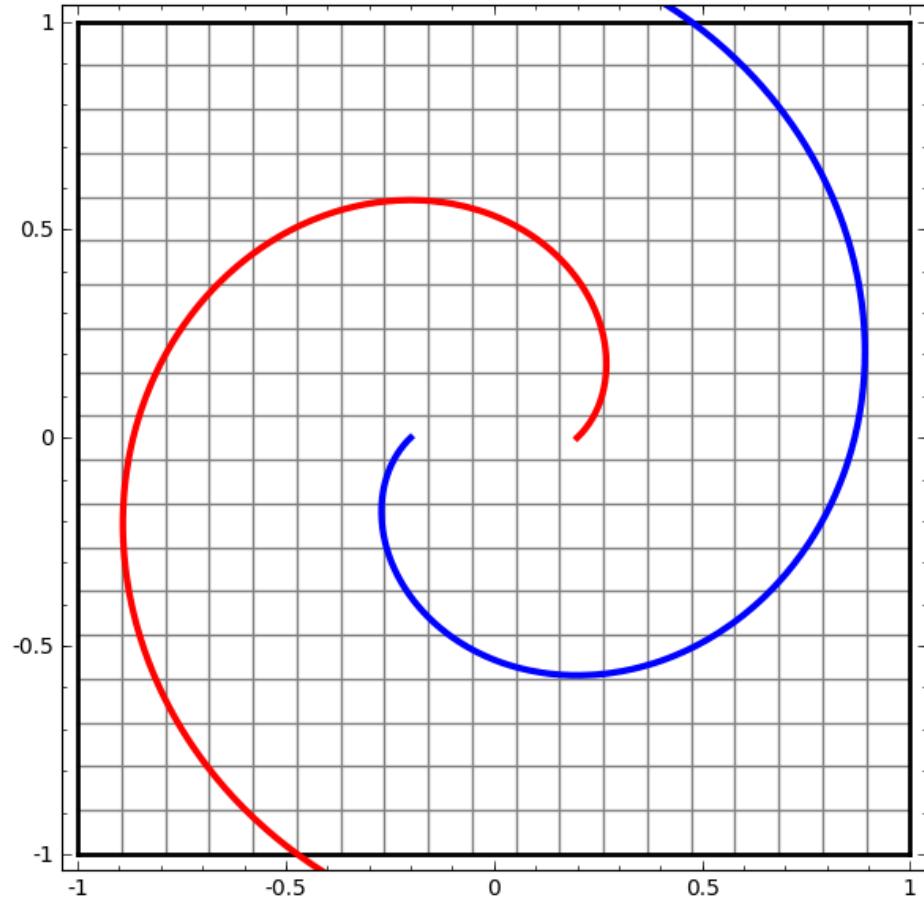


<https://www.quora.com/What-is-the-difference-between-Neural-Networks-and-Deep-Learning>

Hyper Plane

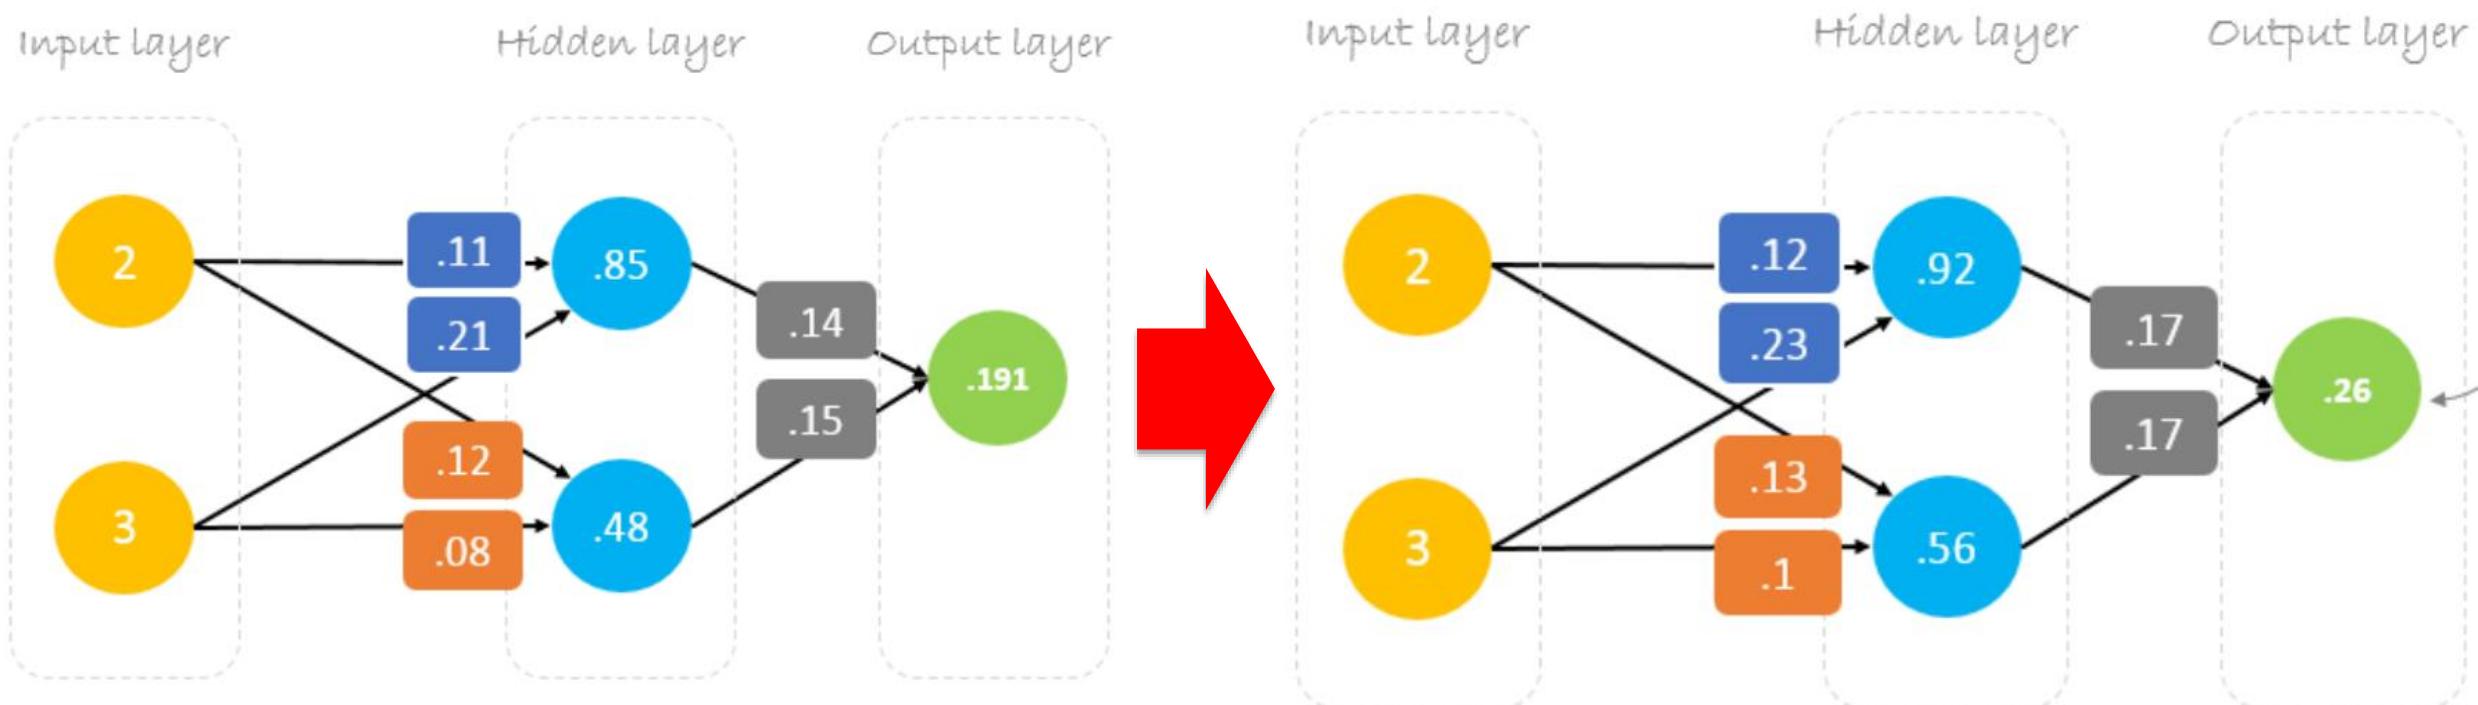


More Hidden Layers



<http://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>

Backward Path



1D Data

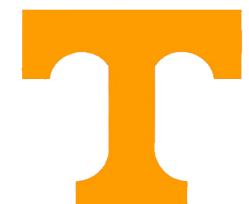
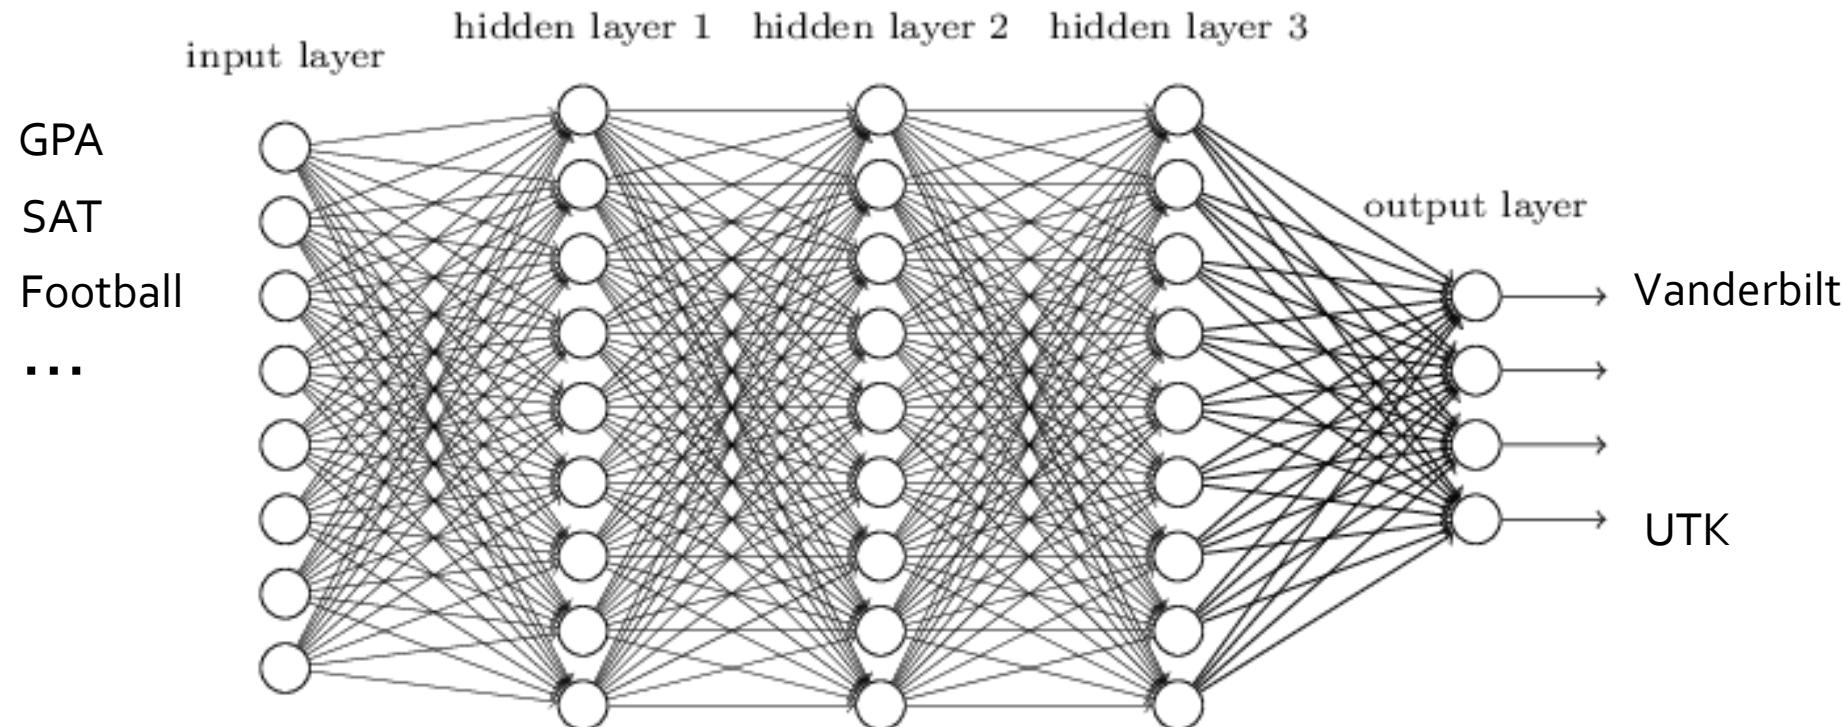
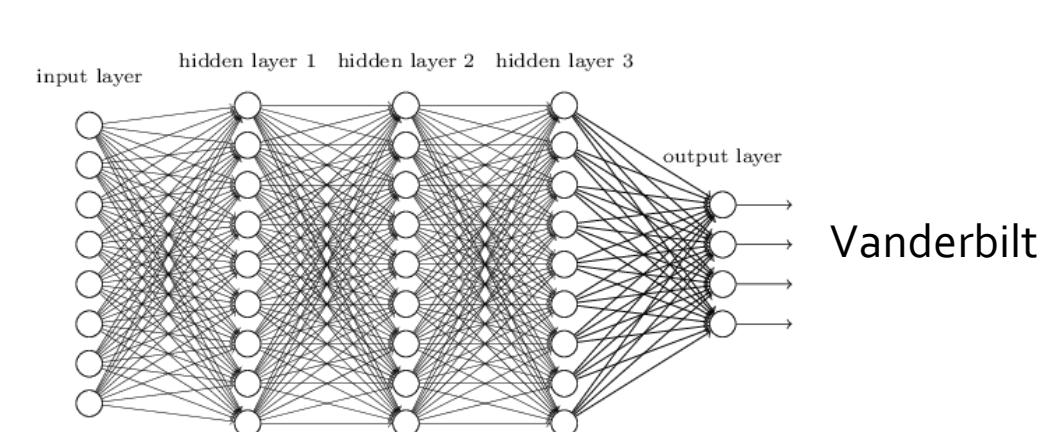
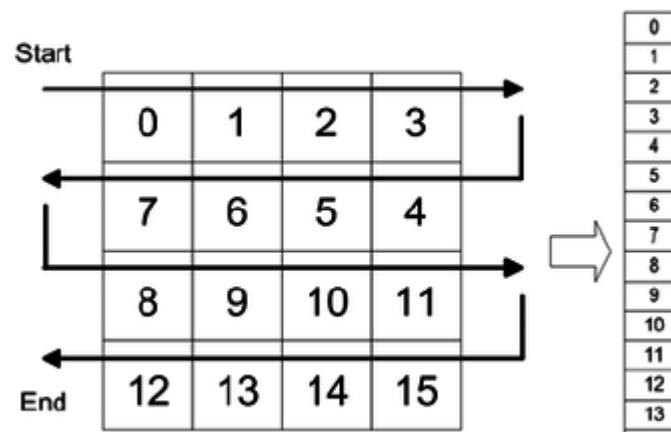
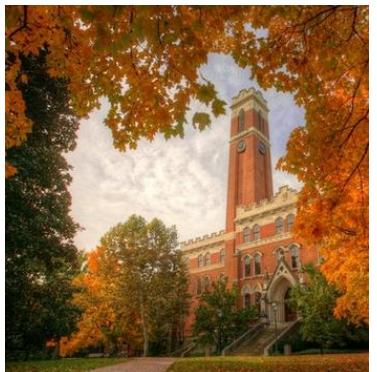


Image Representation 1: Dense



<https://www.usnews.com/best-colleges/vanderbilt-3535>

https://www.researchgate.net/figure/Scanning-2D-blocks-left-into-1D-vector-as-input-into-logical-transform-middle-the_fig4_228345705

Image Representation 2: Convolution

1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0	0
0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1	0
0 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

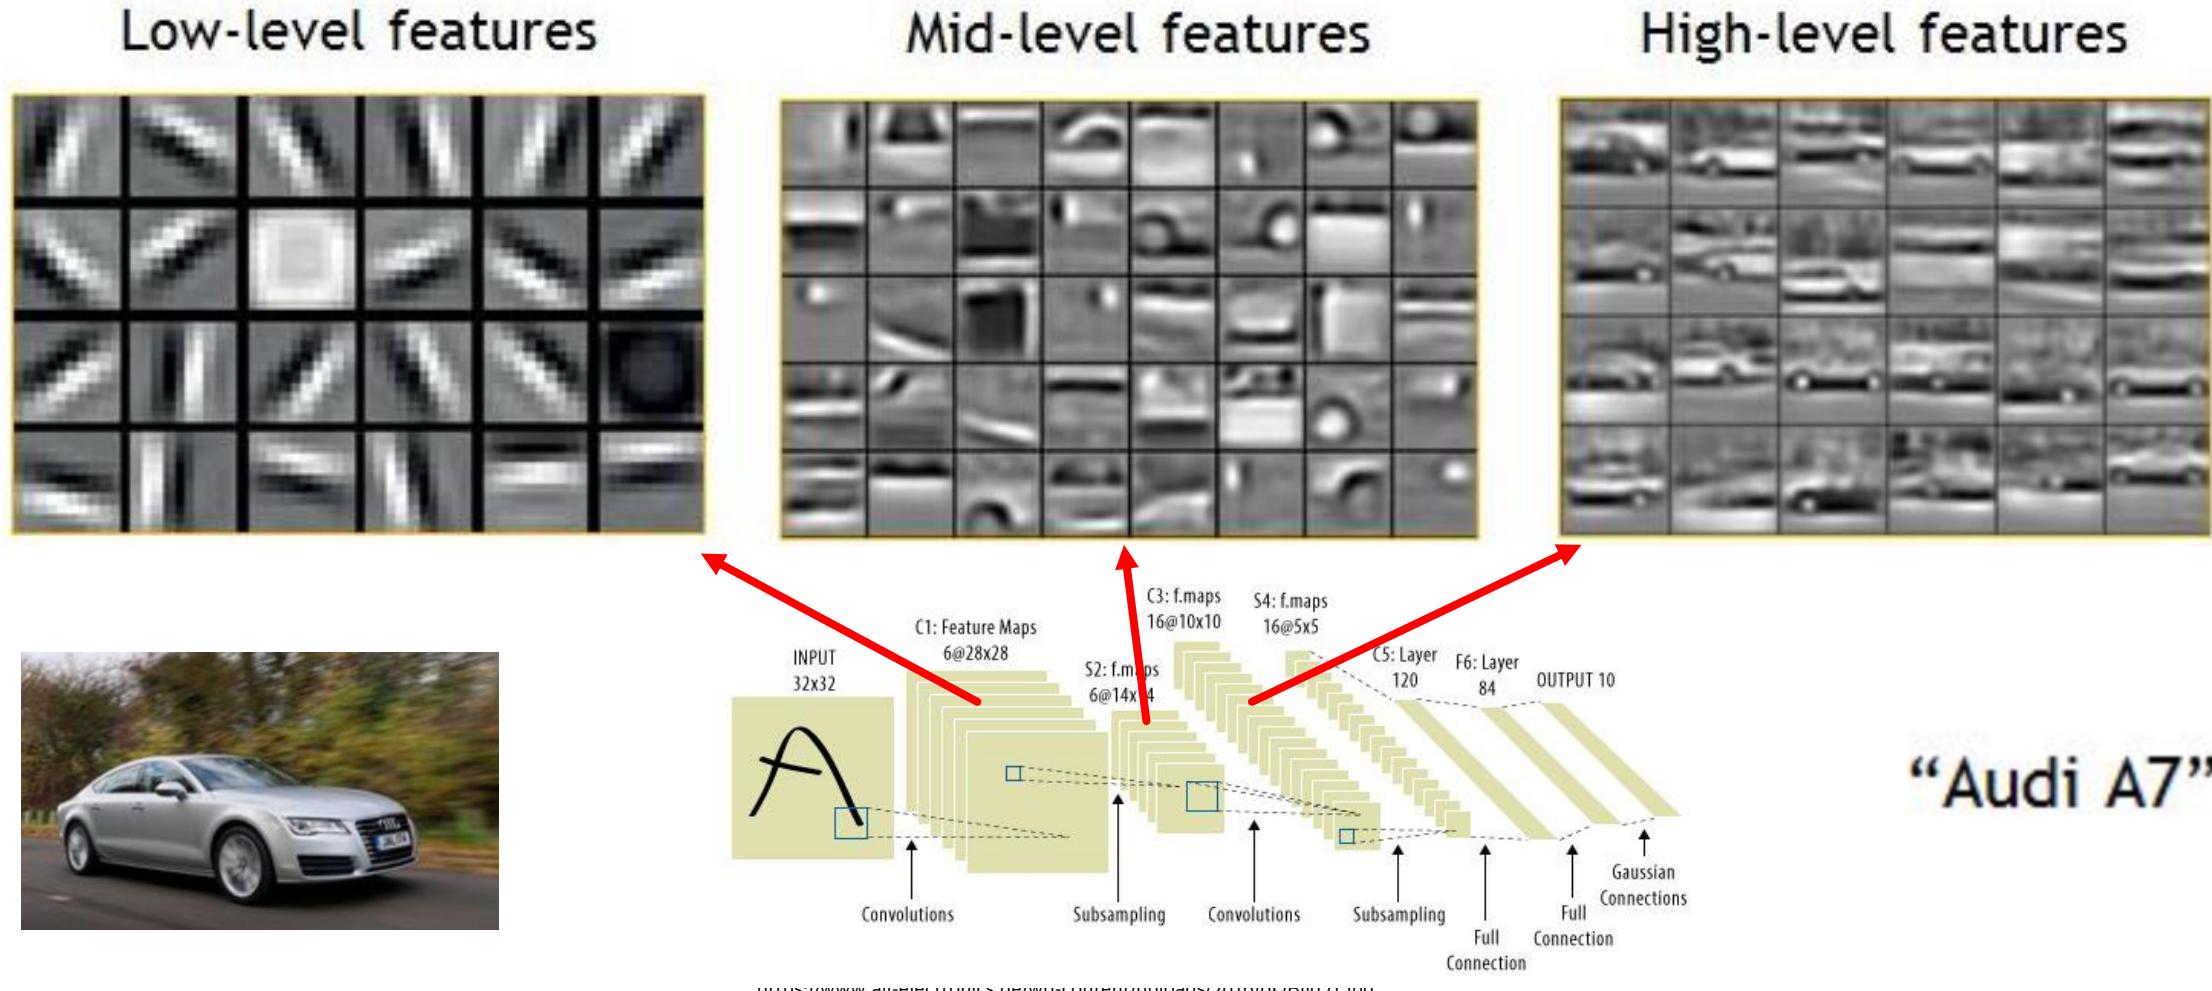
Convolved
Feature



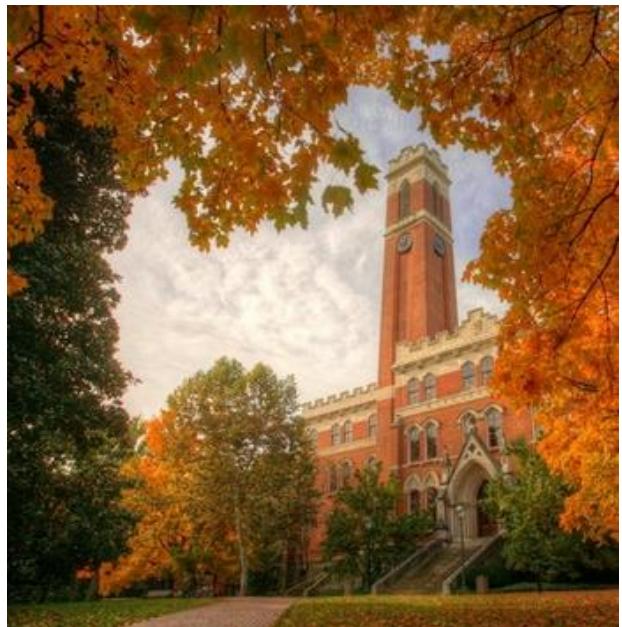
http://ufldl.stanford.edu/wiki/index.php/Feature_extraction_using_convolution

https://leonardoaraujosantos.gitbooks.io/artificial-intelelligence/content/convolutional_neural_networks.html

Different Layers



Parameters



<https://www.usnews.com/best-colleges/vanderbilt-3535>

```
import utilities
```

```
img_rows, img_cols = 224, 224
colors = 3
input_size = img_rows * img_cols * colors
input_shape = (img_rows, img_cols, colors)

num_classes = 10
```

DNN vs CNN

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential([
    Dense(32, activation='relu', input_shape=(input_size,)),
    Dense(64, activation='relu'),
    Dense(128, activation='relu'),
    Dense(num_classes, activation='softmax')
])

model.summary()
```

Using TensorFlow backend.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 32)	4816928
dense_2 (Dense)	(None, 64)	2112
dense_3 (Dense)	(None, 128)	8320
dense_4 (Dense)	(None, 10)	1290

Total params: 4,828,650
Trainable params: 4,828,650
Non-trainable params: 0

DNN

```
from keras.models import Sequential
from keras.layers import Conv2D

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=input_shape),
    Conv2D(64, (3, 3), activation='relu'),
    Conv2D(128, (3, 3), activation='relu'),
    Dense(num_classes, activation='softmax')
])

model.summary()
```

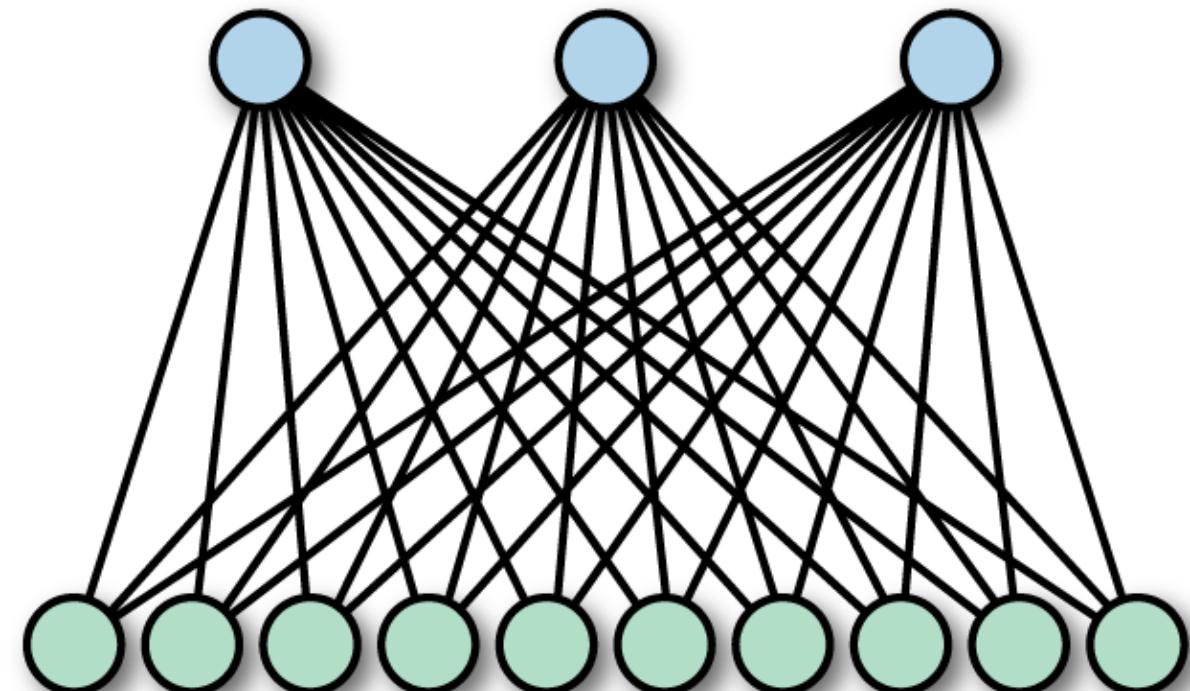
Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 222, 222, 32)	896
conv2d_2 (Conv2D)	(None, 220, 220, 64)	18496
conv2d_3 (Conv2D)	(None, 218, 218, 128)	73856
dense_5 (Dense)	(None, 218, 218, 10)	1290

Total params: 94,538
Trainable params: 94,538
Non-trainable params: 0

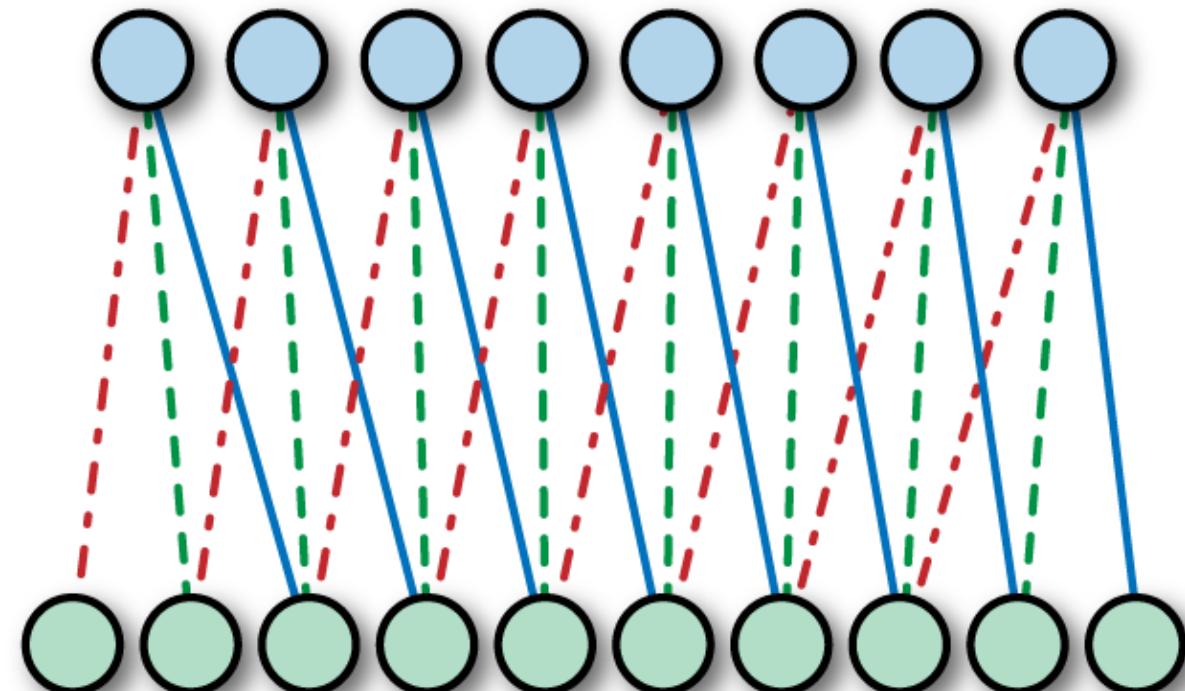
CNN

Reduce Parameters

Fully Connected

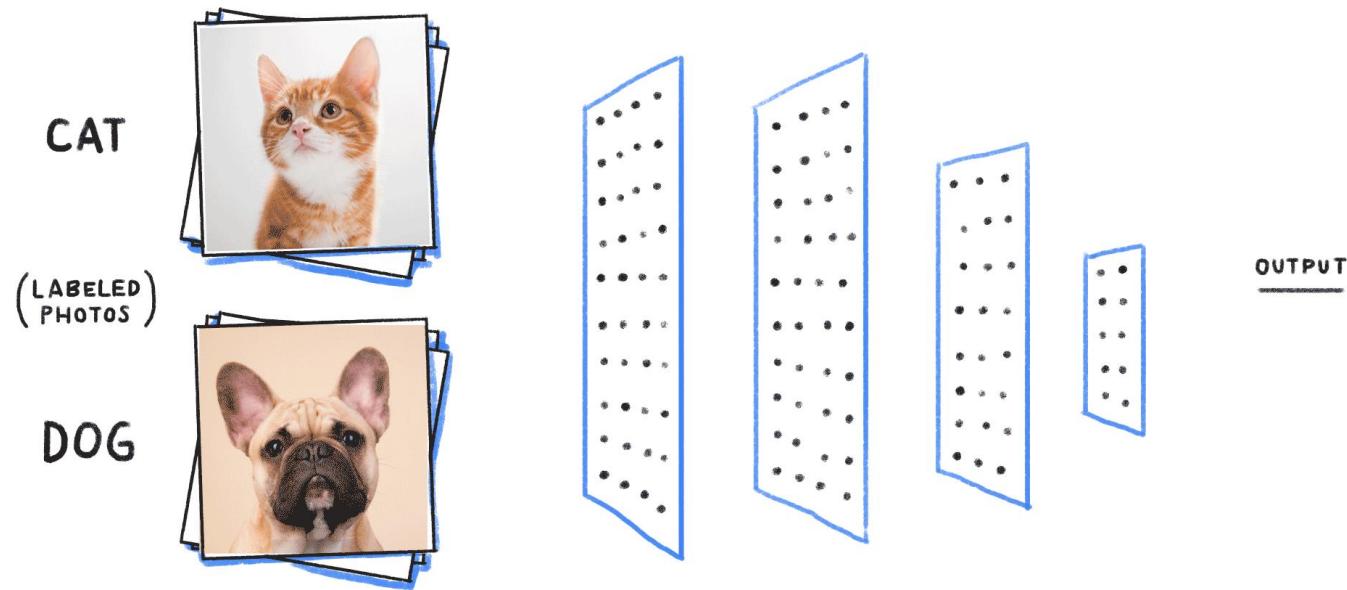
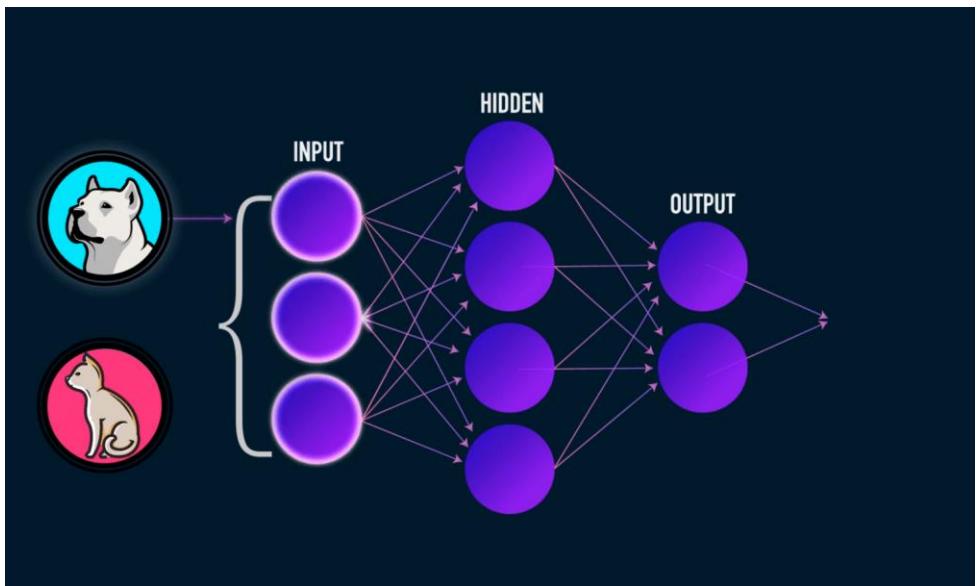


Convolutional Layer



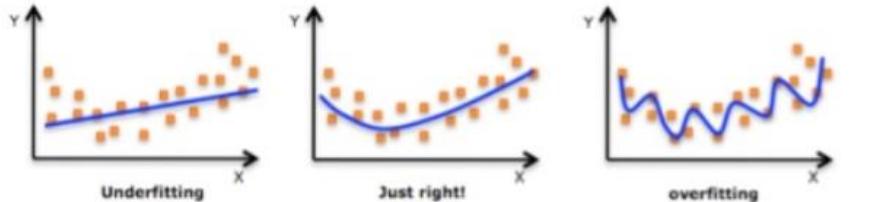
<https://www.safaribooksonline.com/library/view/learning-tensorflow/9781491978504/cho4.html>

Go Deep!

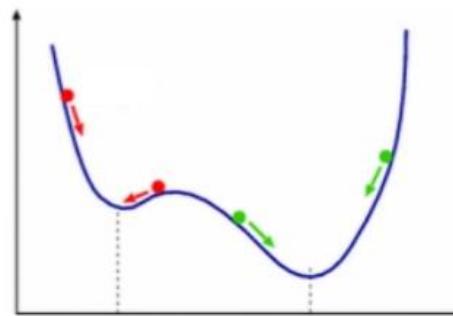
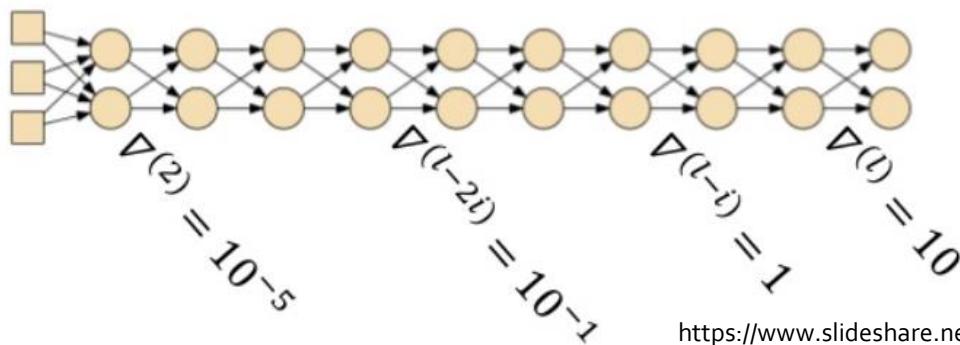


Challenges

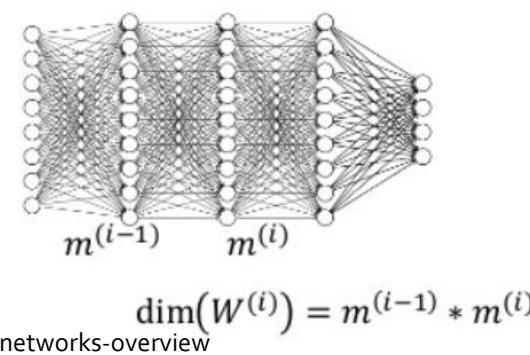
- Big networks => Too huge separating ability => **Overfitting**
- Complex error's surface => **Local minimum**



- **Vanishing gradient problem** during training
- Curse of dimensionality => **memory & computations**

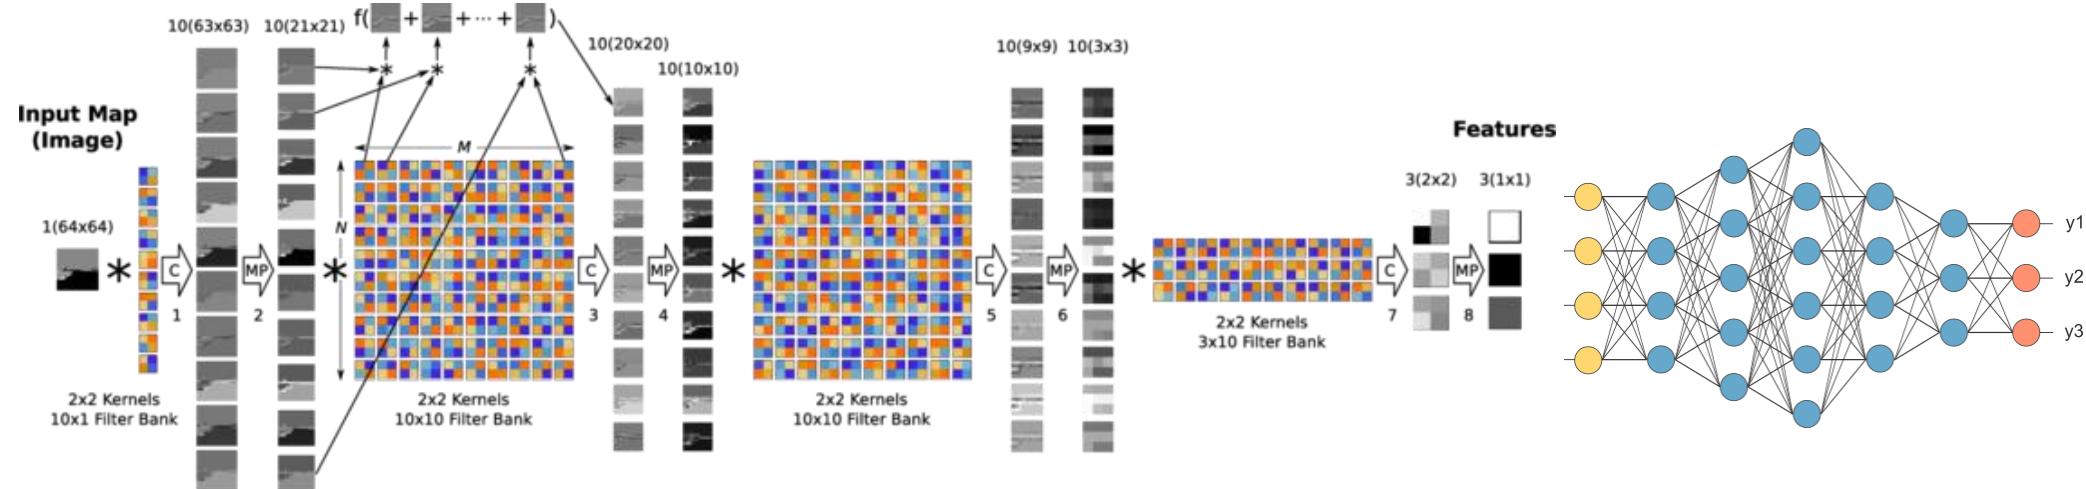


**Overfitting
Local minimum
Gradient Vanishing
Computation & Memory**



<https://www.slideshare.net/baiev1/neural-networks-overview>

Deep Convolutional Network Classification



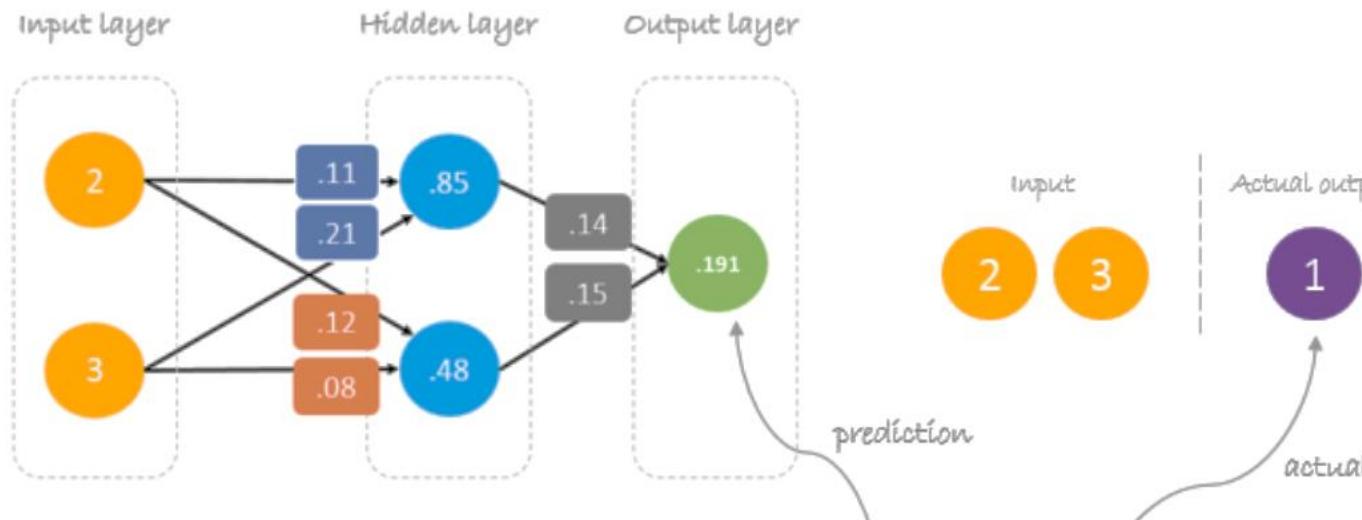
Convolutional Layer

Loss
Function

Fully Connected Layer
(Dense Layer)

https://www.researchgate.net/figure/Max-Pooling-Convolutional-Neural-Network-MPCNN-with-8-layers-alternating-between_fig3_266656356

How To Get Loss



Error = 0, if prediction = actual

$$\text{Error} = \frac{1}{2}(\text{prediction} - \text{actual})^2$$

Error is always positive because of the square

$\frac{1}{2}$ is added to ease the calculation of the derivative

$$\text{Error} = \frac{1}{2}(0.191 - 1.0)^2 = 0.327$$

Three Examples

computed			targets			correct?

0.3	0.3	0.4	0	0	1	yes
0.3	0.4	0.3	0	1	0	yes
0.1	0.2	0.7	1	0	0	no

computed			targets			correct?

0.1	0.2	0.7	0	0	1	yes
0.1	0.7	0.2	0	1	0	yes
0.3	0.4	0.3	1	0	0	no

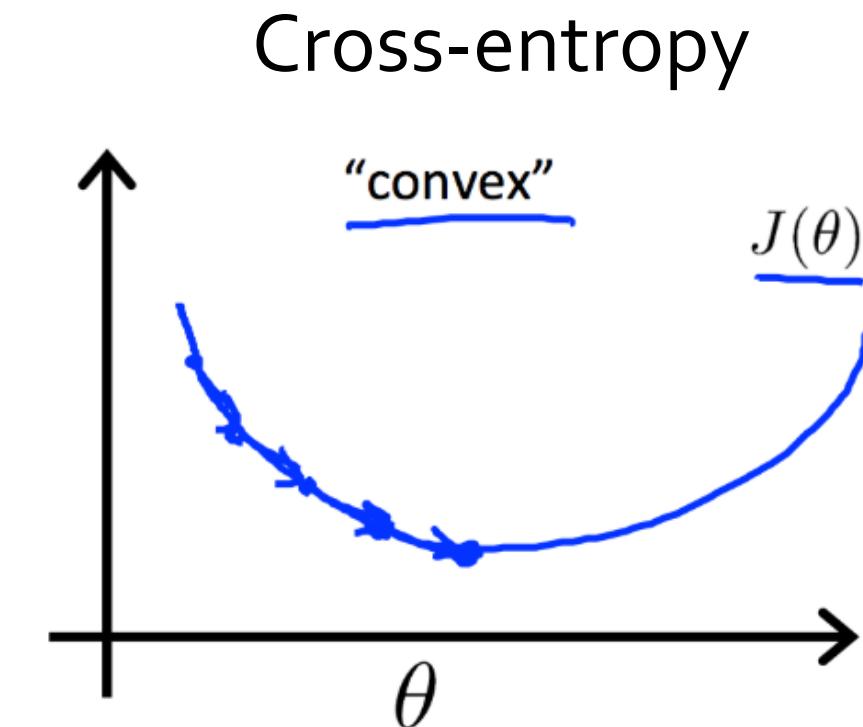
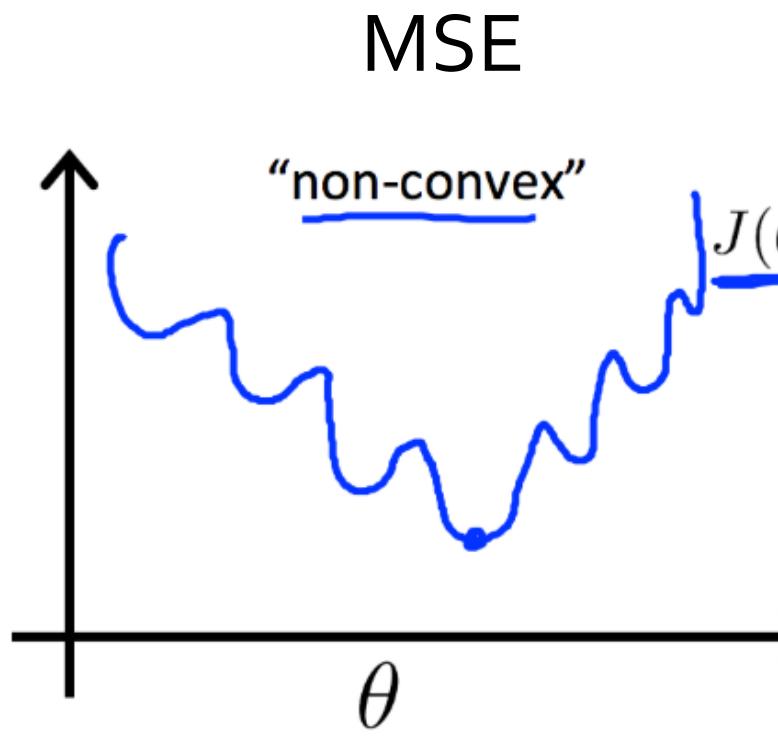
Classification Accuracy: $CA = (\# \text{correct} / \# \text{all})$

Mean square Error (MSE)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$$

$$\text{Cross-entropy} \quad H(p, q) = - \sum_x p(x) \log q(x)$$

MSE vs Cross-entropy



Andrew Ng

Cross Entropy

Entropy

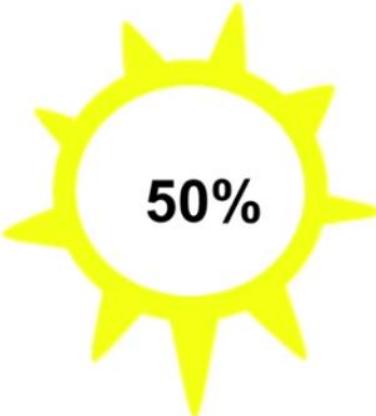
$$H(y, \hat{y}) = -\sum_i y_i \log \frac{1}{\hat{y}_i} = -\sum_i y_i \log \hat{y}_i$$



Claude Shannon

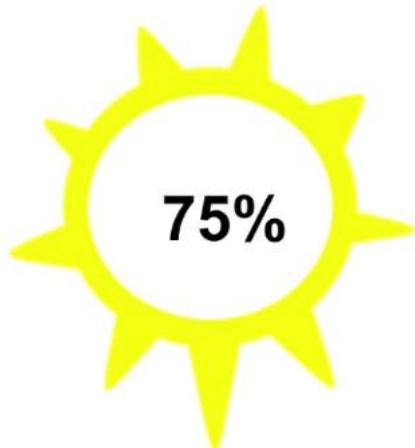
<https://rdipietro.github.io/friendly-intro-to-cross-entropy-loss/>

Entropy



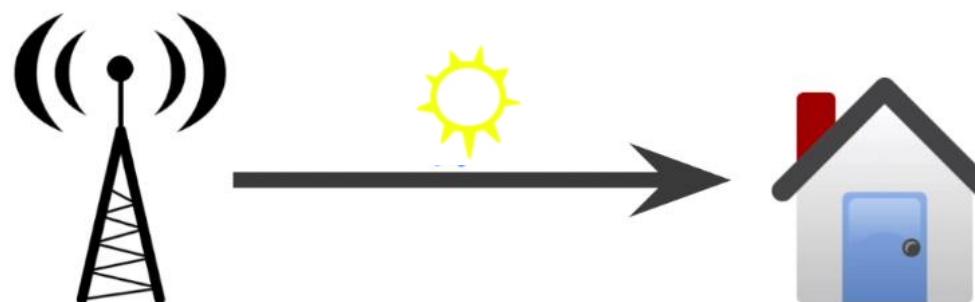
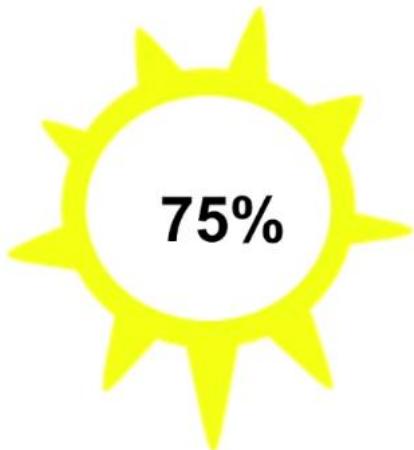
Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Entropy



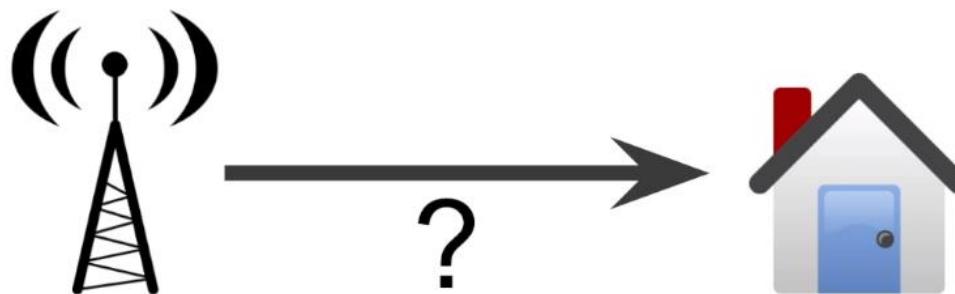
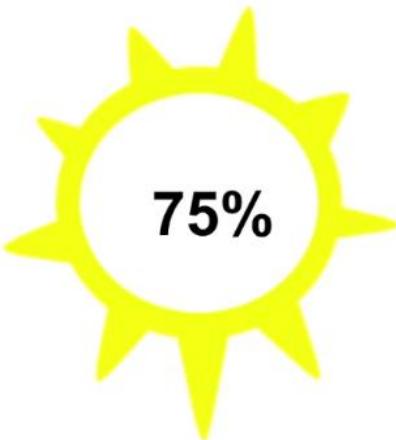
Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Entropy



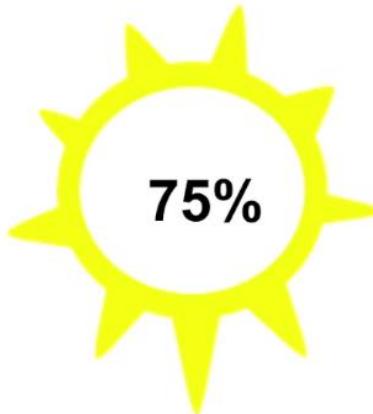
Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Entropy

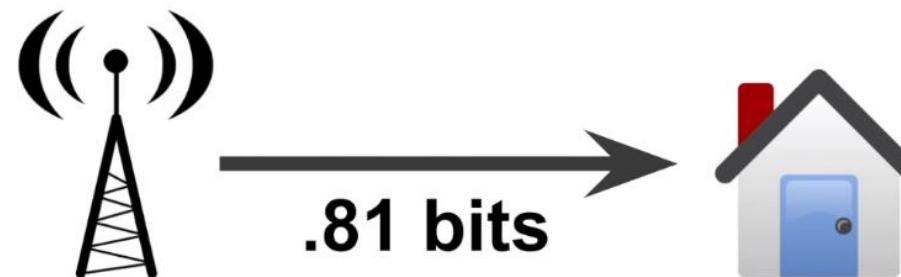


Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Entropy

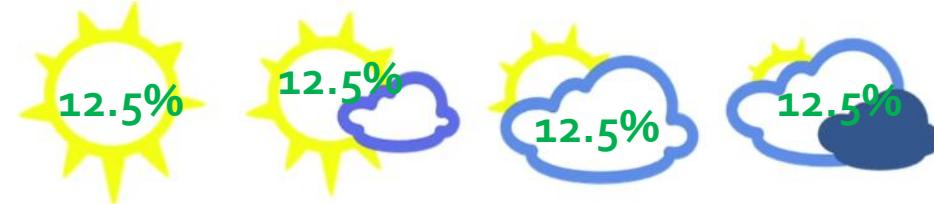


$$\begin{aligned}75\% \times 0.41 \\+ 25\% \times 2 \\= 0.81 \text{ bits}\end{aligned}$$

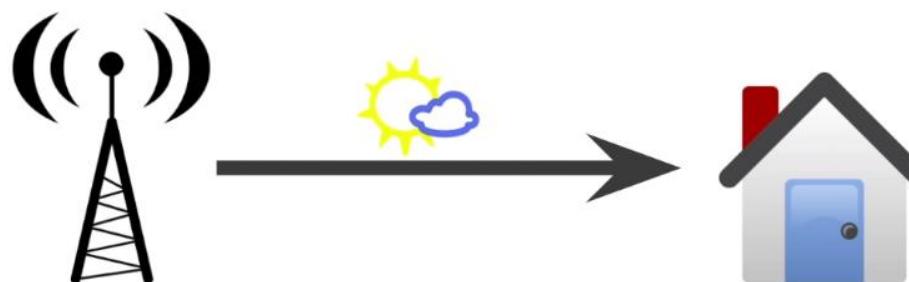


Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Entropy



Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

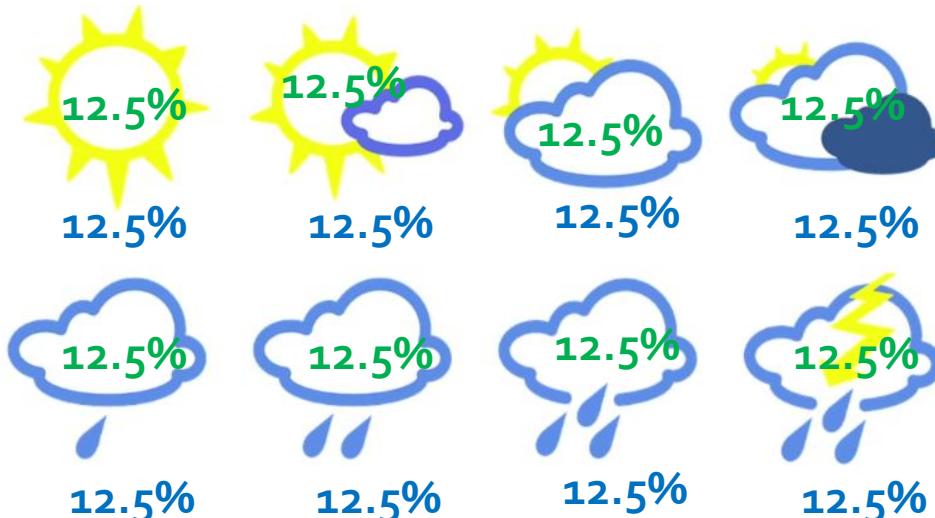


Cross Entropy

Cross Entropy: Average Message Length

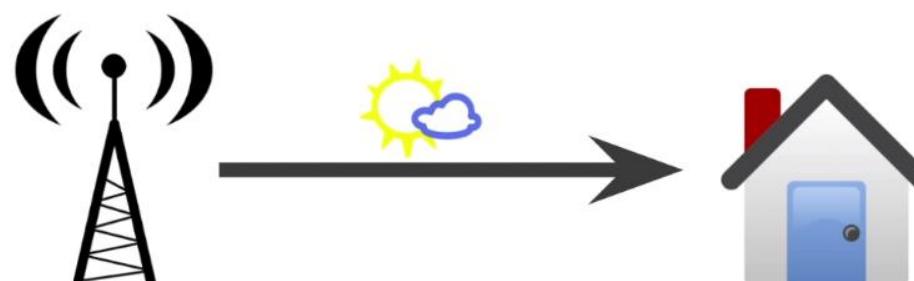
p = true distribution

q = predicted distribution



Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

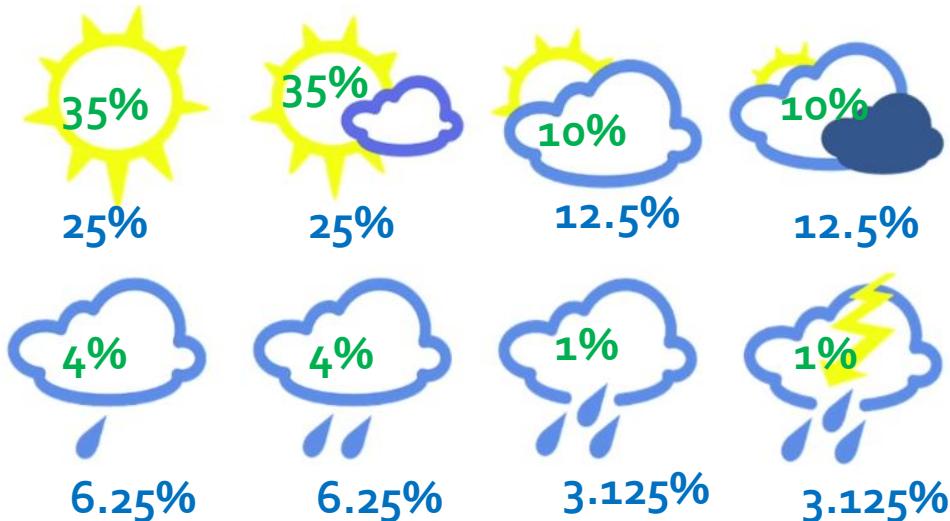
Cross-Entropy:
 $H(p, q) = -\sum_i p_i \log_2(q_i)$



<https://www.youtube.com/watch?v=ErfnhcEV1O8>

Cross Entropy

Decent Weather Report

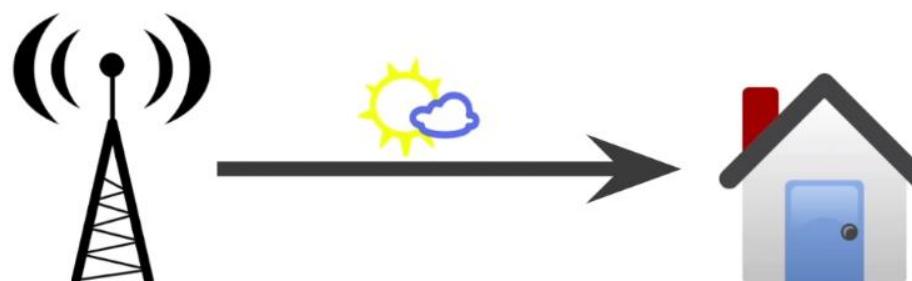


p = true distribution

q = predict distribution

Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

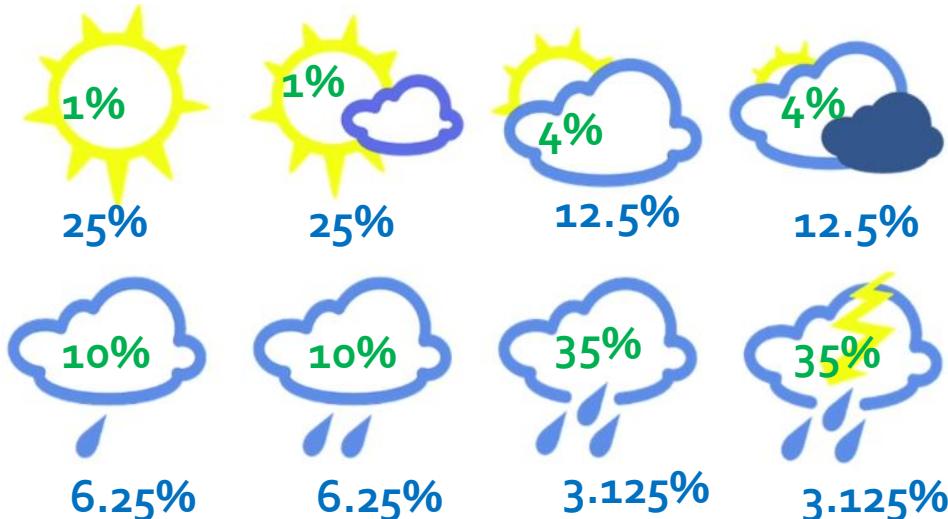
Cross-Entropy:
 $H(p, q) = -\sum_i p_i \log_2(q_i)$



<https://www.youtube.com/watch?v=ErfnhcEV1O8>

Cross Entropy

Bad Weather Report

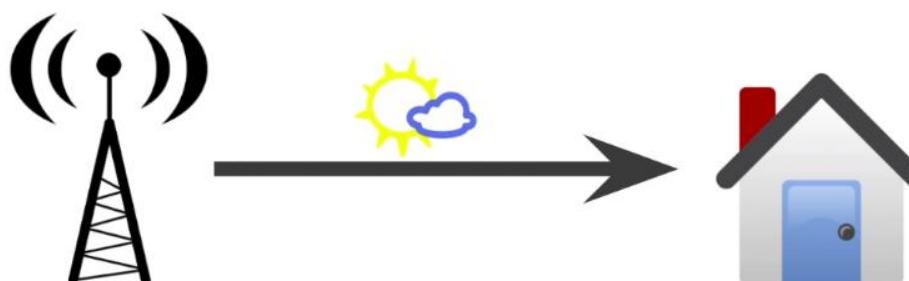


p = true distribution

q = predict distribution

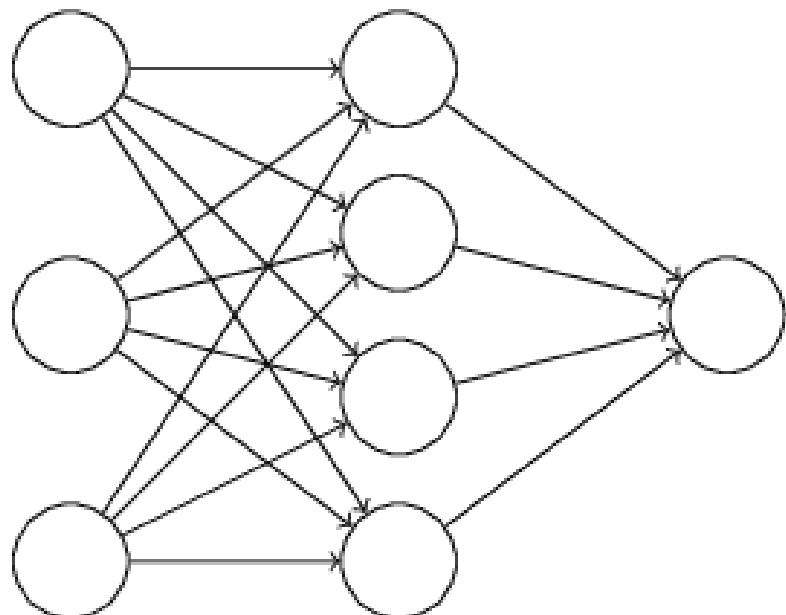
Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Cross-Entropy:
 $H(p, q) = -\sum_i p_i \log_2(q_i)$



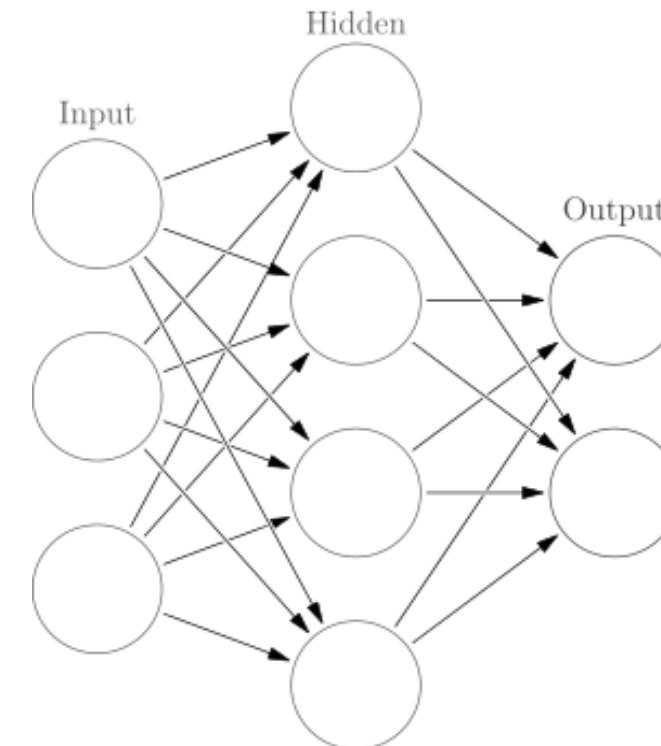
<https://www.youtube.com/watch?v=ErfnhcEV1O8>

How to Code the Output



<http://neuralnetworksanddeeplearning.com/chap1.html>

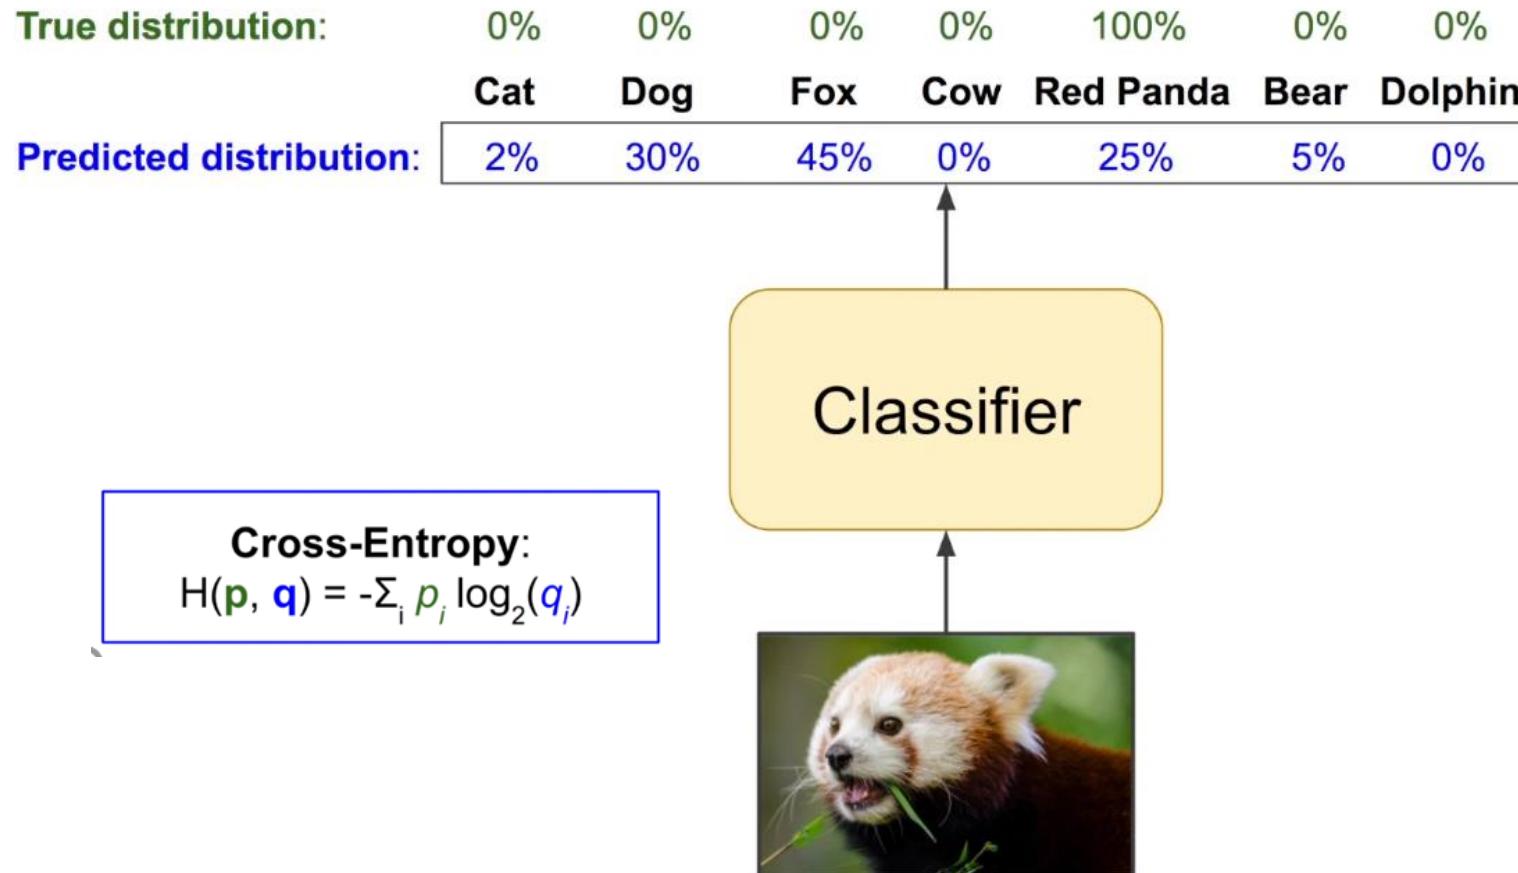
1, 2



https://en.wikipedia.org/wiki/Artificial_neural_network

[1 0],[0 1]

Classification Example



<https://www.youtube.com/watch?v=ErfnhcEV1O8>

Classification Example

True distribution:	0%	0%	0%	0%	100%	0%	0%
	Cat	Dog	Fox	Cow	Red Panda	Bear	Dolphin
Predicted distribution:	2%	30%	45%	0%	25%	5%	0%

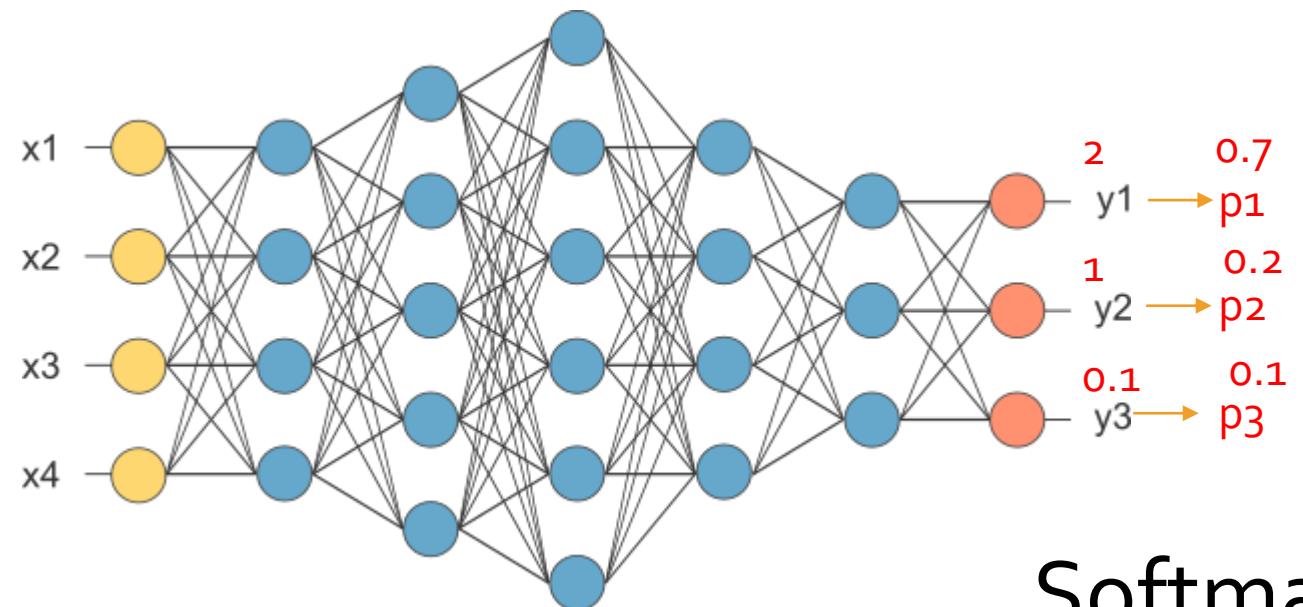
Classifier



Cross-Entropy Loss:
 $H(p, q) = -\sum_i p_i \log(q_i)$
 $= -\log(0.25) = 1.386$

$$\log_2(x) = \log(x) / \log(2)$$

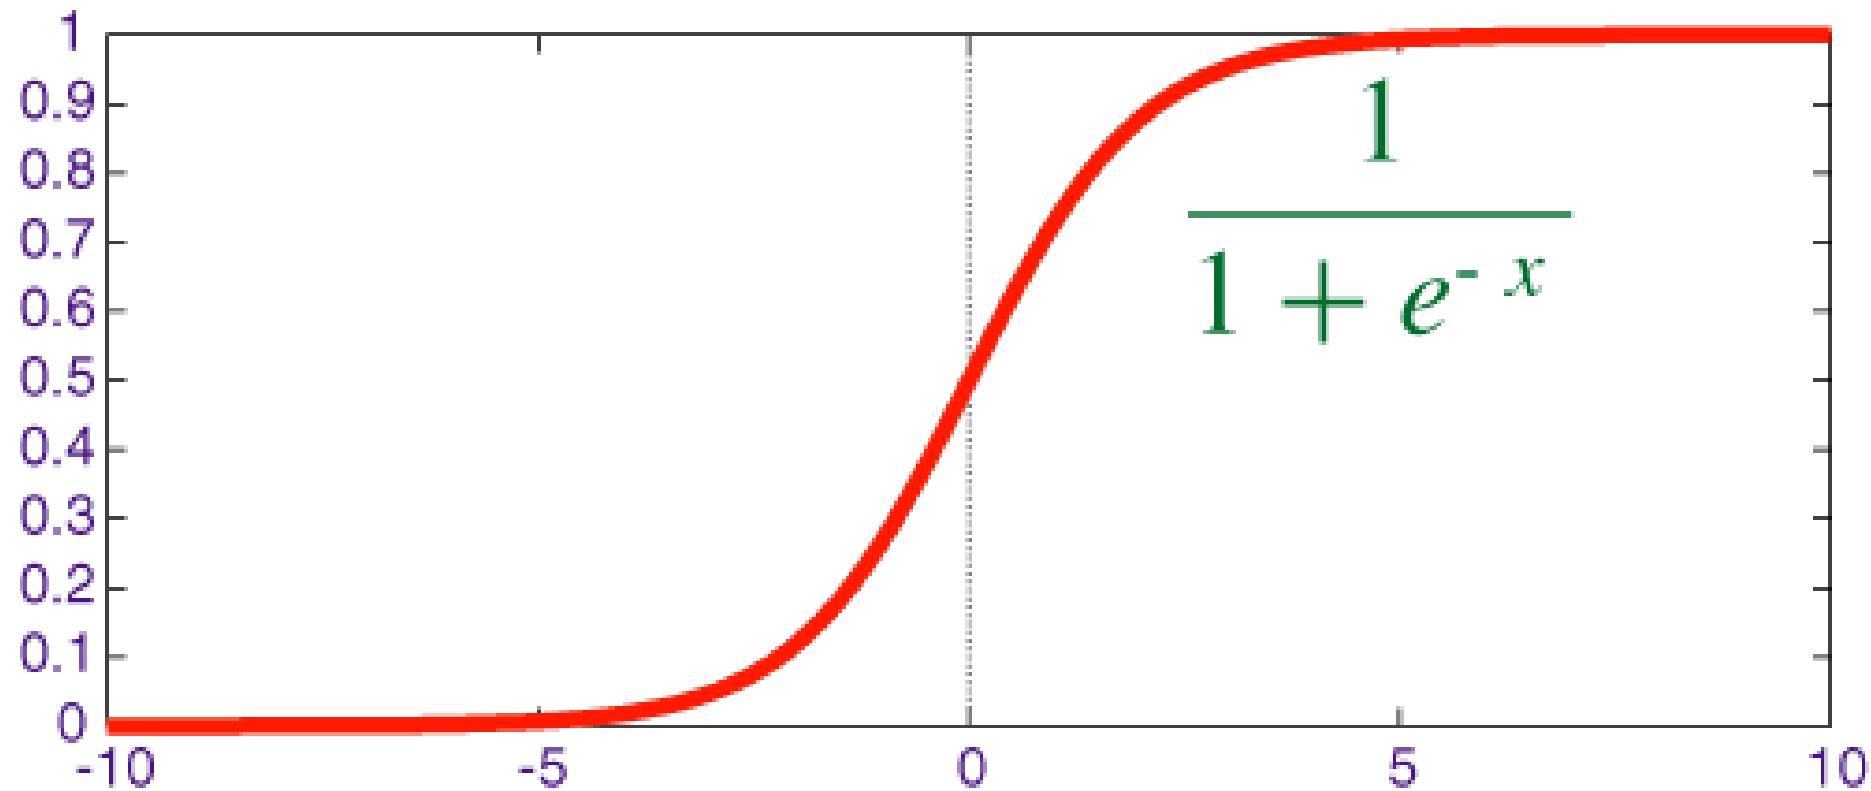
One Step Need



computed	p_1	p_2	p_3	targets	correct?	
0.3	0.3	0.4	0	0	1	yes
0.3	0.4	0.3	0	1	0	yes
0.1	0.2	0.7	1	0	0	no

Softmax

sigmoid

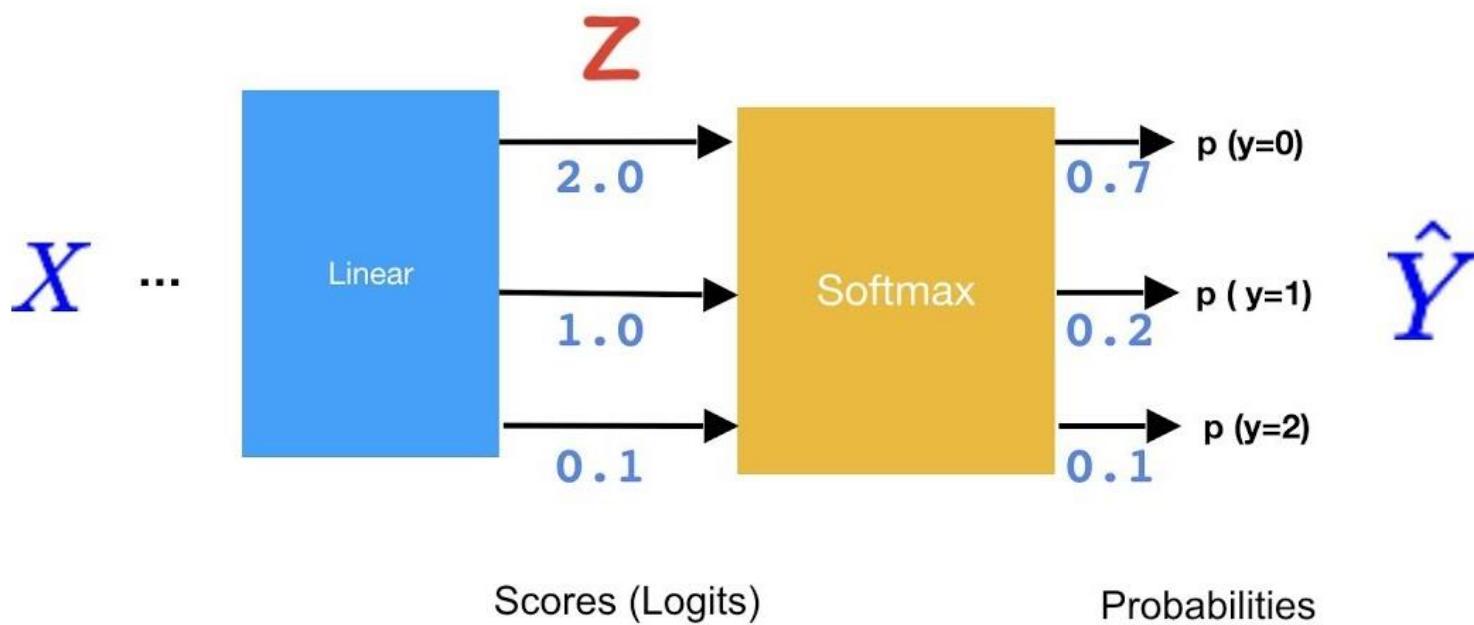


<https://www.kaggle.com/vinay6666/sigmoid-function-for-logisticregression-neuralnets>

Softmax

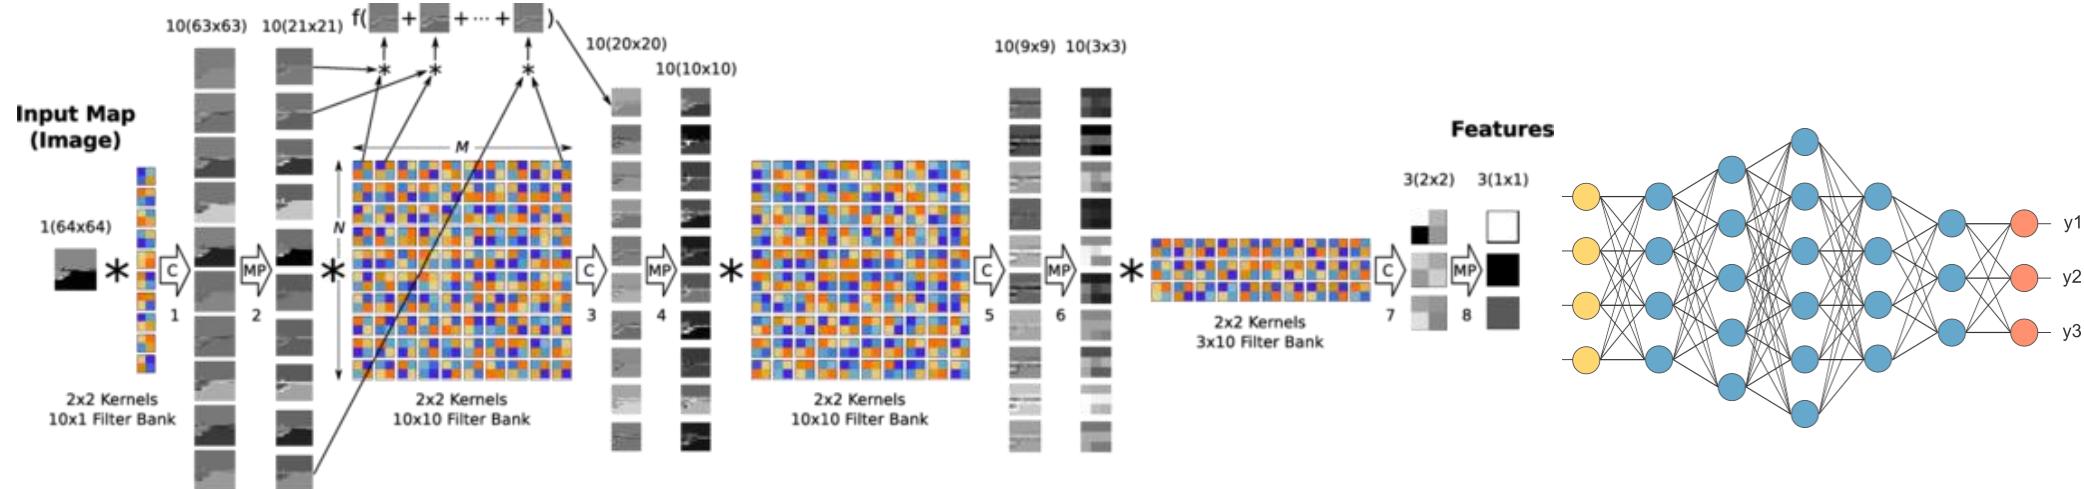
Meet Softmax

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$



<https://www.youtube.com/watch?v=lvNdl7yg4Pg>

Deep Convolutional Network Classification

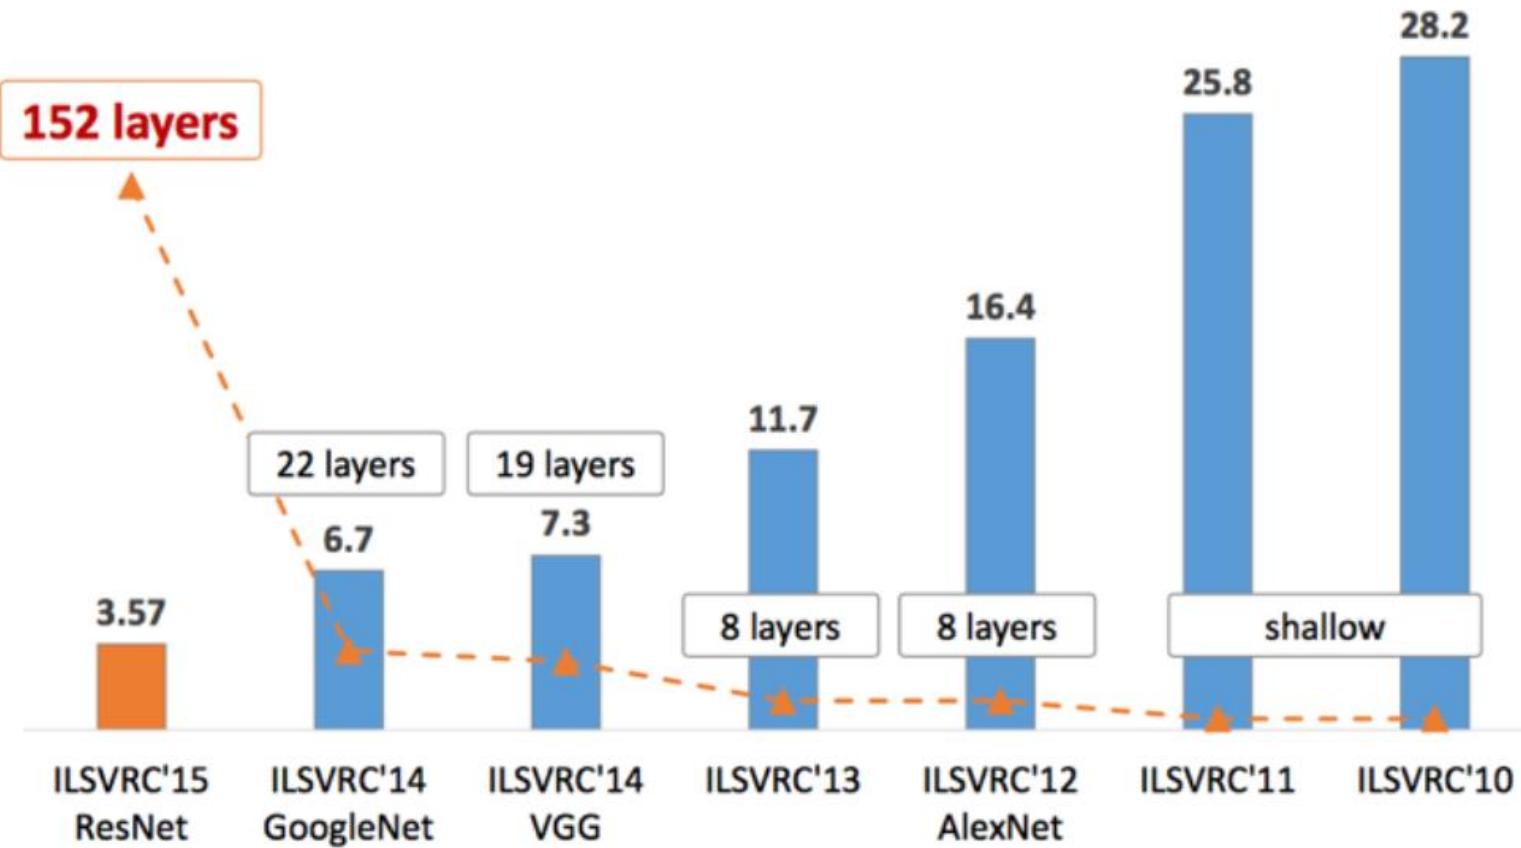


Convolutional Layer

Fully Connected Layer
(Dense Layer)

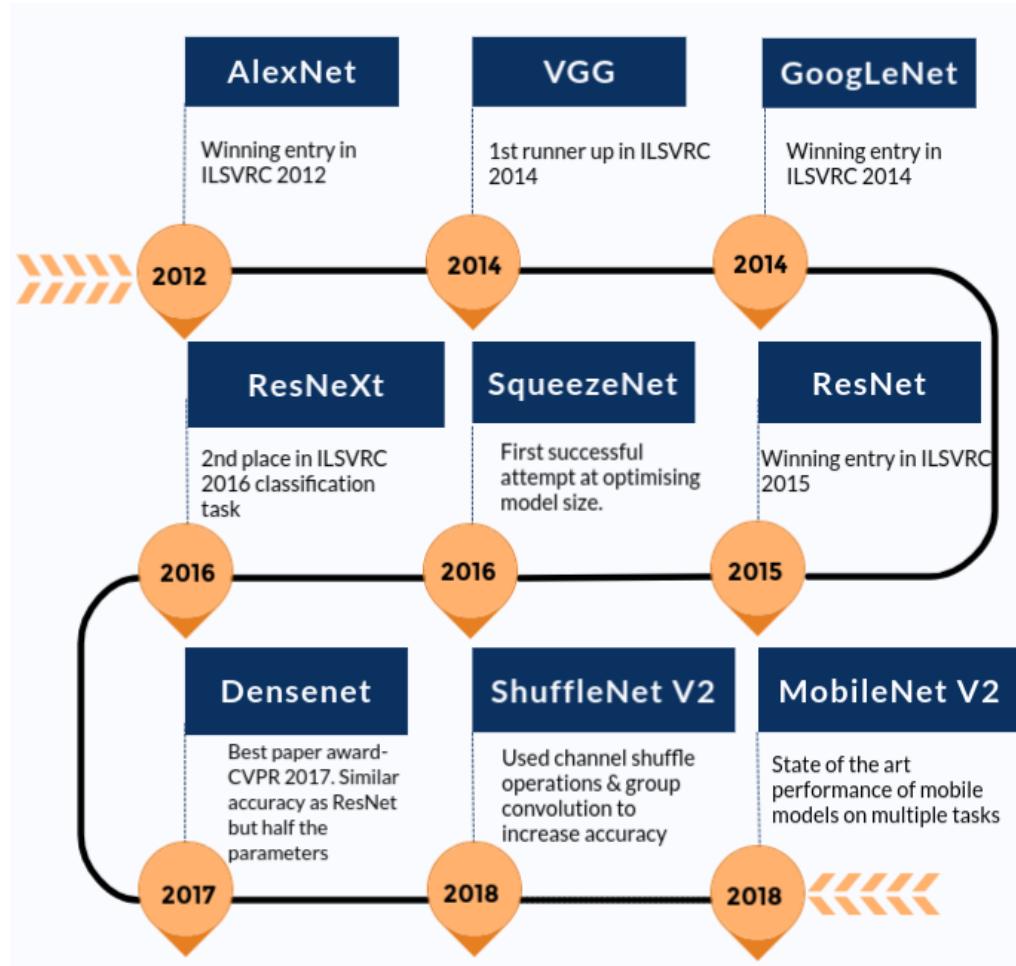
https://www.researchgate.net/figure/Max-Pooling-Convolutional-Neural-Network-MPCNN-with-8-layers-alternating-between_fig3_266656356

Winners



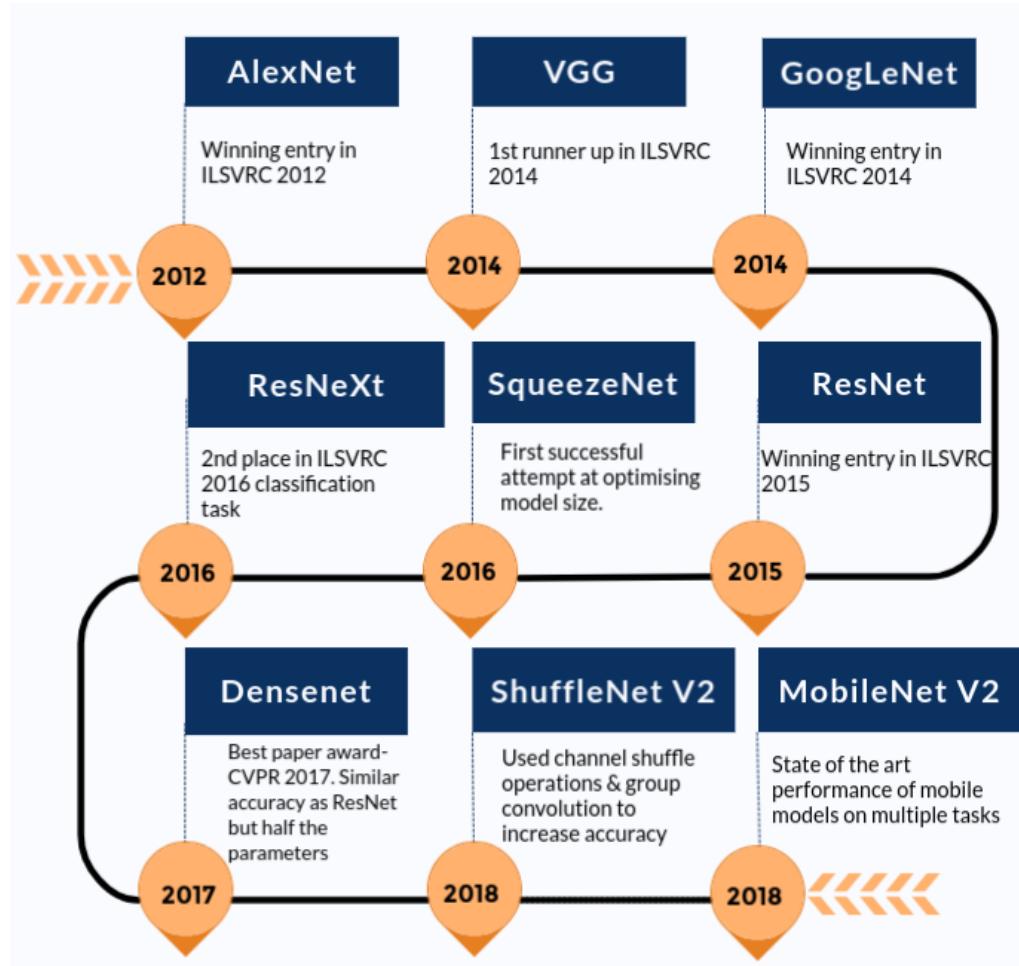
<https://medium.com/@sidereal/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5>

State-of-the-art



<https://www.learnopencv.com/pytorch-for-beginners-image-classification-using-pre-trained-models/>

State-of-the-art



AlexNet

VGG

GoogleNet

ResNet

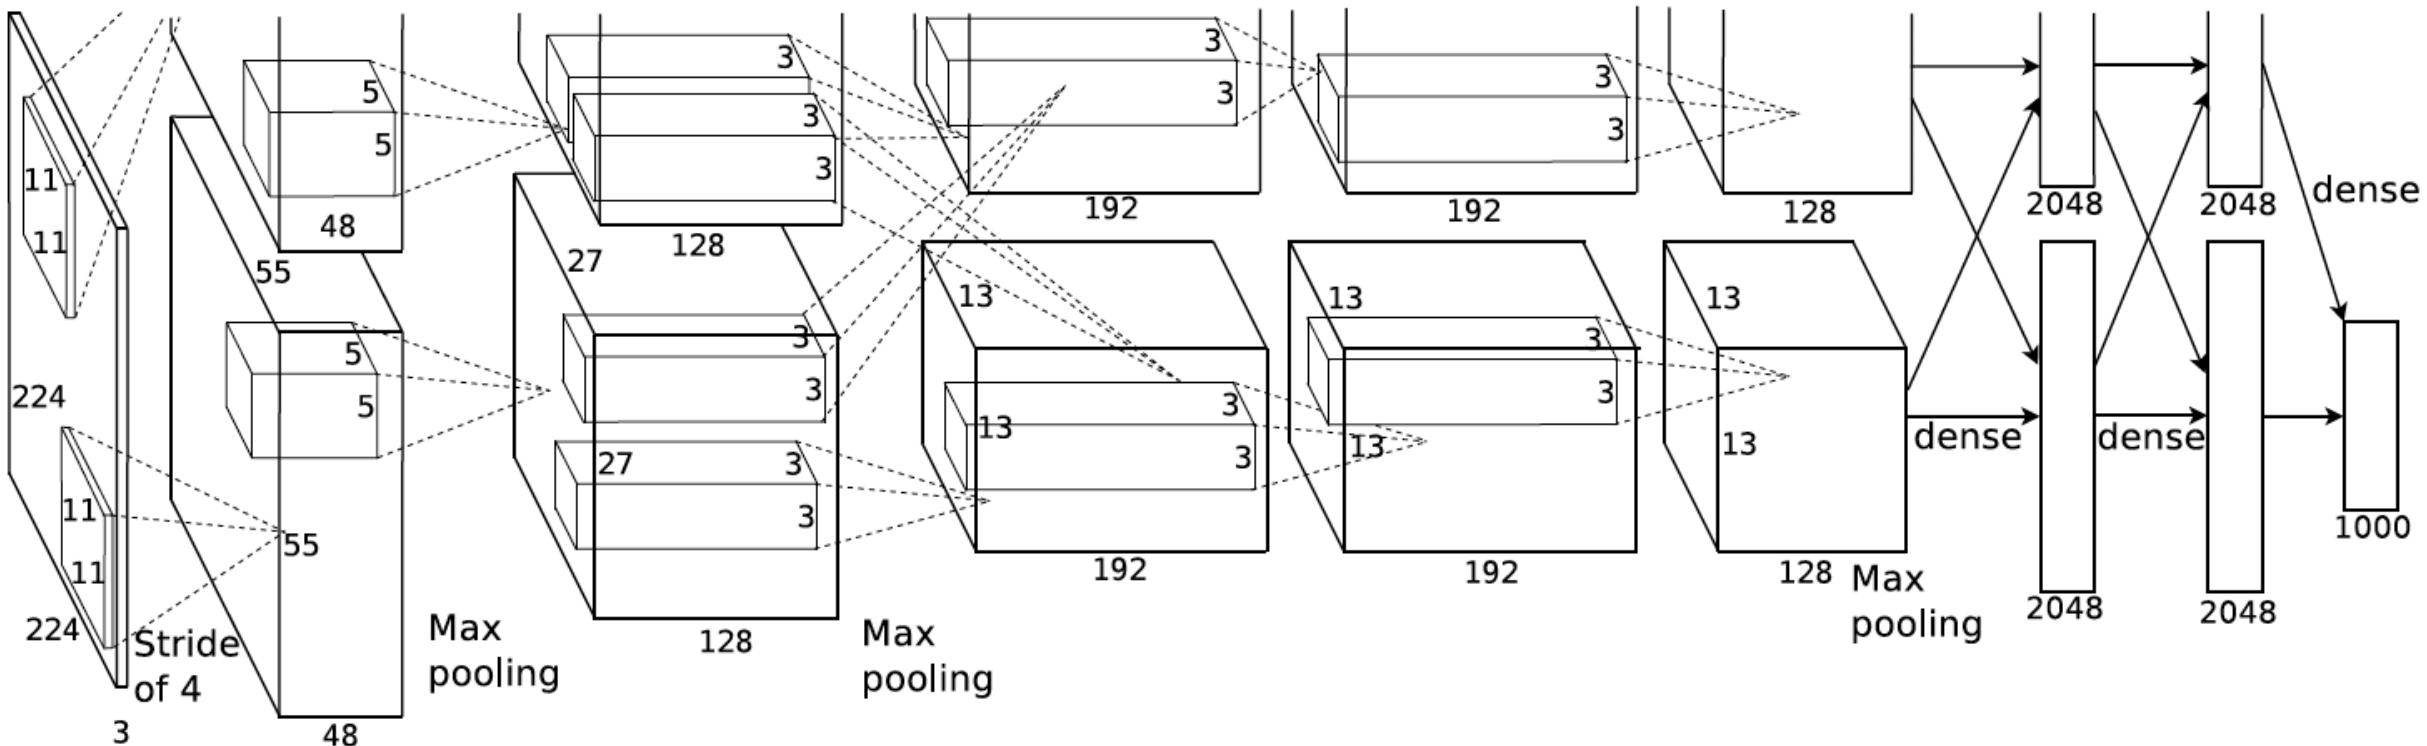
DenseNet

AlexNet

[PDF] ImageNet Classification with Deep Convolutional ... - NIPS Proceedings

<https://papers.nips.cc/.../4824-imagenet-classification-with-deep-convolutional-neural-...> ▾

by A Krizhevsky - 2012 - Cited by 27351 - Related articles



AlexNet

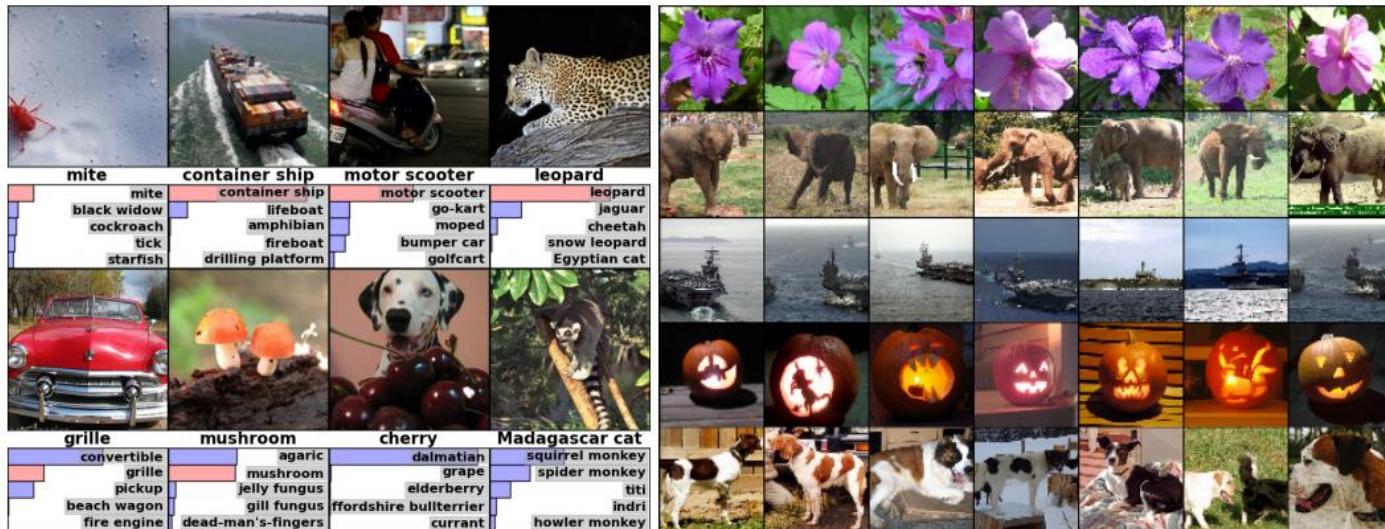


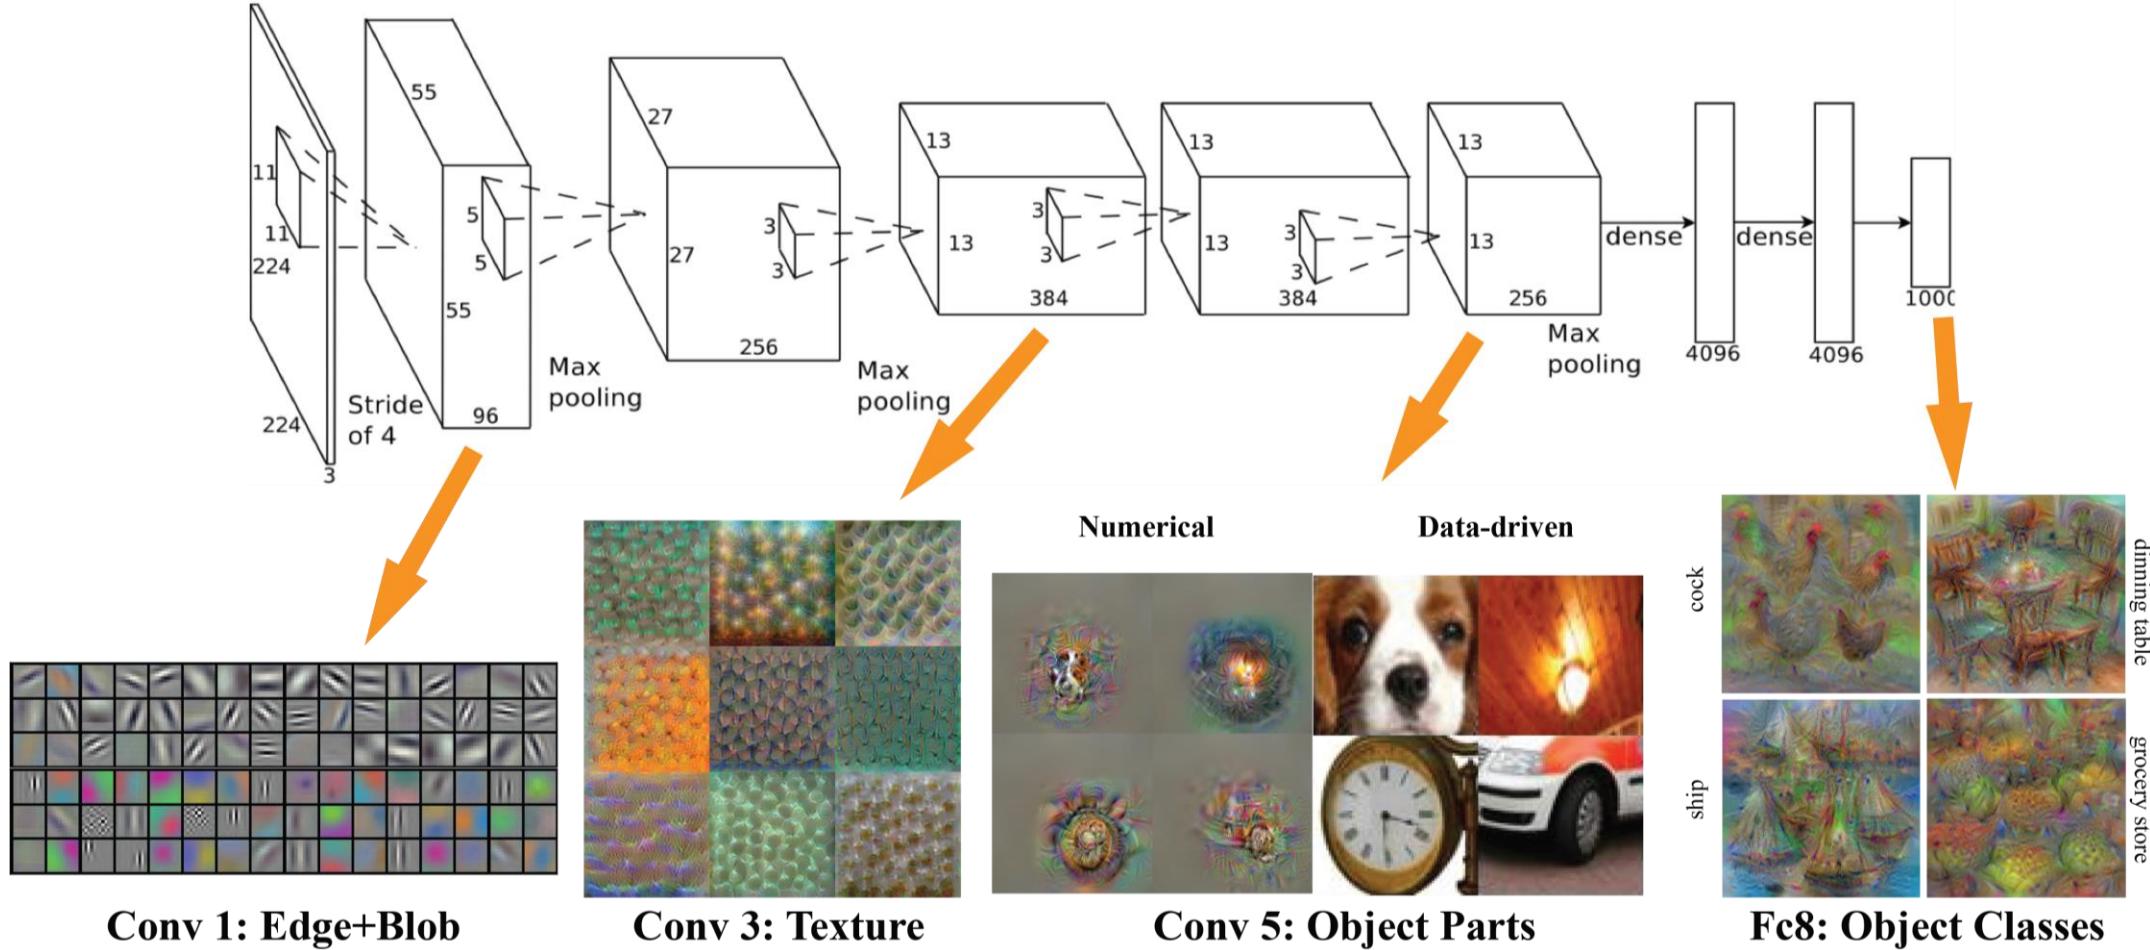
Figure 4: (Left) Eight ILSVRC-2010 test images and the five labels considered most probable by our model. The correct label is written under each image, and the probability assigned to the correct label is also shown with a red bar (if it happens to be in the top 5). (Right) Five ILSVRC-2010 test images in the first column. The remaining columns show the six training images that produce feature vectors in the last hidden layer with the smallest Euclidean distance from the feature vector for the test image.

Krizhevsky et al. 2012

Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs [7]</i>	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	16.4%
1 CNN*	39.0%	16.6%	—
7 CNNs*	36.7%	15.4%	15.3%

Table 2: Comparison of error rates on ILSVRC-2012 validation and test sets. In *italics* are best results achieved by others. Models with an asterisk* were “pre-trained” to classify the entire ImageNet 2011 Fall release. See Section 6 for details.

Full AlexNet



<https://www.cc.gatech.edu/~hays/compvision/proj6/>

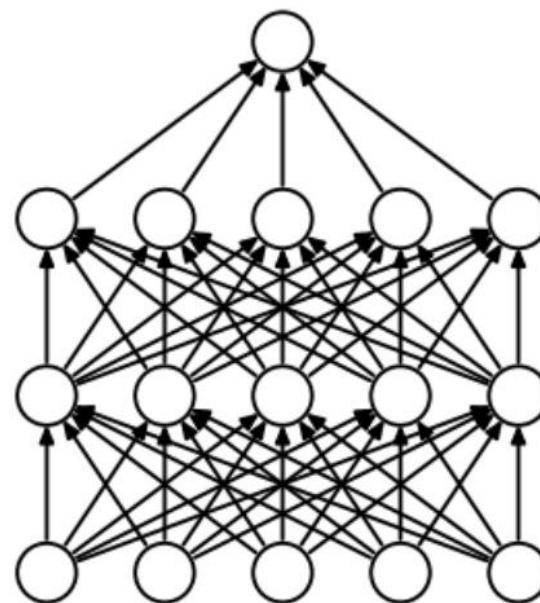
Dropout

Dropout: A Simple Way to Prevent Neural Networks from Overfitting

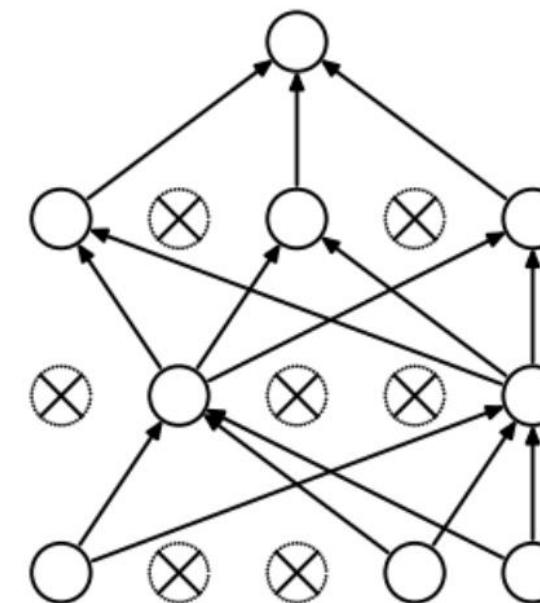
jmlr.org/papers/v15/srivastava14a.html ▾

by N Srivastava - 2014 - Cited by 7597 - Related articles

Dropout is a technique for addressing this problem. The key idea is to randomly drop units (along with their connections) from the neural network during training.

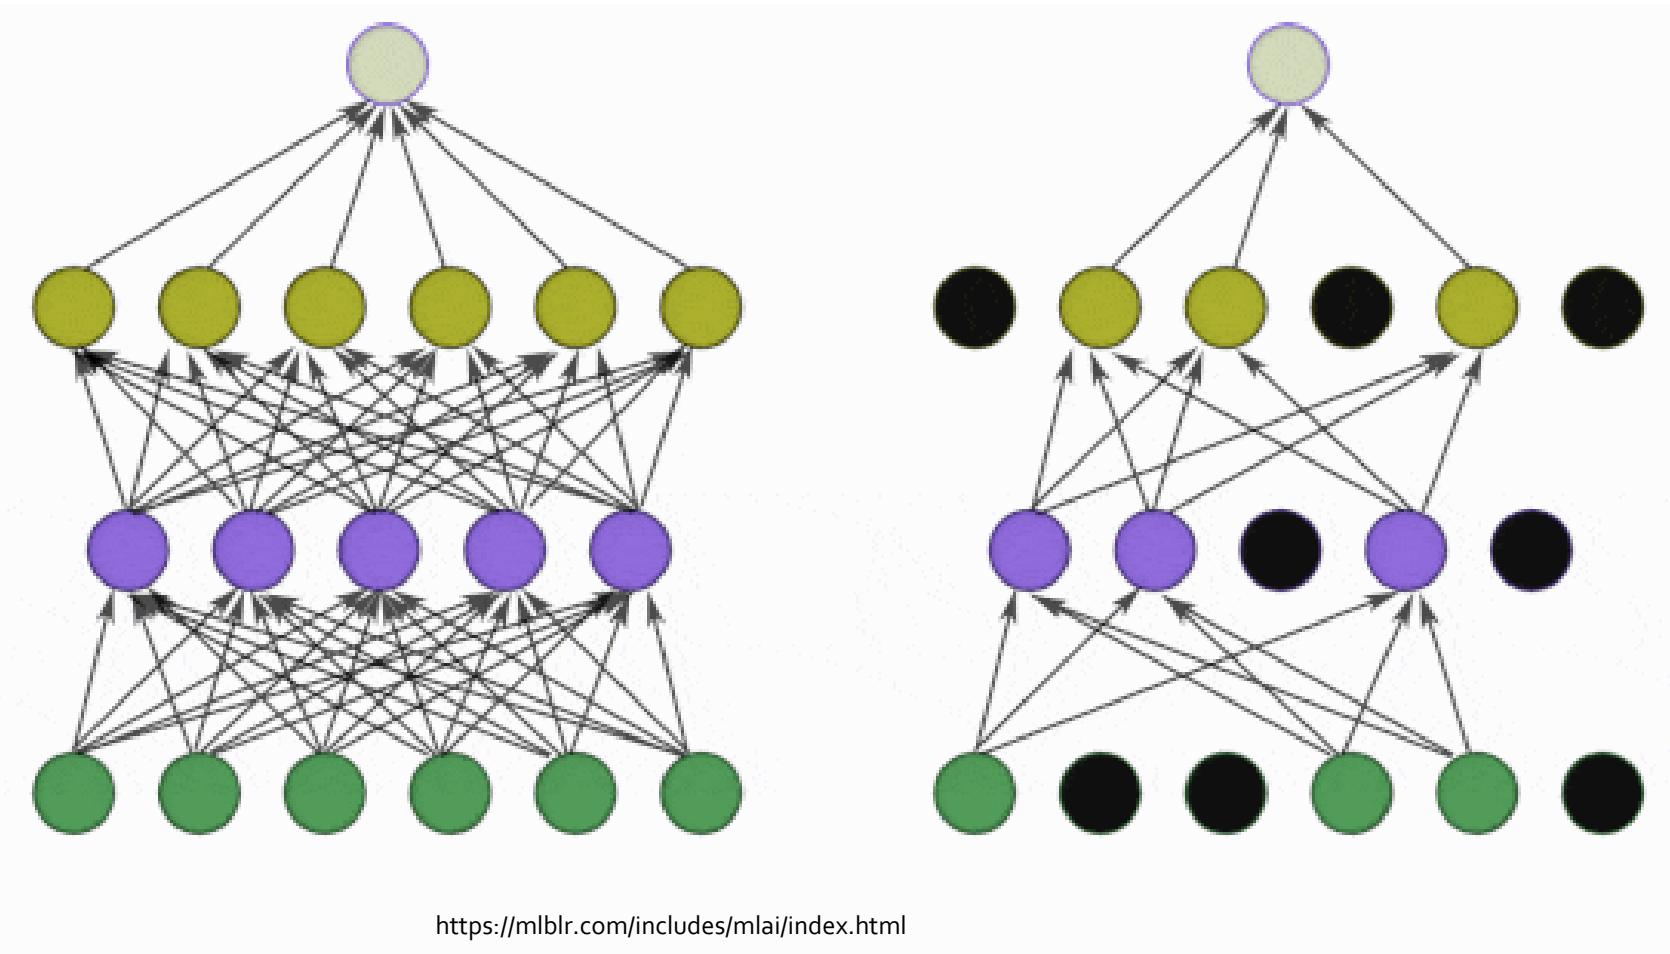


(a) Standard Neural Net

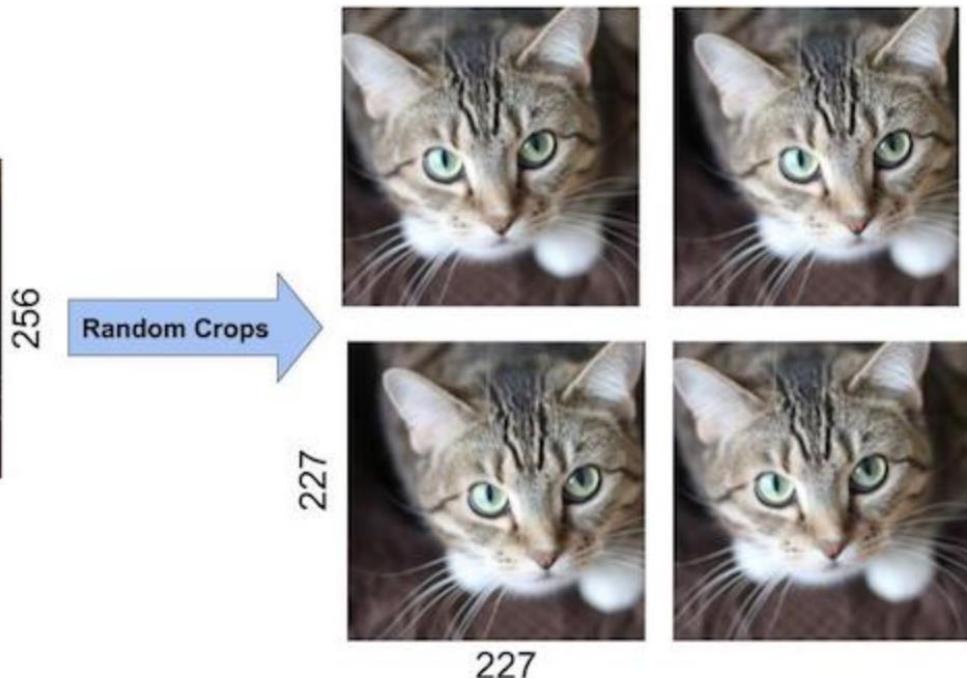
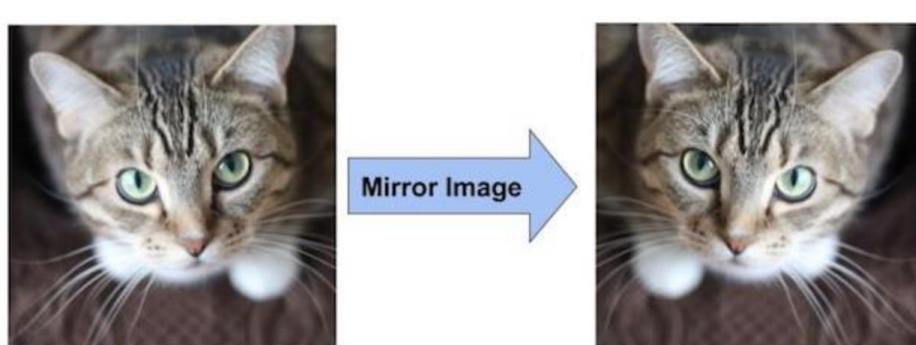


(b) After applying dropout.

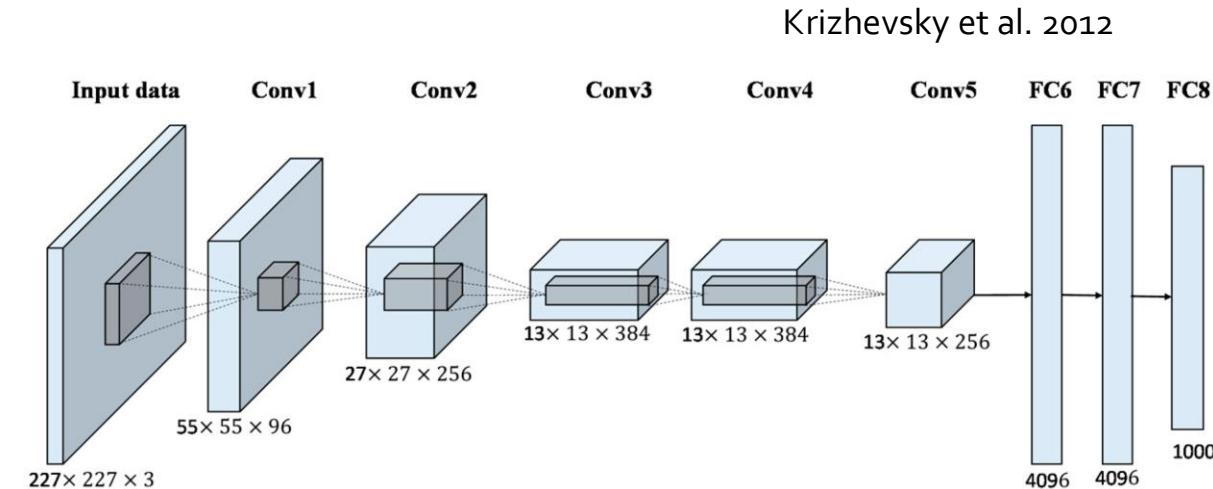
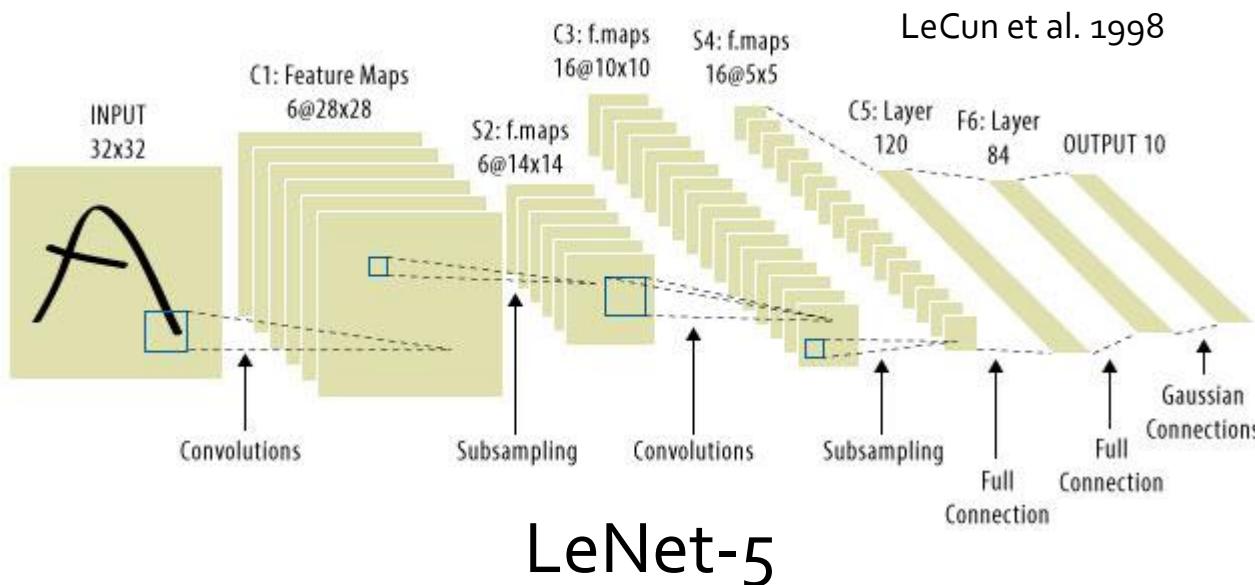
Dropout



Data Augmentation



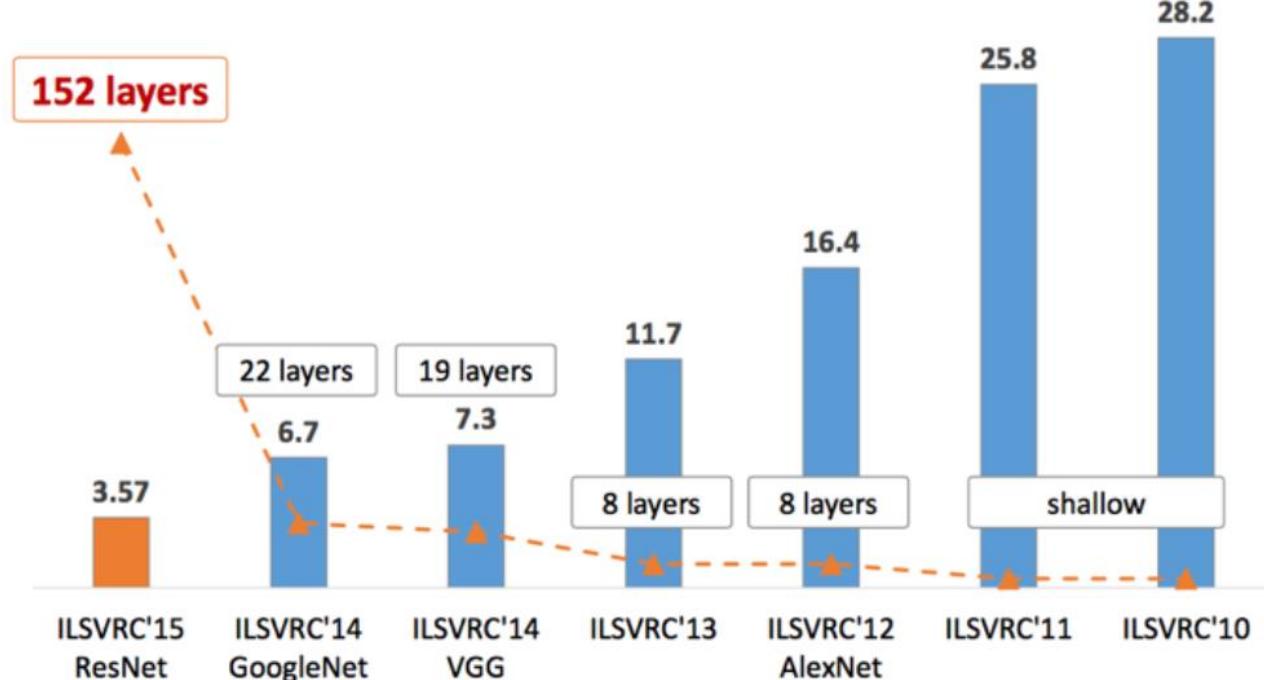
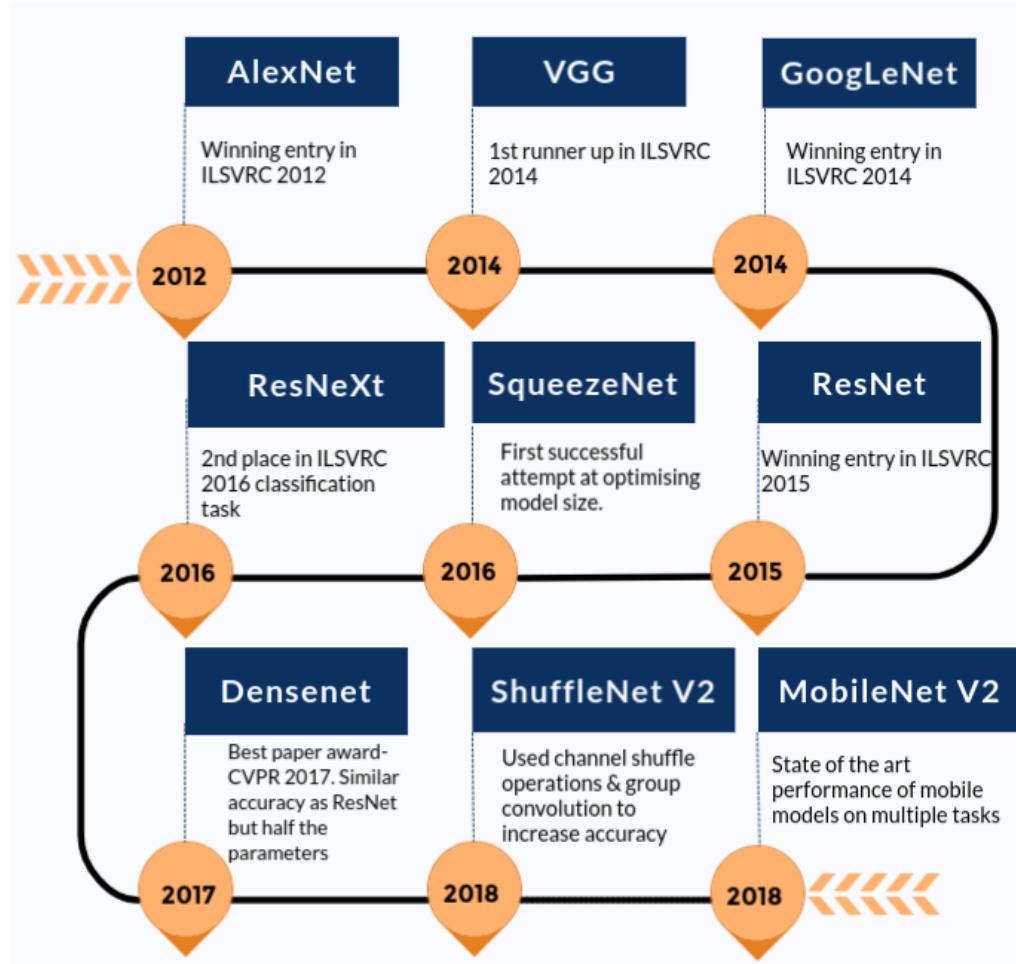
AlexNet vs LeNet-5



The network had a very similar architecture as LeNet by Yann LeCun et al but was deeper, with more filters per layer, and with stacked convolutional layers. It consisted 11×11 , 5×5 , 3×3 , convolutions, max pooling, dropout, data augmentation, ReLU activations, SGD with momentum. It attached ReLU activations after every convolutional and fully-connected layer. AlexNet was trained for 6 days simultaneously on two Nvidia Geforce GTX 580 GPUs which is the reason for why their network is split into two pipelines.

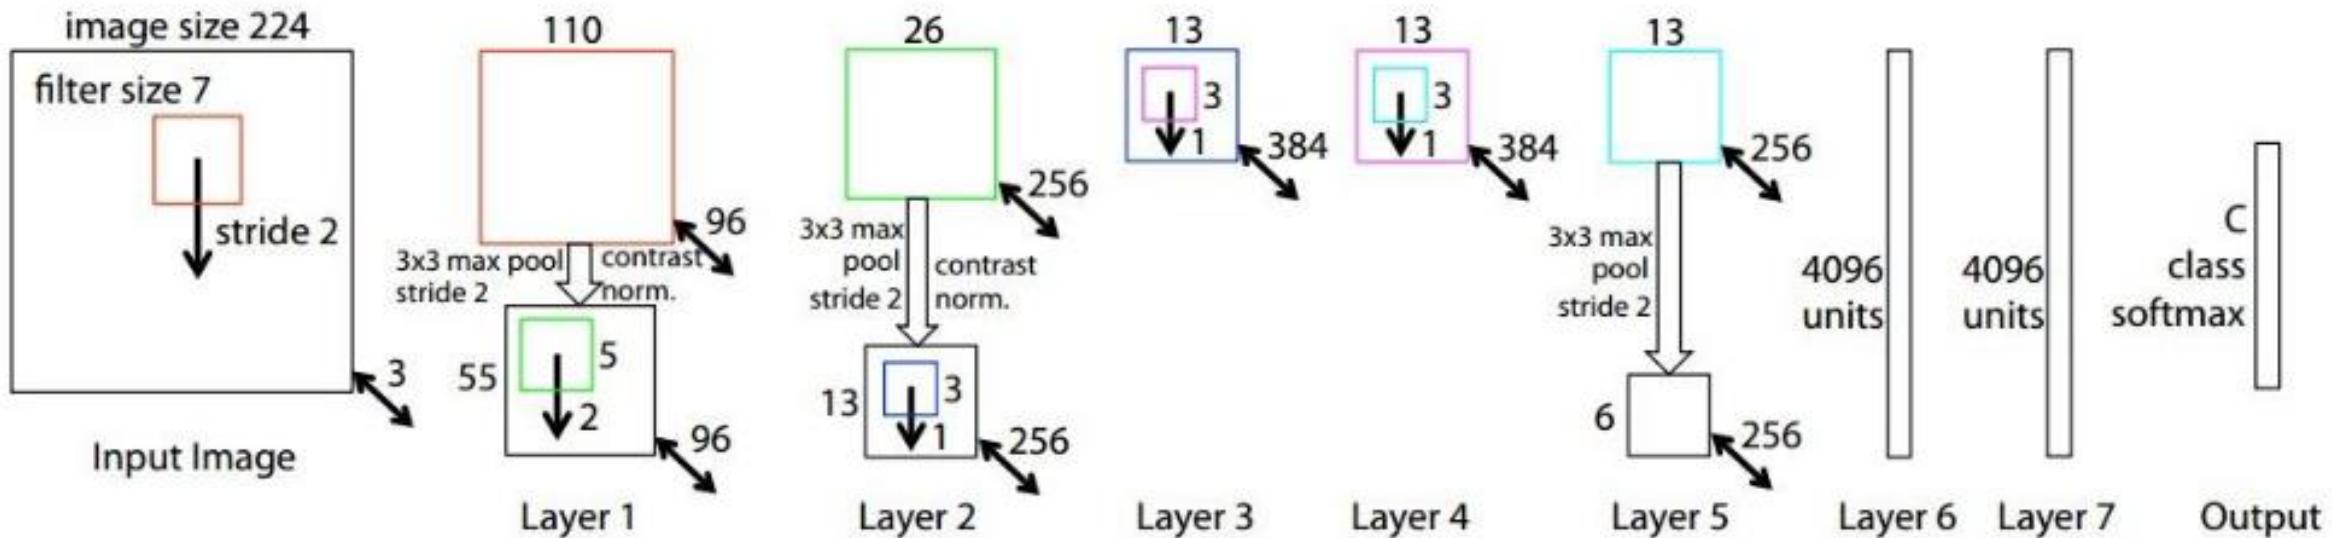
<https://medium.com/@sidereal/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5>

State-of-the-art



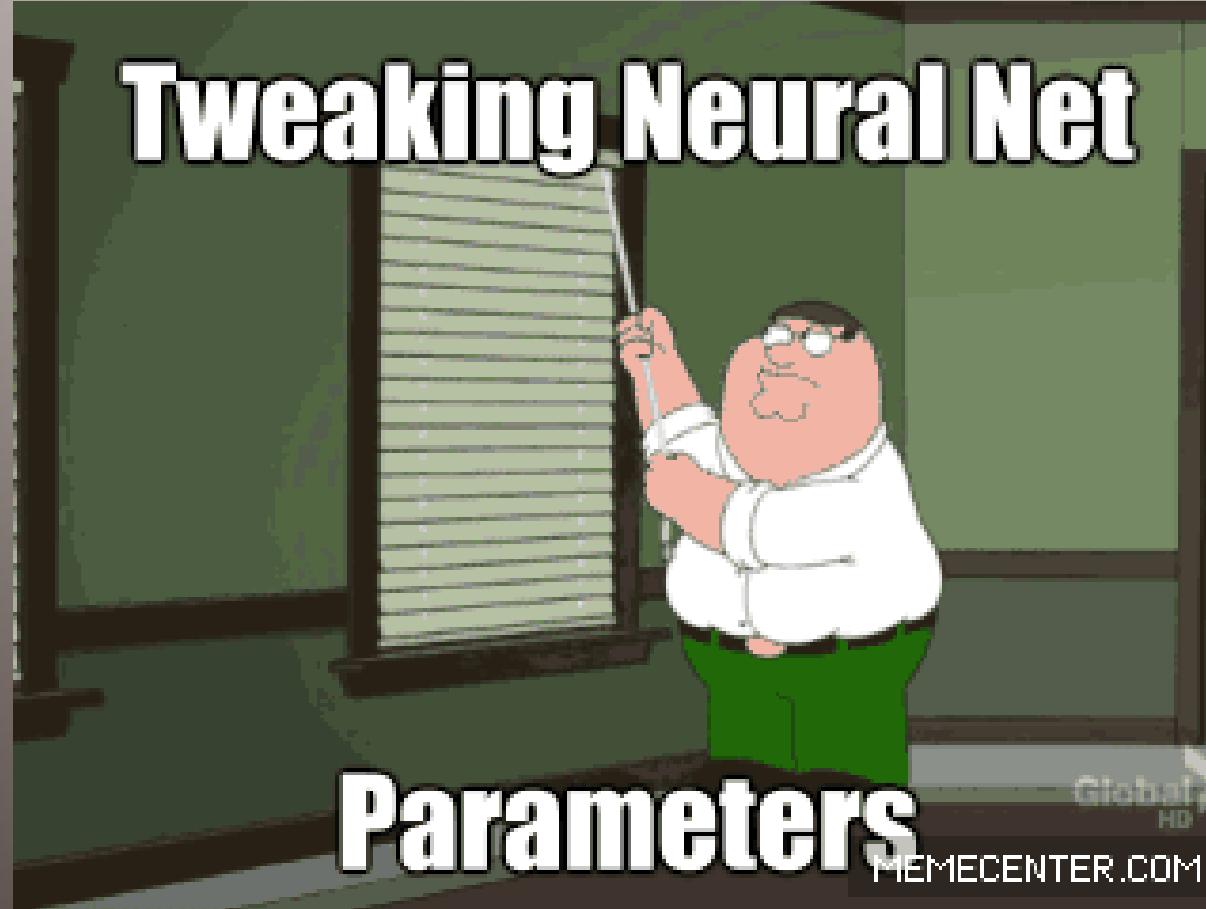
<https://www.learnopencv.com/pytorch-for-beginners-image-classification-using-pre-trained-models/>

ZFNet (2013)



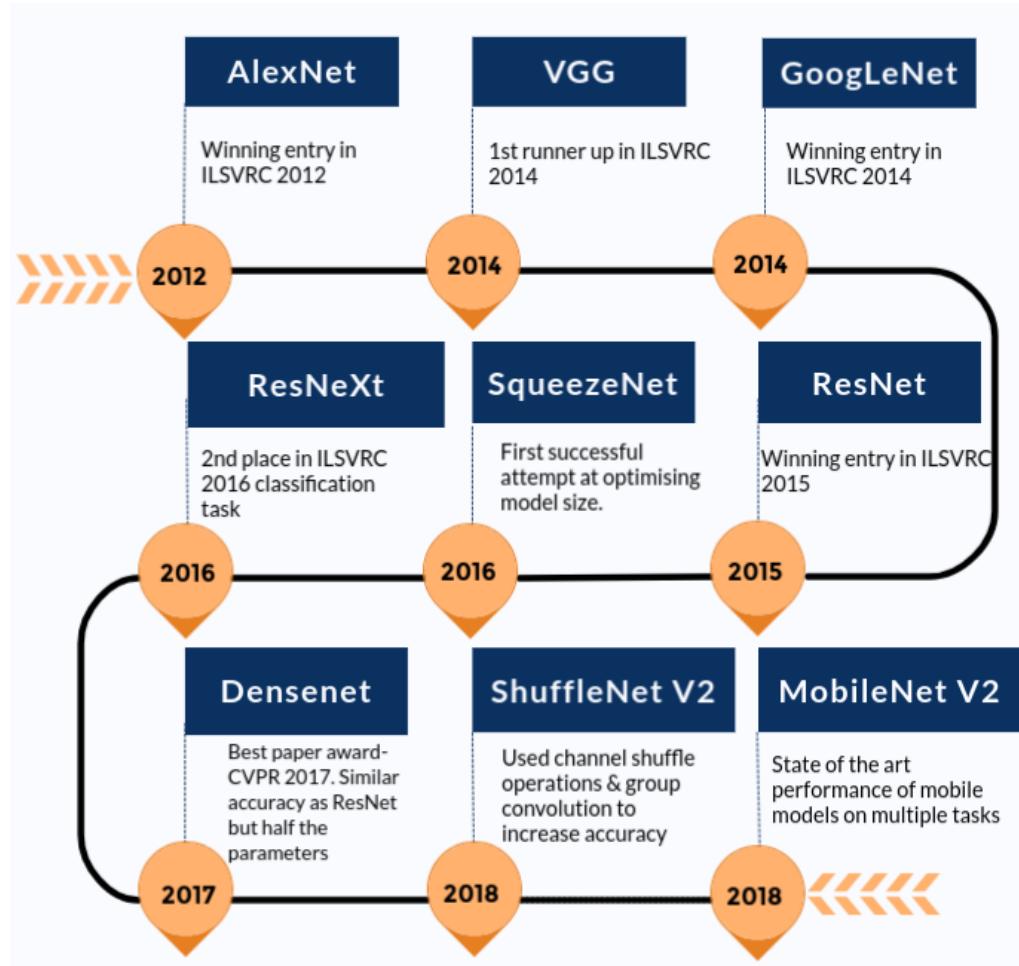
Not surprisingly, the ILSVRC 2013 winner was also a CNN which became known as ZFNet. It achieved a top-5 error rate of 14.8% which is now already half of the prior mentioned non-neural error rate. It was mostly an achievement by tweaking the hyper-parameters of AlexNet while maintaining the same structure with additional Deep Learning elements as discussed earlier in this essay.

ZFNet : That's it



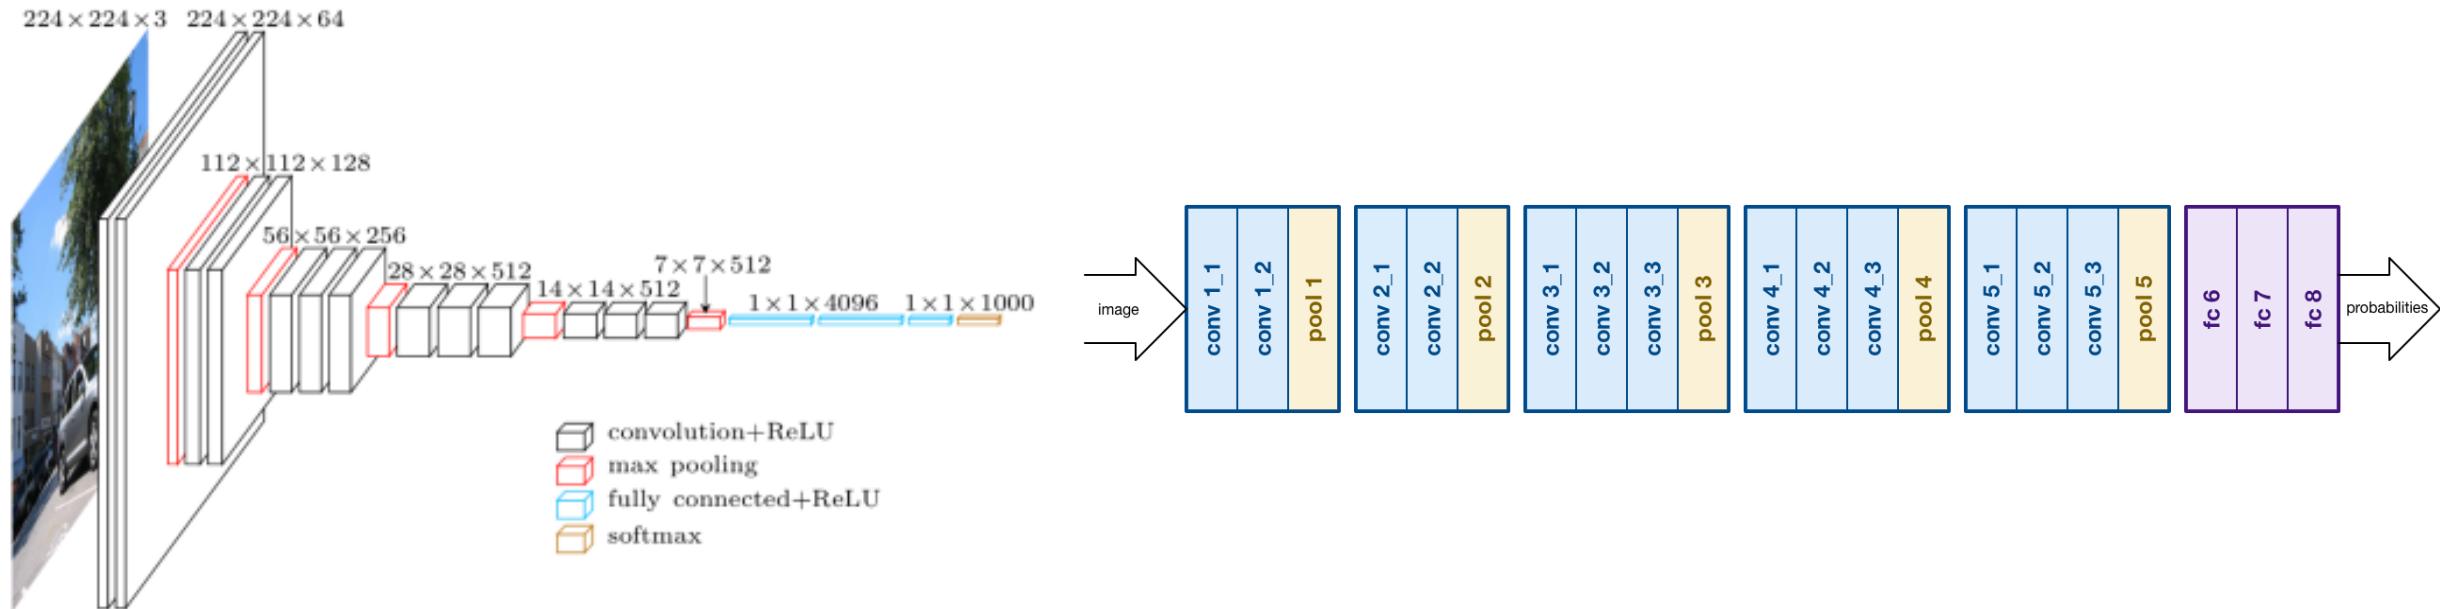
<http://bbabenko.tumblr.com/post/83319141207/convolutional-learnings-things-i-learned-by>

State-of-the-art



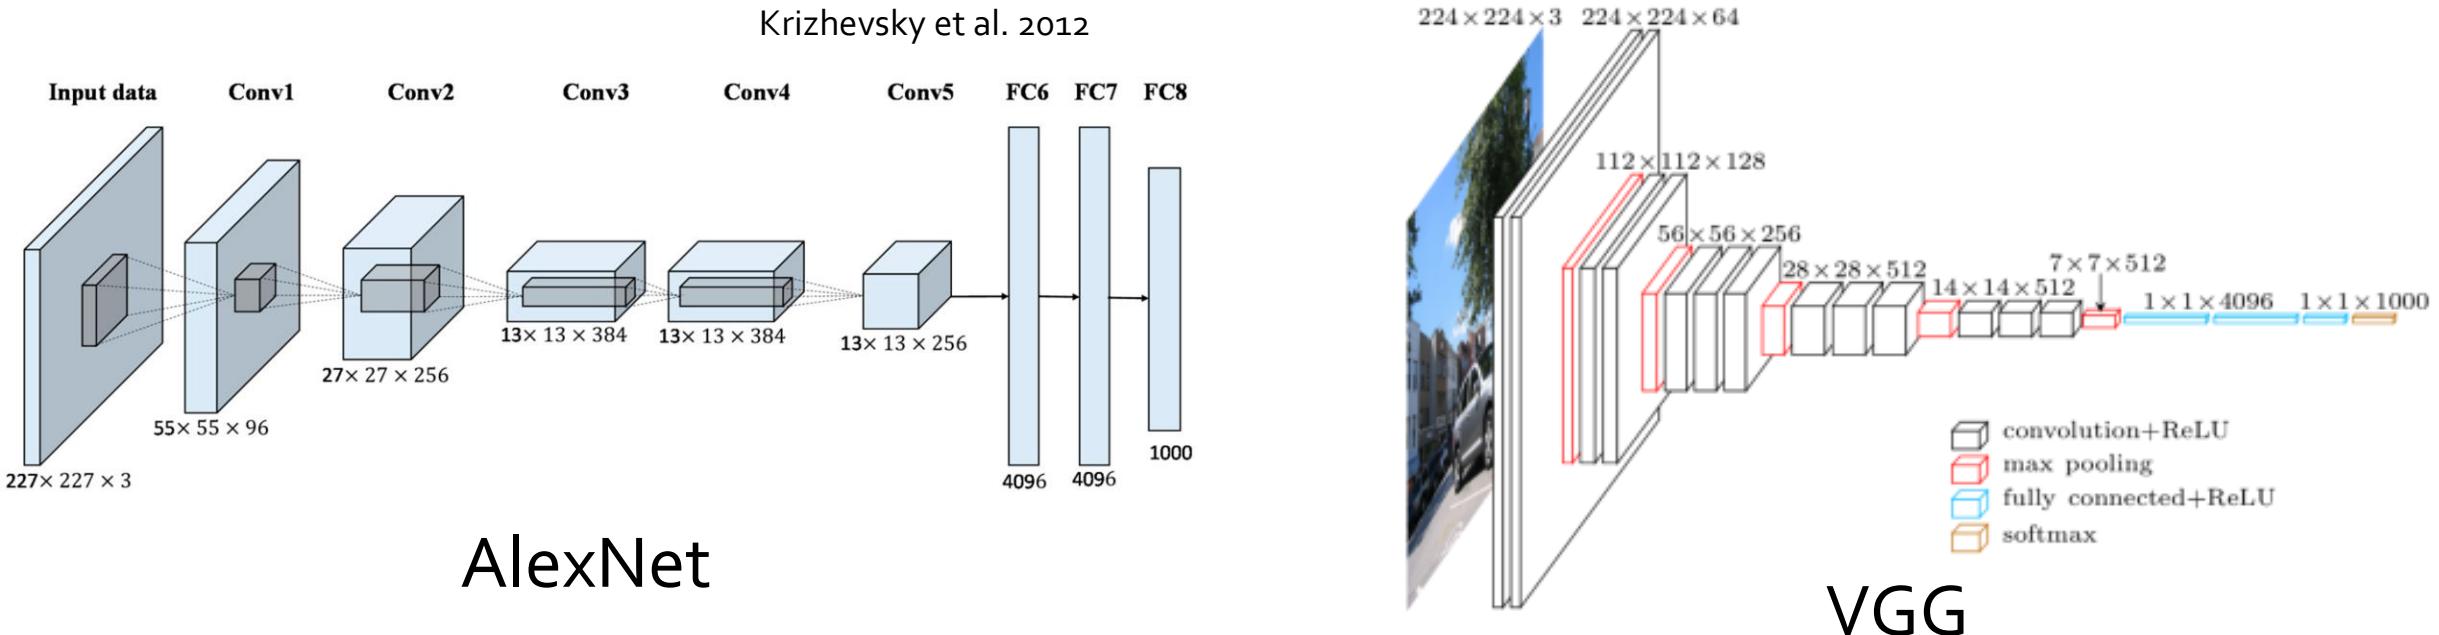
AlexNet
VGG
GoogleNet
ResNet
DenseNet

VGGNet (2014)



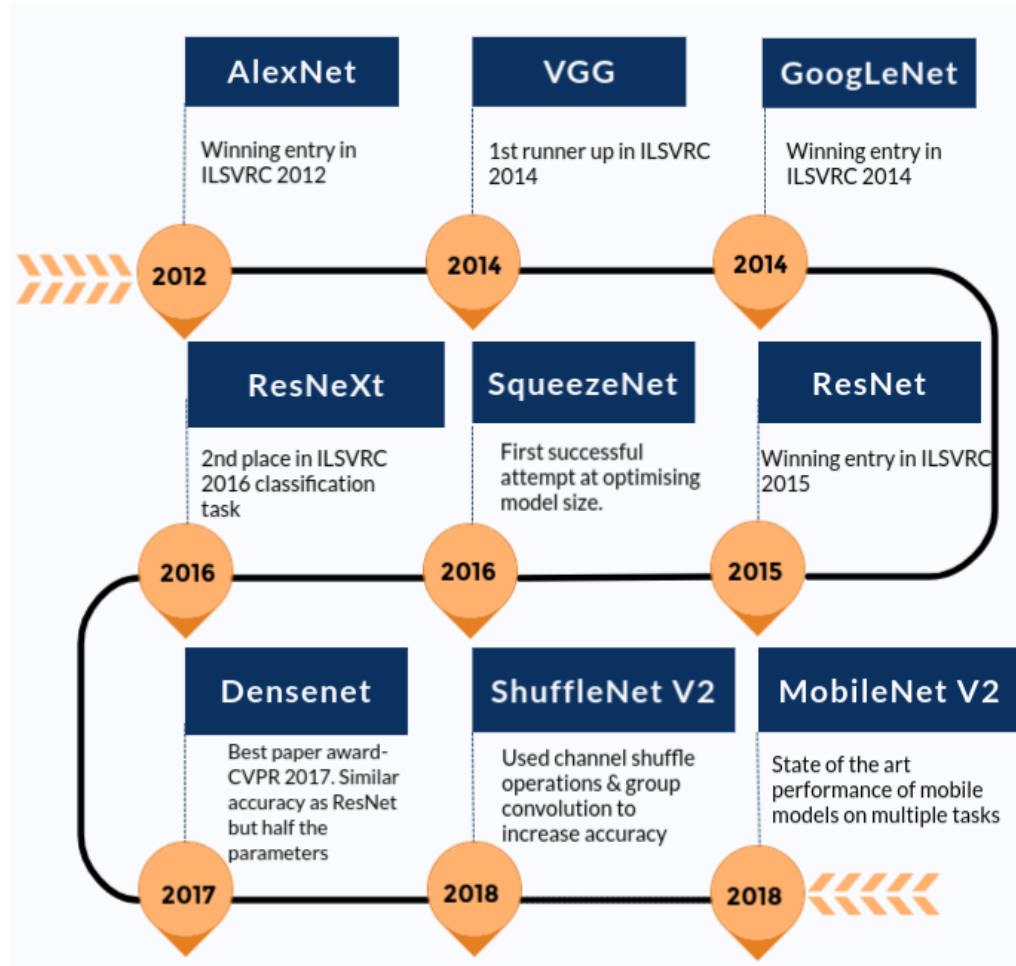
<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>
<https://medium.com/@sidereal/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5>

AlexNet vs VGG



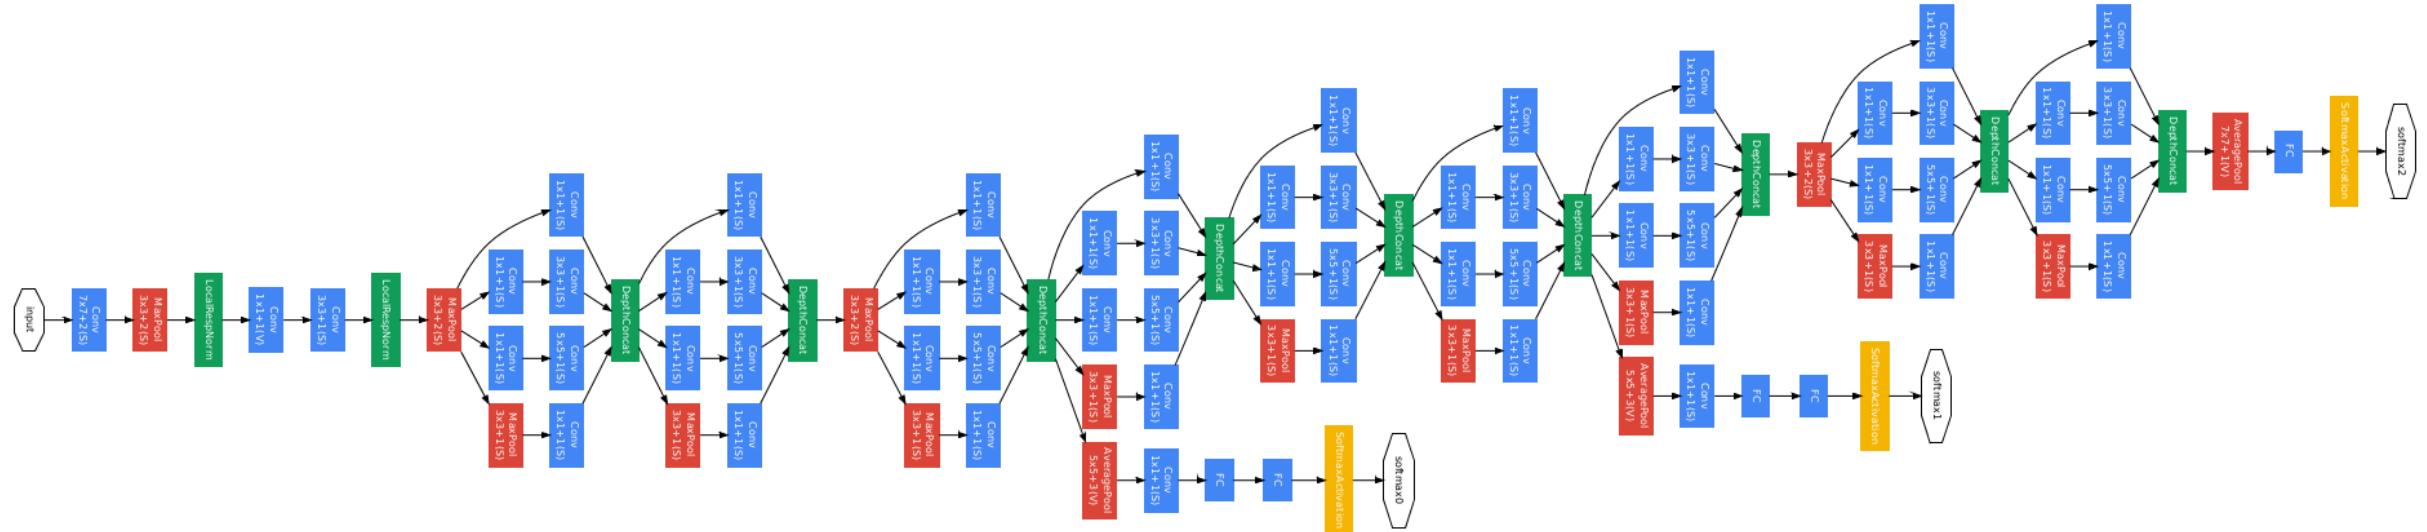
The idea behind having fixed size kernels is that all the variable size convolutional kernels used in Alexnet (11×11 , 5×5 , 3×3) can be replicated by making use of multiple 3×3 kernels as building blocks. The replication is in terms of the receptive field covered by the kernels. This reduces the number of trainable variables by 44.9% (62.8%). Reduced number of trainable variables means faster learning and more robust to overfitting. max poolings should be placed after each 2 convolutions and the number of filters should be doubled after each max-pooling.

State-of-the-art



AlexNet
VGG
GoogleNet
ResNet
DenseNet

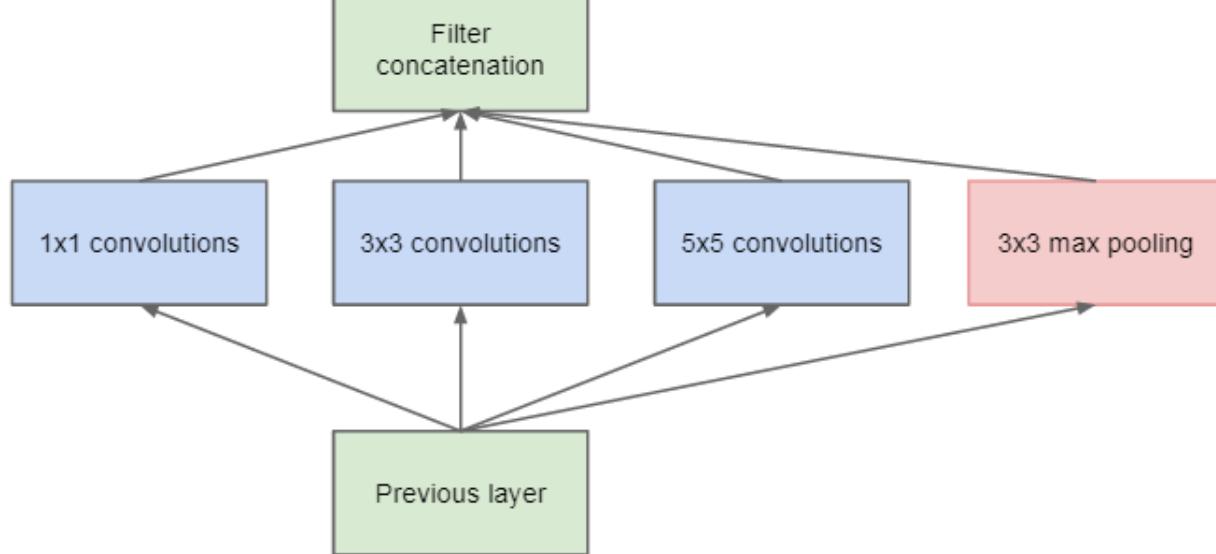
GoogleNet/Inception(2014)



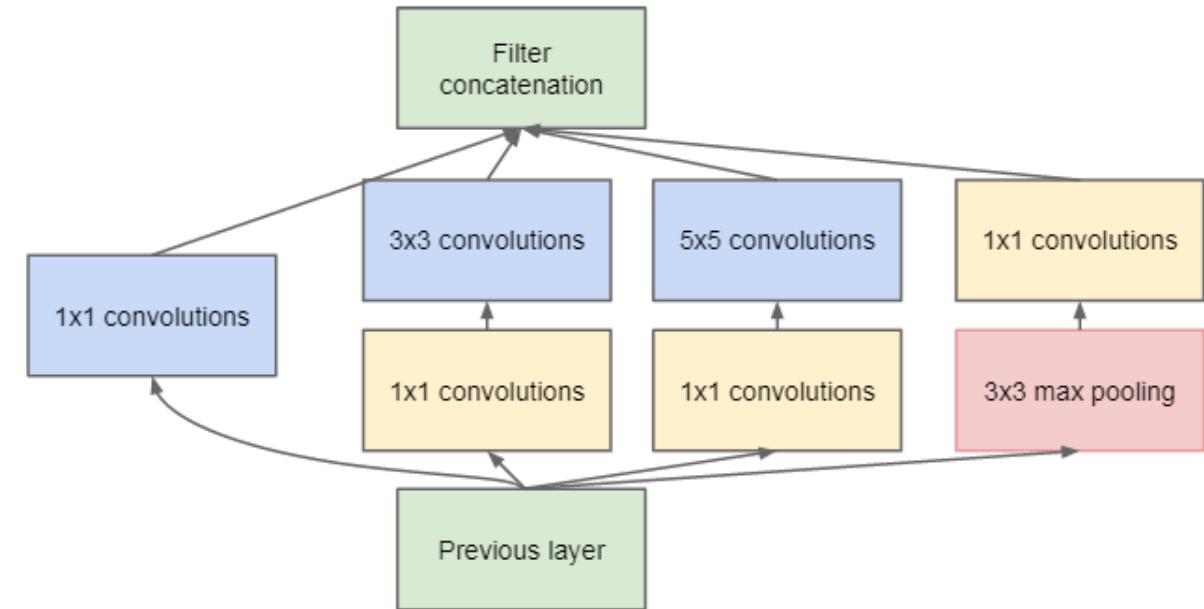
The winner of the ILSVRC 2014 competition was GoogleNet(a.k.a. Inception V1) from Google. It achieved a top-5 error rate of 6.67%! This was very close to human level performance which the organisers of the challenge were now forced to evaluate. As it turns out, this was actually rather hard to do and required some human training in order to beat GoogLeNets accuracy. After a few days of training, the human expert (Andrej Karpathy) was able to achieve a top-5 error rate of 5.1%(single model) and 3.6%(ensemble). The network used a CNN inspired by LeNet but implemented a novel element which is dubbed an inception module. It used batch normalization, image distortions and RMSprop. This module is based on several very small convolutions in order to drastically reduce the number of parameters. Their architecture consisted of a 22 layer deep CNN but reduced the number of parameters from 60 million (AlexNet) to 4 million.

<https://medium.com/@sidereal/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5>

Inception

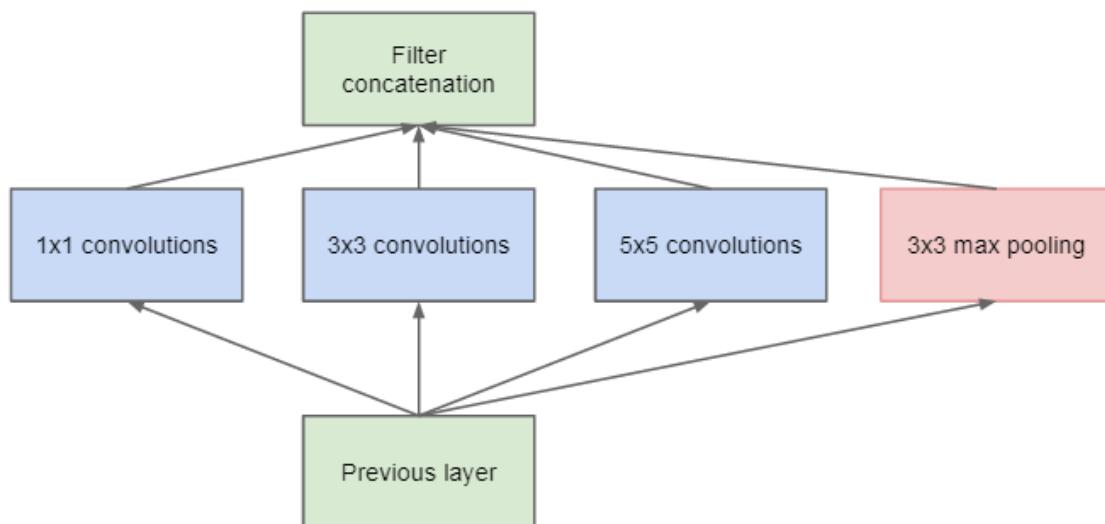


(a) Inception module, naïve version

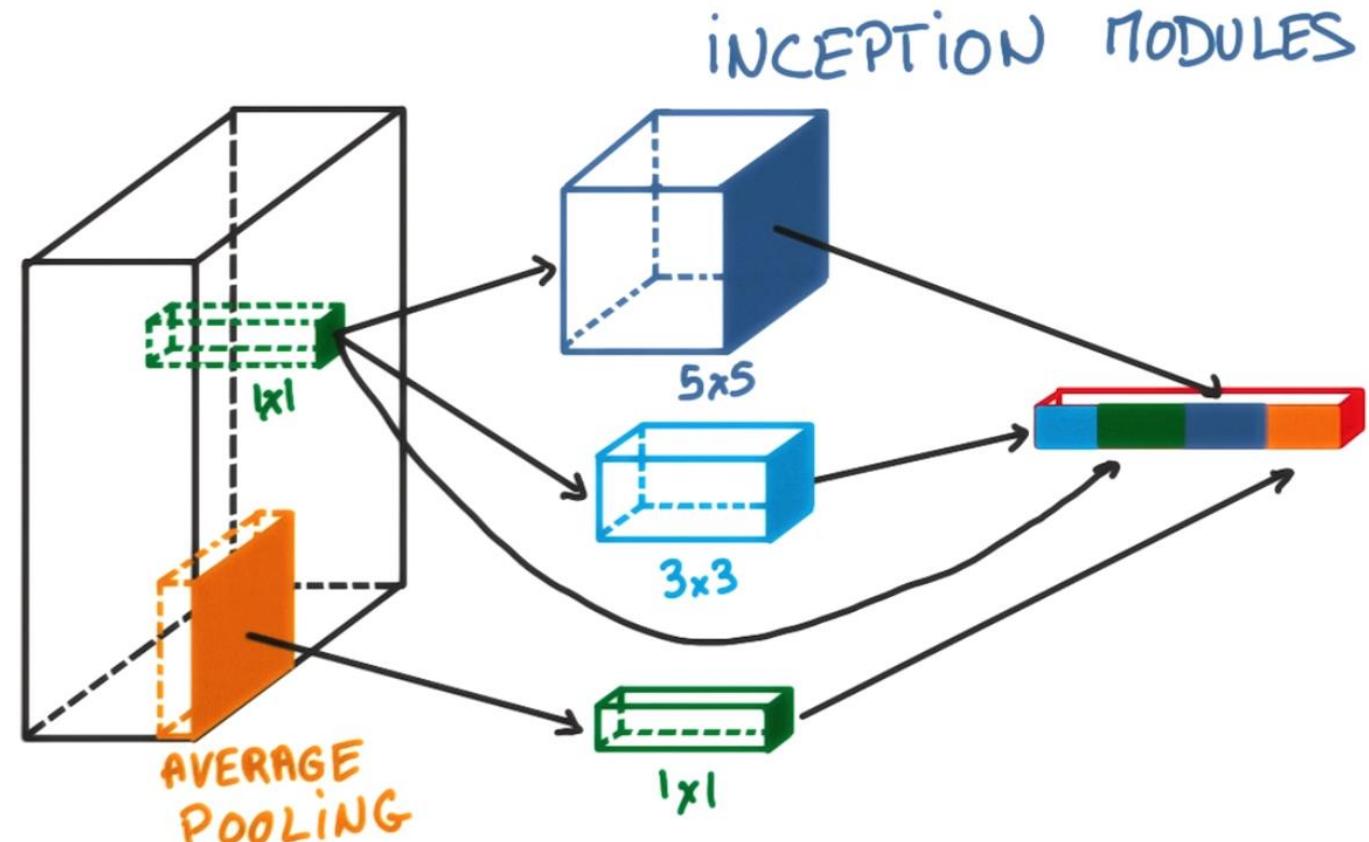


(b) Inception module with dimension reductions

Inception

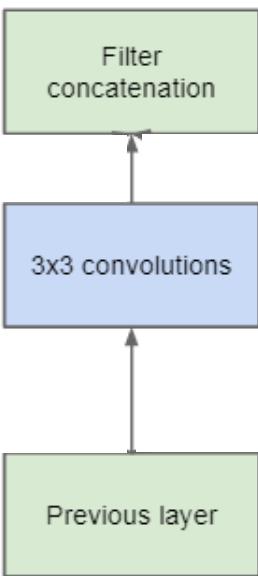


(a) Inception module, naïve version

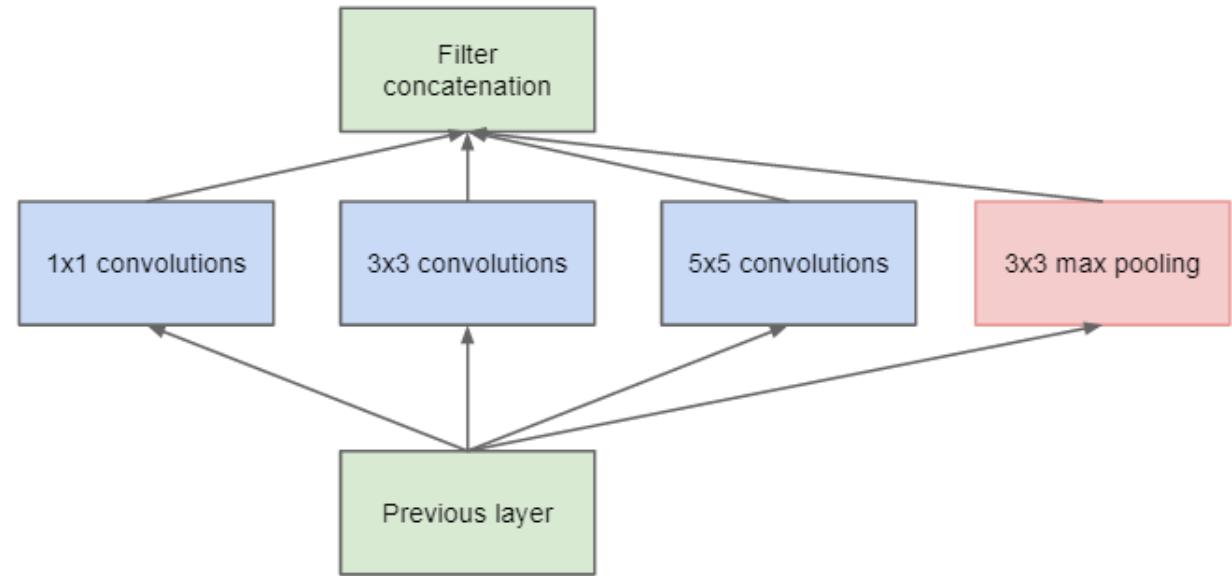


INCEPTION MODULES

Difference from VGG

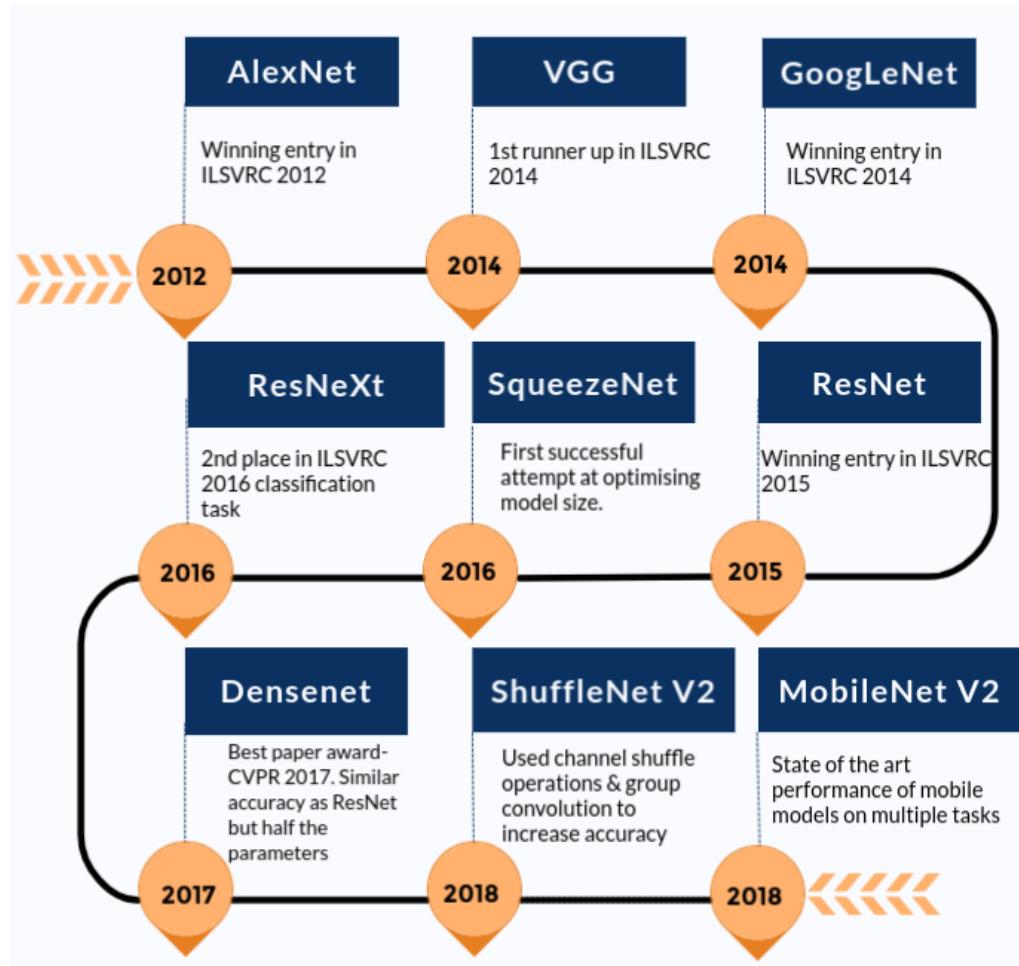


VGG



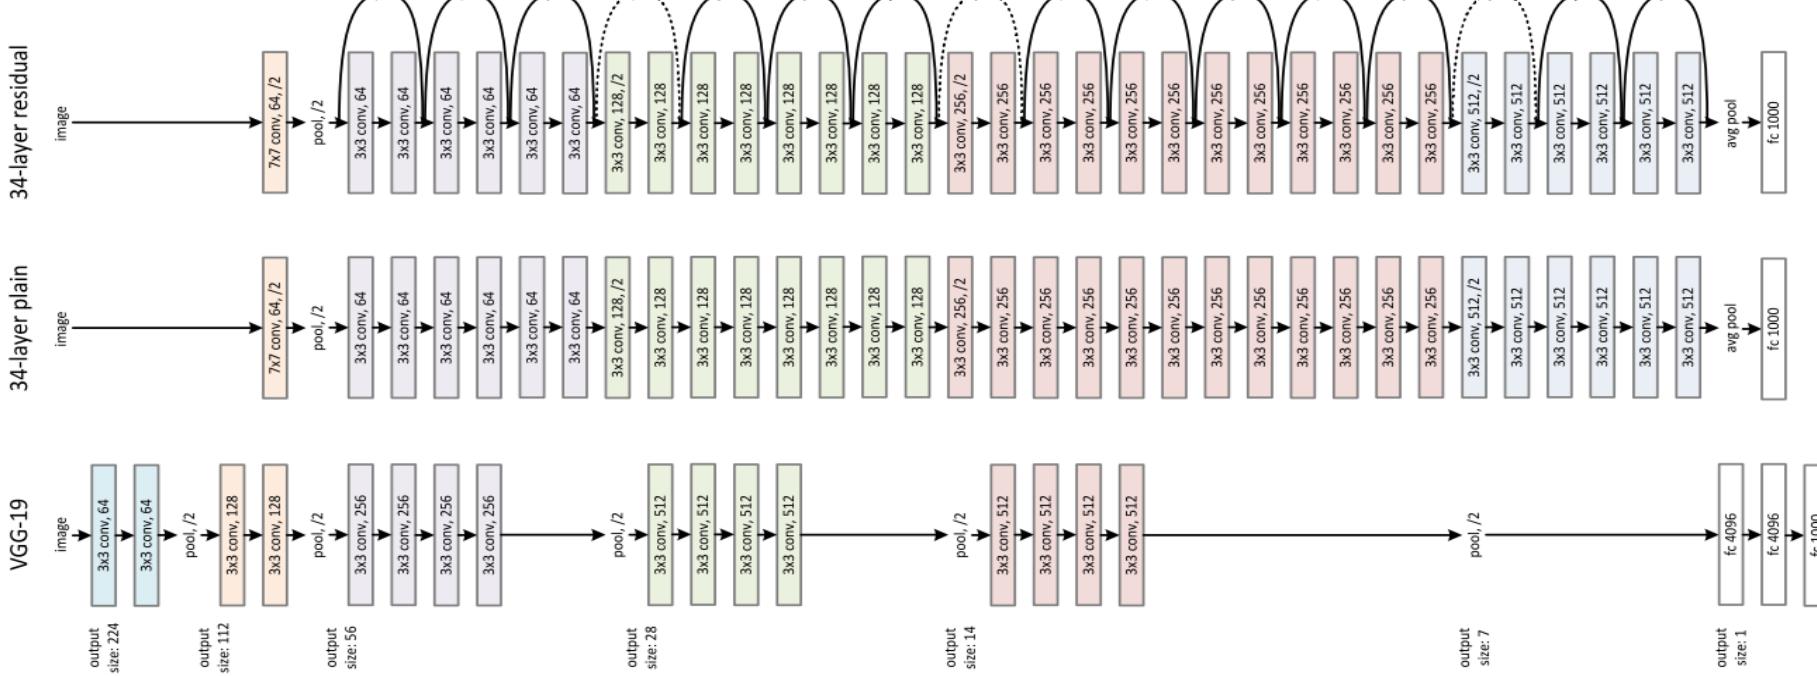
Google

State-of-the-art



AlexNet
VGG
GoogleNet
ResNet
DenseNet

ResNet(2015) vs VGG



At the ILSVRC 2015, the so-called Residual Neural Network (ResNet) by Kaiming He et al introduced a novel architecture with “skip connections” and features heavy batch normalization. Such skip connections are also known as gated units or gated recurrent units and have a strong similarity to recent successful elements applied in RNNs. Thanks to this technique they were able to train a NN with 152 layers while still having lower complexity than VGGNet. It achieves a top-5 error rate of 3.57% which beats human-level performance on this dataset.

https://medium.com/@pierre_guilhou/understand-how-works-resnet-without-talking-about-residual-64698f157eoc

<https://medium.com/@sidereal/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5>

ResNet

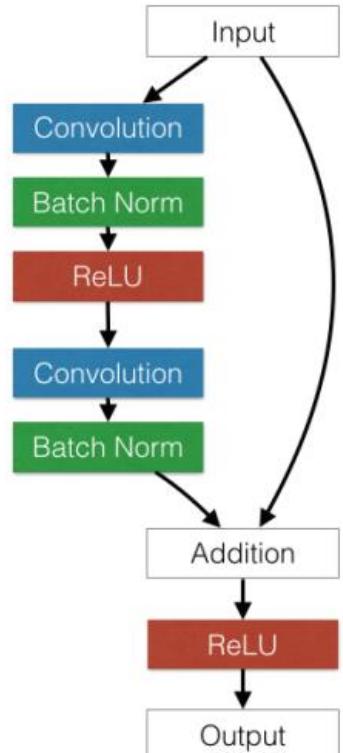
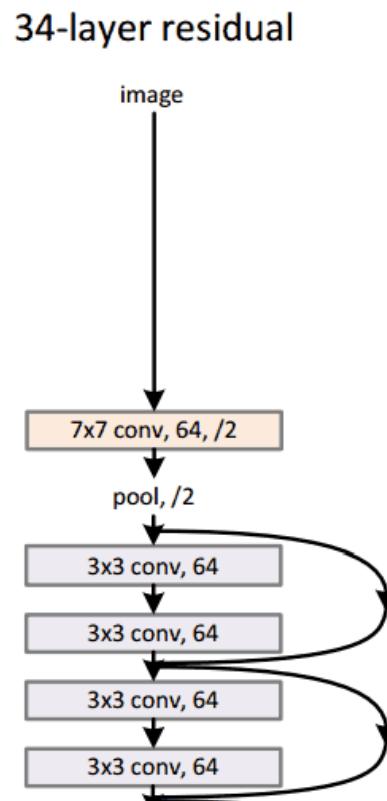


Figure 1. A RestNet basic block



34-layer residual

Netscope CNN Analyzer

A web-based tool for visualizing and analyzing convolutional neural network architectures (or technically, any directed acyclic graph). Currently supports Caffe's prototxt format.

Basis by ethereon. Extended for CNN Analysis by dgschwend.

Gist Support

If your `.prototxt` file is part of a GitHub Gist, you can visualize it by visiting this URL:

`http://dgschwend.github.io/netscope/#/gist/your-gist-id`

After ResNet

IMAGENET Large Scale Visual Recognition Challenge 2016 (ILSVRC2016)

[DET](#) [LOC](#) [VID](#) [Scene](#) [Team information](#)

Legend:

Yellow background = winner in this task according to this metric; authors are willing to reveal the method

White background = authors are willing to reveal the method

Grey background = authors chose not to reveal the method

Italics = authors requested entry not participate in competition

Object detection (DET)^[top]

Task 1a: Object detection with provided training data

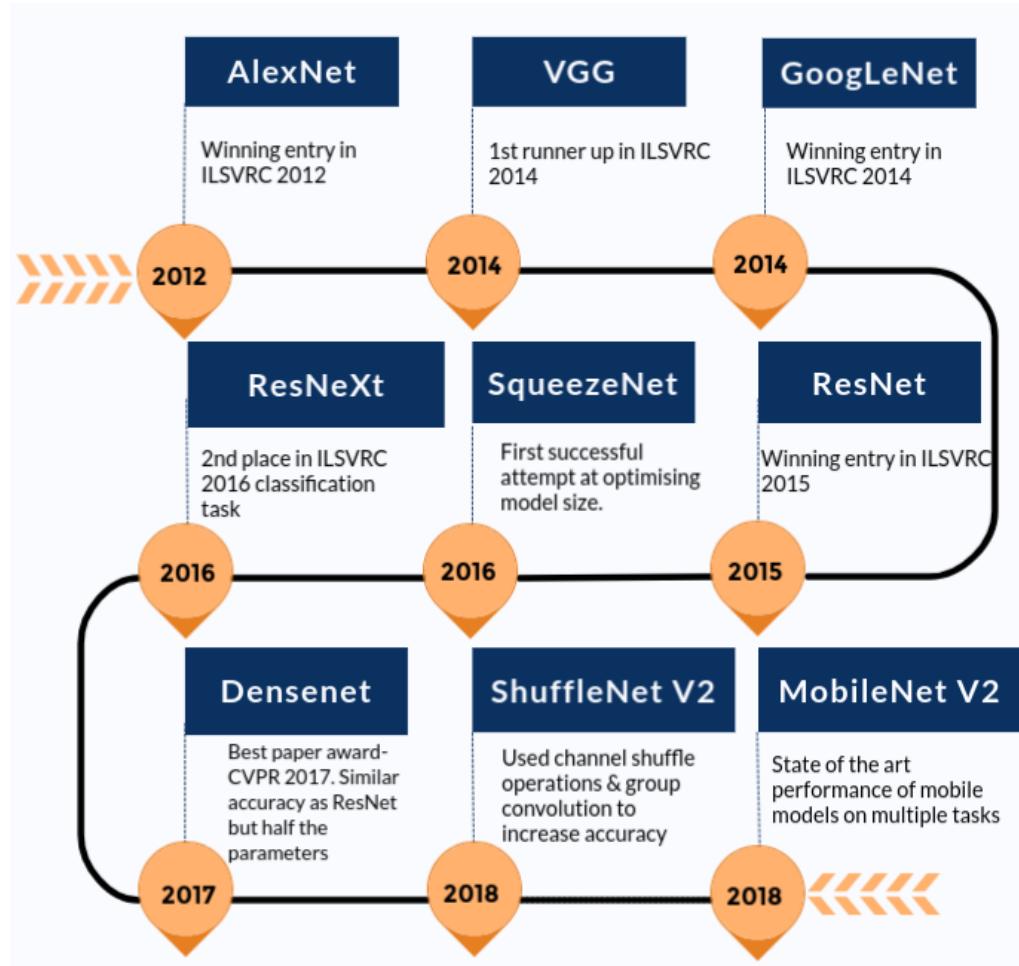
Ordered by number of categories won

Team name	Entry description	Number of object categories won	mean AP
CULimage	Ensemble of 6 models using provided data	109	0.662751
Hikvision	Ensemble A of 3 RPN and 6 FRCN models, mAP is 67 on val2	30	0.652704
Hikvision	Ensemble B of 3 RPN and 5 FRCN models, mean AP is 66.9, median AP is 69.3 on val2	18	0.652003

ImageNet 2017

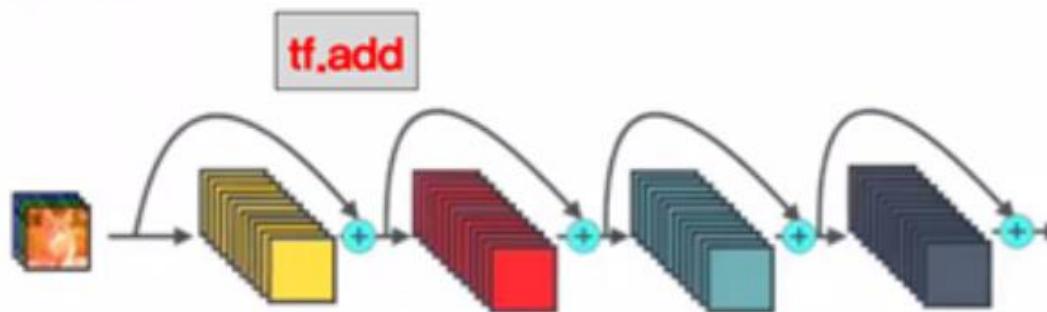


State-of-the-art



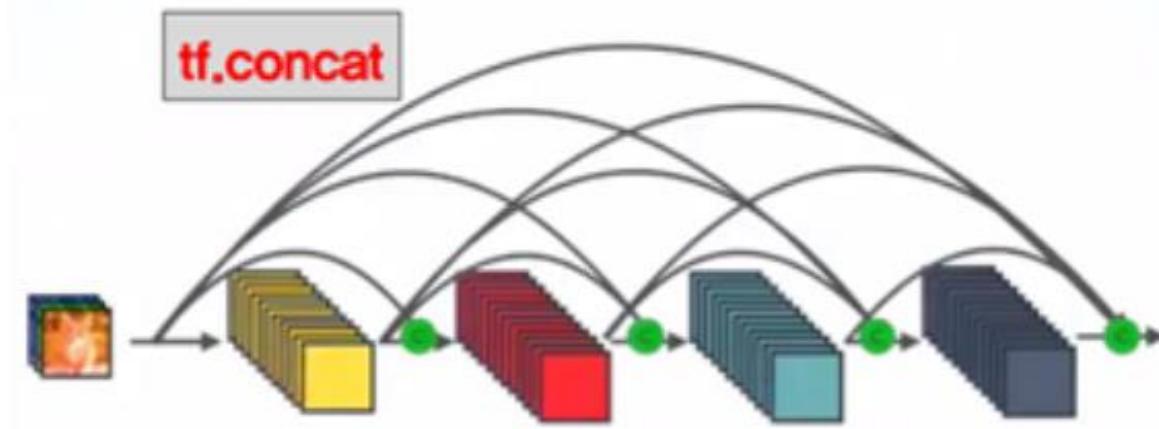
AlexNet
VGG
GoogleNet
ResNet
DenseNet

More Connections



+: Element-wise addition

ResNet Connectivity



● : Channel-wise concatenation

Dense Connectivity

Densenet

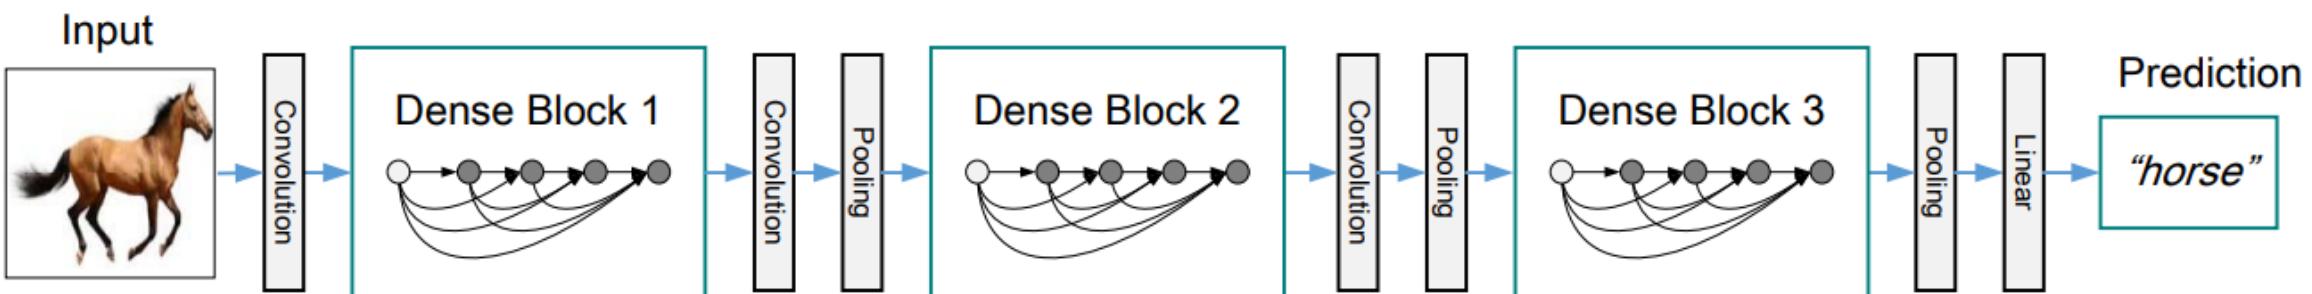
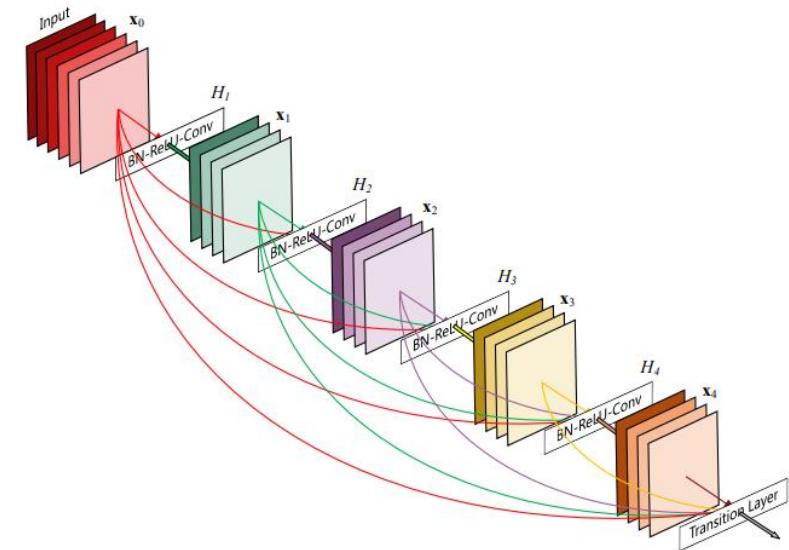
Densely Connected Convolutional Networks

Gao Huang*
Cornell University
gh349@cornell.edu

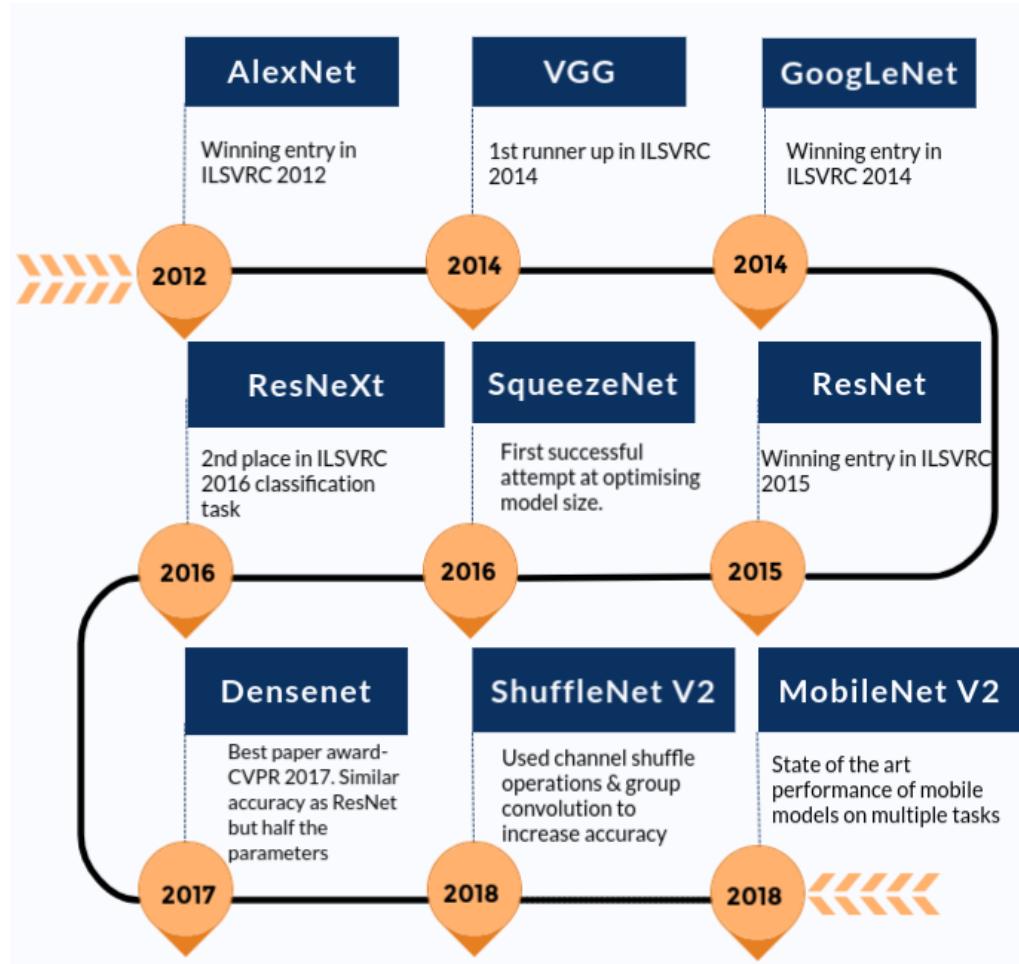
Zhuang Liu*
Tsinghua University
liuzhuang13@mails.tsinghua.edu.cn

Kilian Q. Weinberger
Cornell University
kqw4@cornell.edu

Laurens van der Maaten
Facebook AI Research
lvdmaaten@fb.com



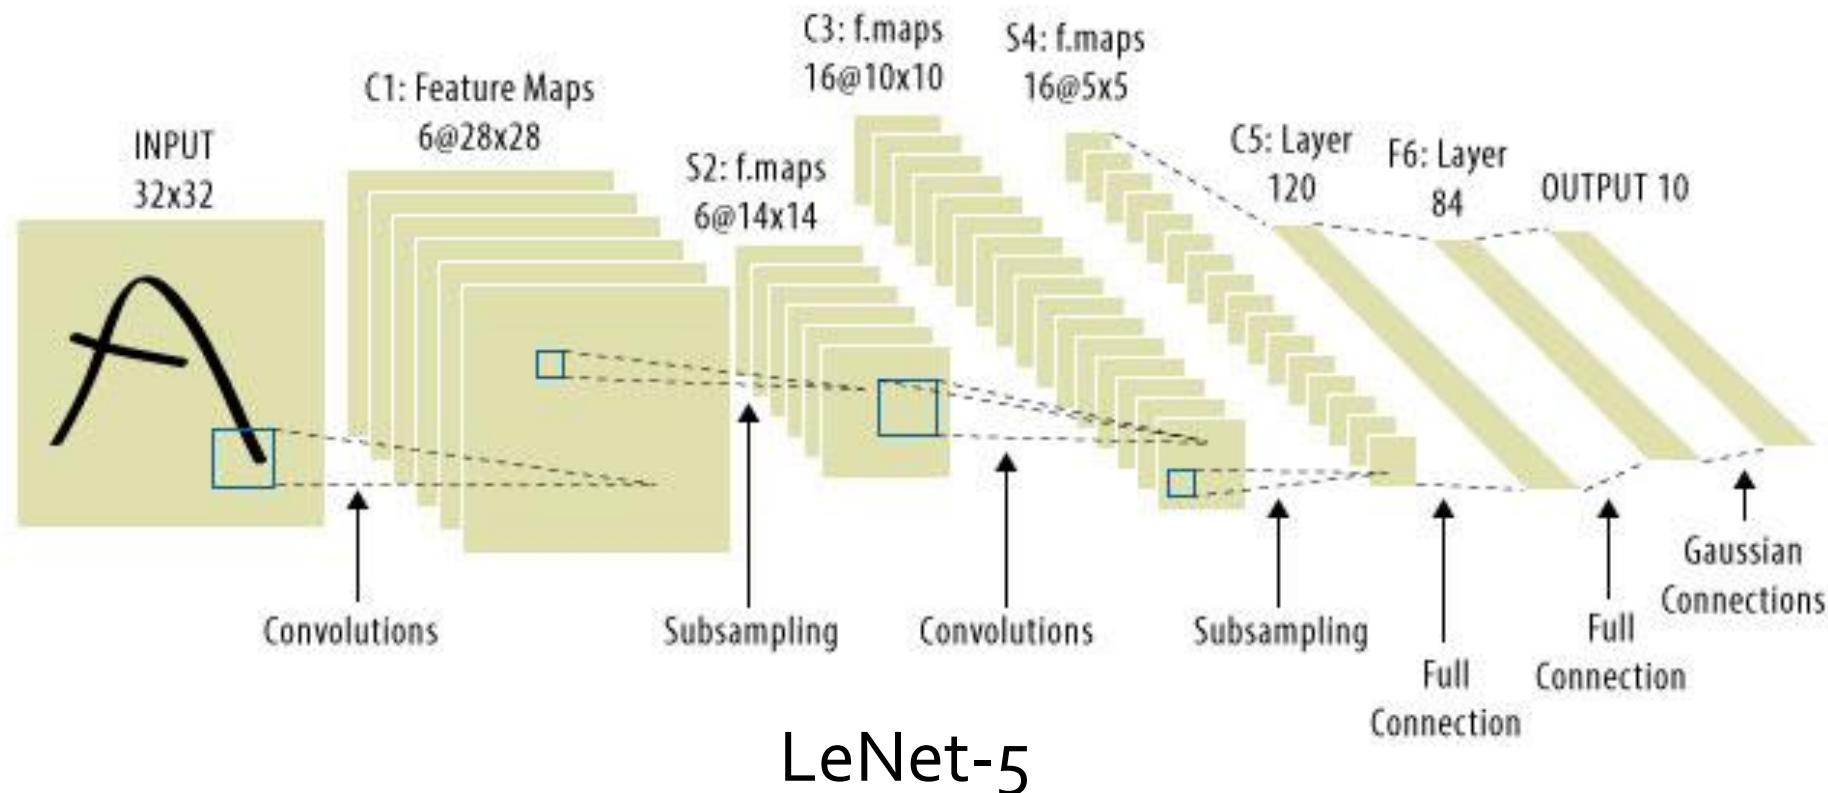
State-of-the-art



AlexNet
VGG
GoogleNet
ResNet
DenseNet

Summary of Classification

“CONV-POOL-FC”



Assignment 0

```
from __future__ import print_function
import argparse
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
```

Assignment 0

```
def main():
    # Training settings
    parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
    parser.add_argument('--batch-size', type=int, default=64, metavar='N',
                        help='input batch size for training (default: 64)')
    parser.add_argument('--test-batch-size', type=int, default=1000, metavar='N',
                        help='input batch size for testing (default: 1000)')
    parser.add_argument('--epochs', type=int, default=10, metavar='N',
                        help='number of epochs to train (default: 10)')
    parser.add_argument('--lr', type=float, default=0.01, metavar='LR',
                        help='learning rate (default: 0.01)')
    parser.add_argument('--momentum', type=float, default=0.5, metavar='M',
                        help='SGD momentum (default: 0.5)')
    parser.add_argument('--no-cuda', action='store_true', default=False,
                        help='disables CUDA training')
    parser.add_argument('--seed', type=int, default=1, metavar='S',
                        help='random seed (default: 1)')
    parser.add_argument('--log-interval', type=int, default=10, metavar='N',
                        help='how many batches to wait before logging training status')
    args = parser.parse_args()
    use_cuda = not args.no_cuda and torch.cuda.is_available()
```

Assignment o

```
torch.manual_seed(args.seed)

device = torch.device("cuda" if use_cuda else "cpu")

kwargs = {'num_workers': 1, 'pin_memory': True} if use_cuda else {}
train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=True, download=True,
                   transform=transforms.Compose([
                       transforms.ToTensor(),
                       transforms.Normalize((0.1307,), (0.3081,)))
                ]),
    batch_size=args.batch_size, shuffle=True, **kwargs)
test_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=False, transform=transforms.Compose([
                       transforms.ToTensor(),
                       transforms.Normalize((0.1307,), (0.3081,)))
                ]),
    batch_size=args.test_batch_size, shuffle=True, **kwargs)
```

Assignment 0

```
model = Net().to(device)
optimizer = optim.SGD(model.parameters(), lr=args.lr, momentum=args.momentum)

for epoch in range(1, args.epochs + 1):
    train(args, model, device, train_loader, optimizer, epoch)
    test(args, model, device, test_loader)
```

Assignment 0

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.conv2_drop = nn.Dropout2d()
        self.fc1 = nn.Linear(320, 50)
        self.fc2 = nn.Linear(50, 10)

    def forward(self, x):
        x = F.relu(F.max_pool2d(self.conv1(x), 2))
        x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x))), 2))
        x = x.view(-1, 320)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)
```

Assignment 0

```
def train(args, model, device, train_loader, optimizer, epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

Assignment 0

```
def test(args, model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item() # sum
            pred = output.max(1, keepdim=True)[1] # get the index of the max log-p
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))
```