

Generative Adversarial Network

[Spring 2020 CS-8395 Deep Learning in Medical Image Computing]

Instructor: Yuankai Huo, Ph.D.
Department of Electrical Engineering and Computer Science
Vanderbilt University

Topics



- Review
- Semantic Segmentation
- U-Net
- Instance Segmentation

Tasks

Computer Vision

Classification



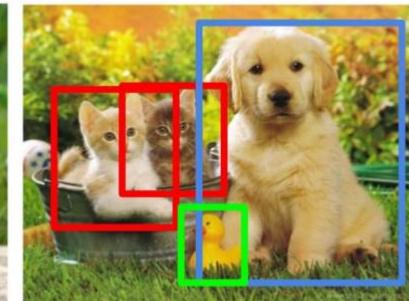
CAT

Classification + Localization



CAT

Object Detection



CAT, DOG, DUCK

Instance Segmentation



CAT, DOG, DUCK

Synthesis



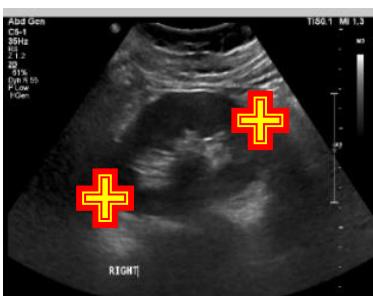
Medical Image Computing



Input



Classification



Detection



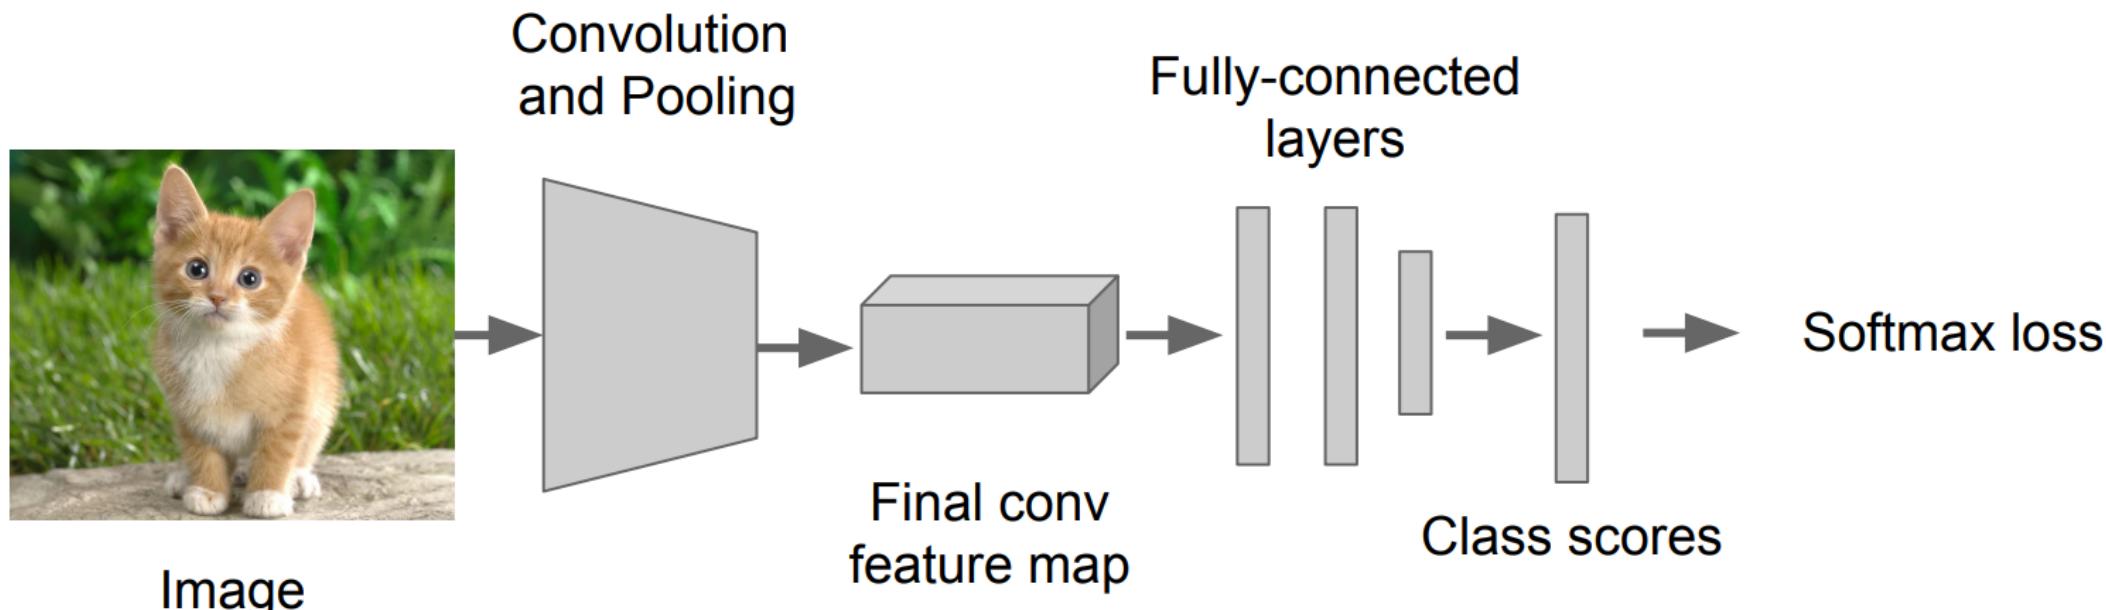
Segmentation



Synthesis

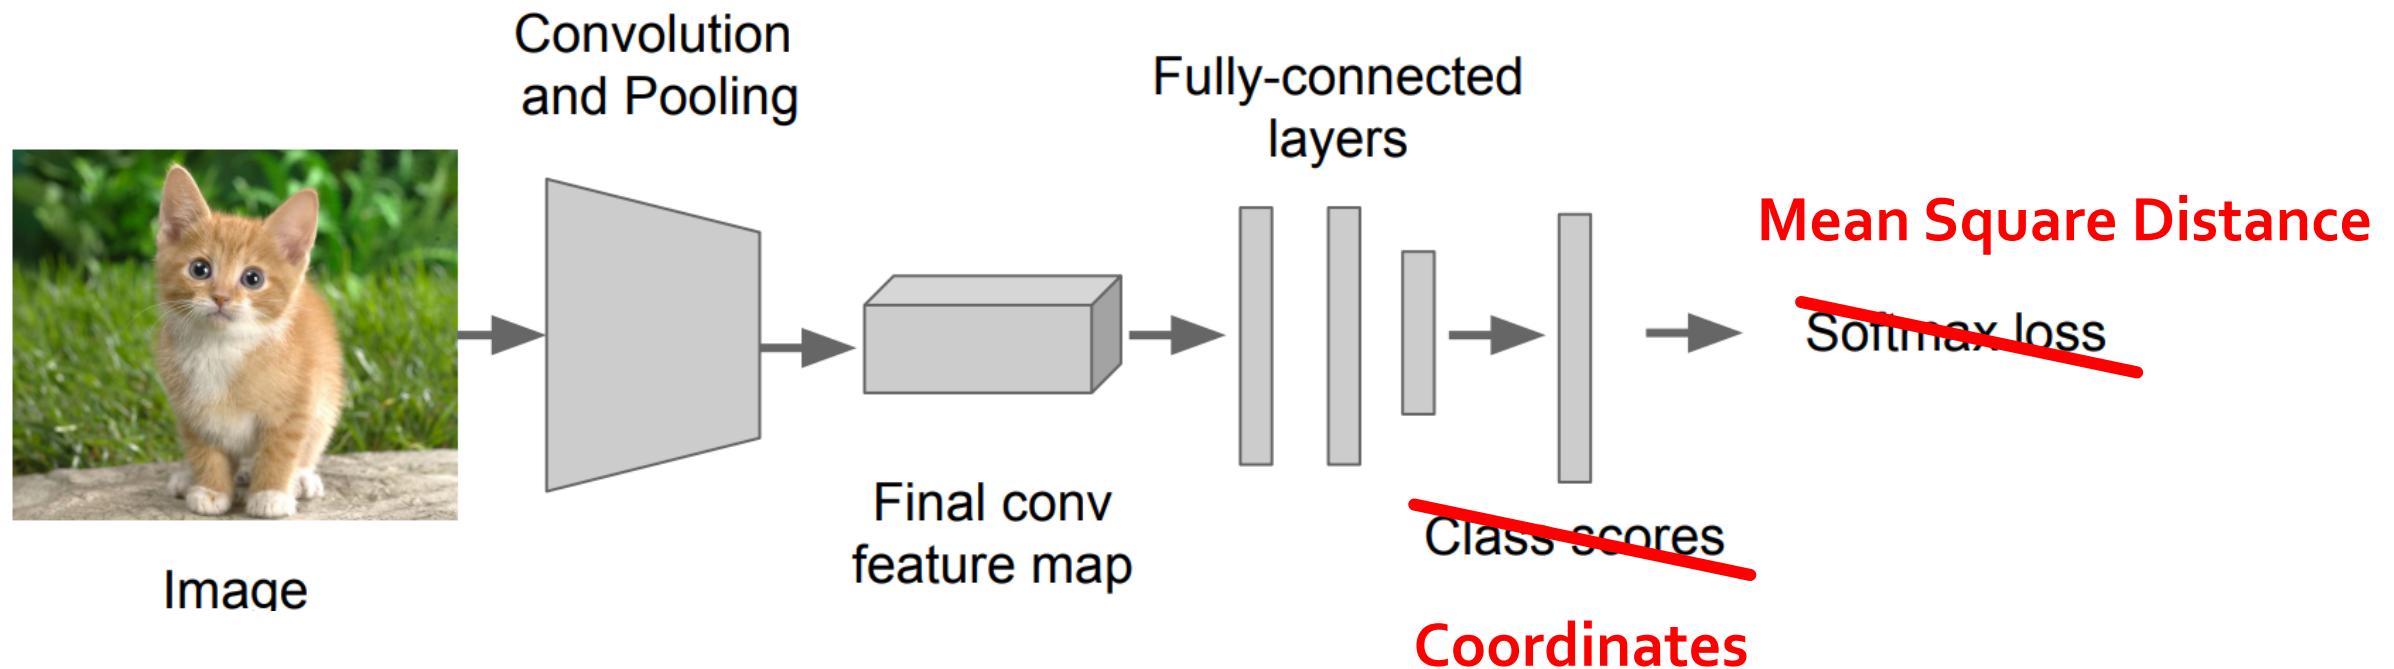
Simple Recipe for Classification

“CONV-POOL-FC”



Simple Recipe for Localization

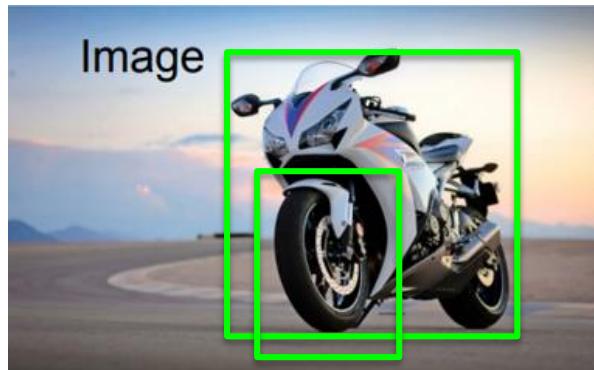
“CONV-POOL-FC”



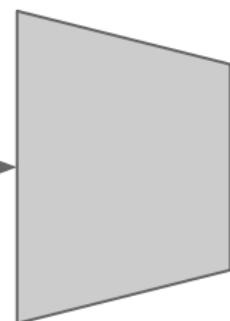
Detection

“CONV-POOL-FC”

Region Proposals



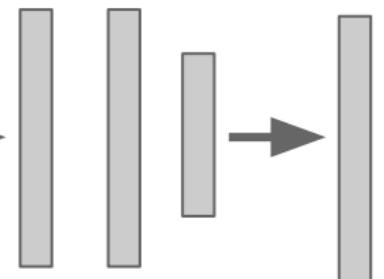
Convolution
and Pooling



Final conv
feature map



Fully-connected
layers



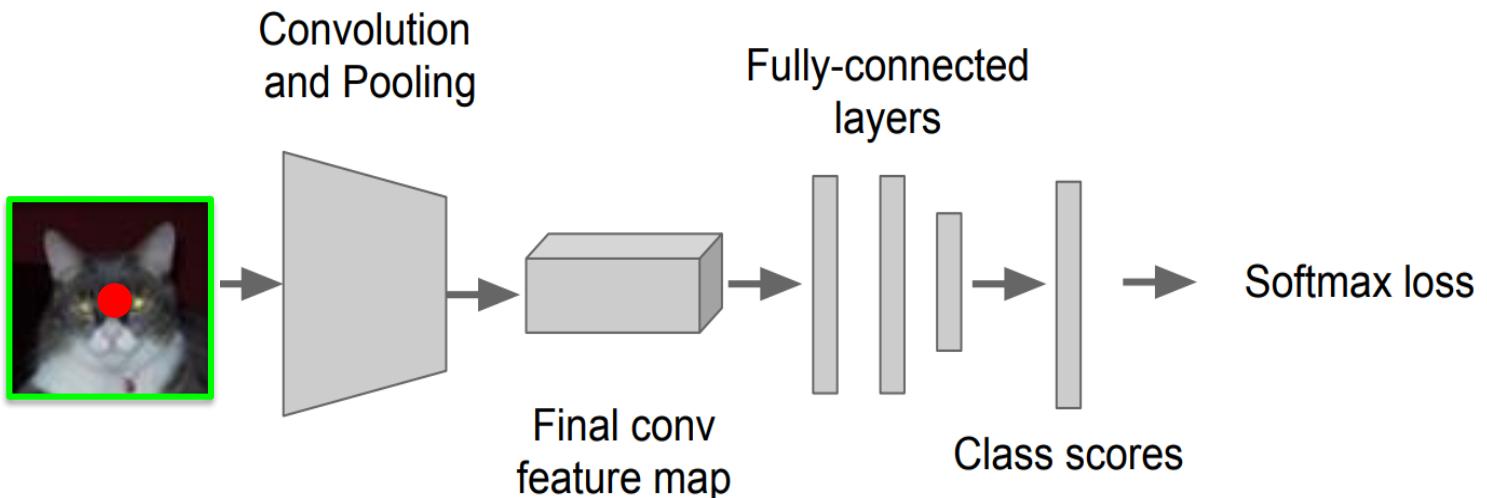
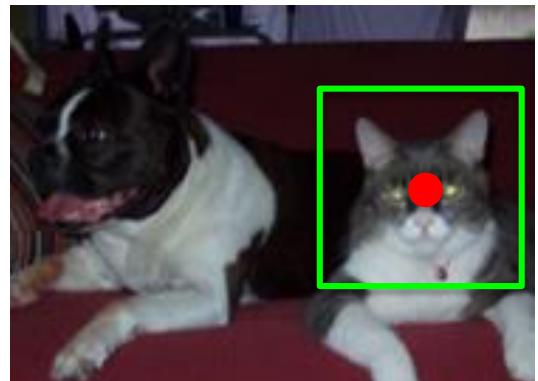
Class scores

Softmax loss

http://slazebni.cs.illinois.edu/spring17/lec07_detection.pdf
https://blog.csdn.net/sum_nap/article/details/80388110
http://cs231n.stanford.edu/slides/2016/winter1516_lecture8.pdf

Segmentation

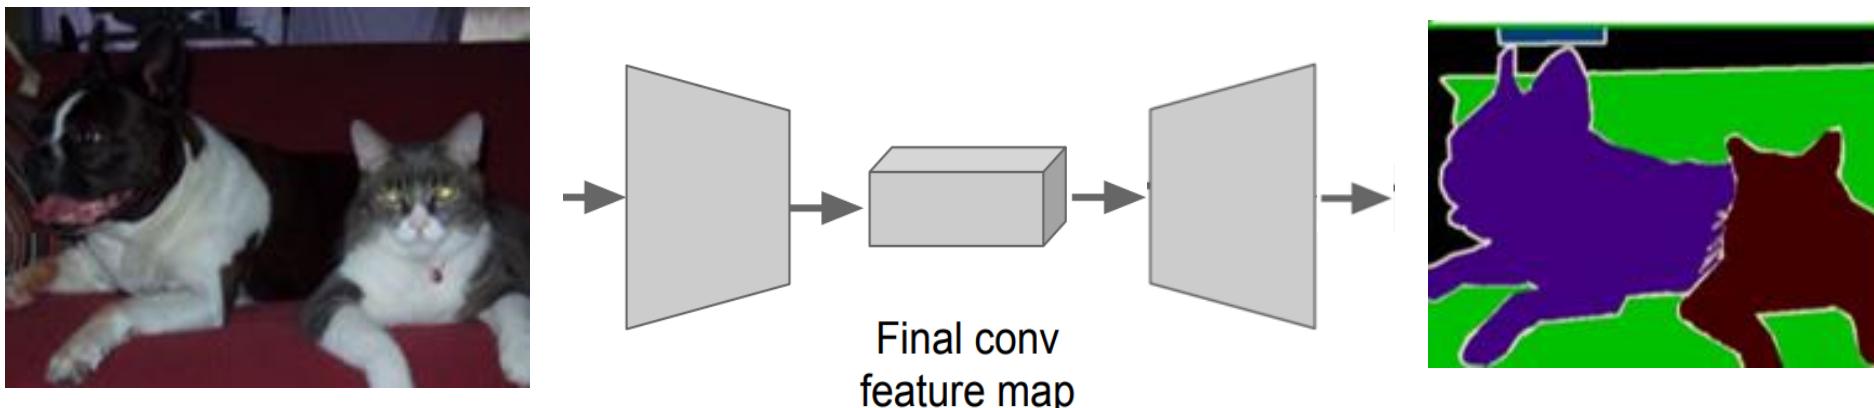
“CONV-POOL-FC”



http://slazebni.cs.illinois.edu/spring17/lec07_detection.pdf
https://blog.csdn.net/sum_nap/article/details/80388110
http://cs231n.stanford.edu/slides/2016/winter1516_lecture8.pdf

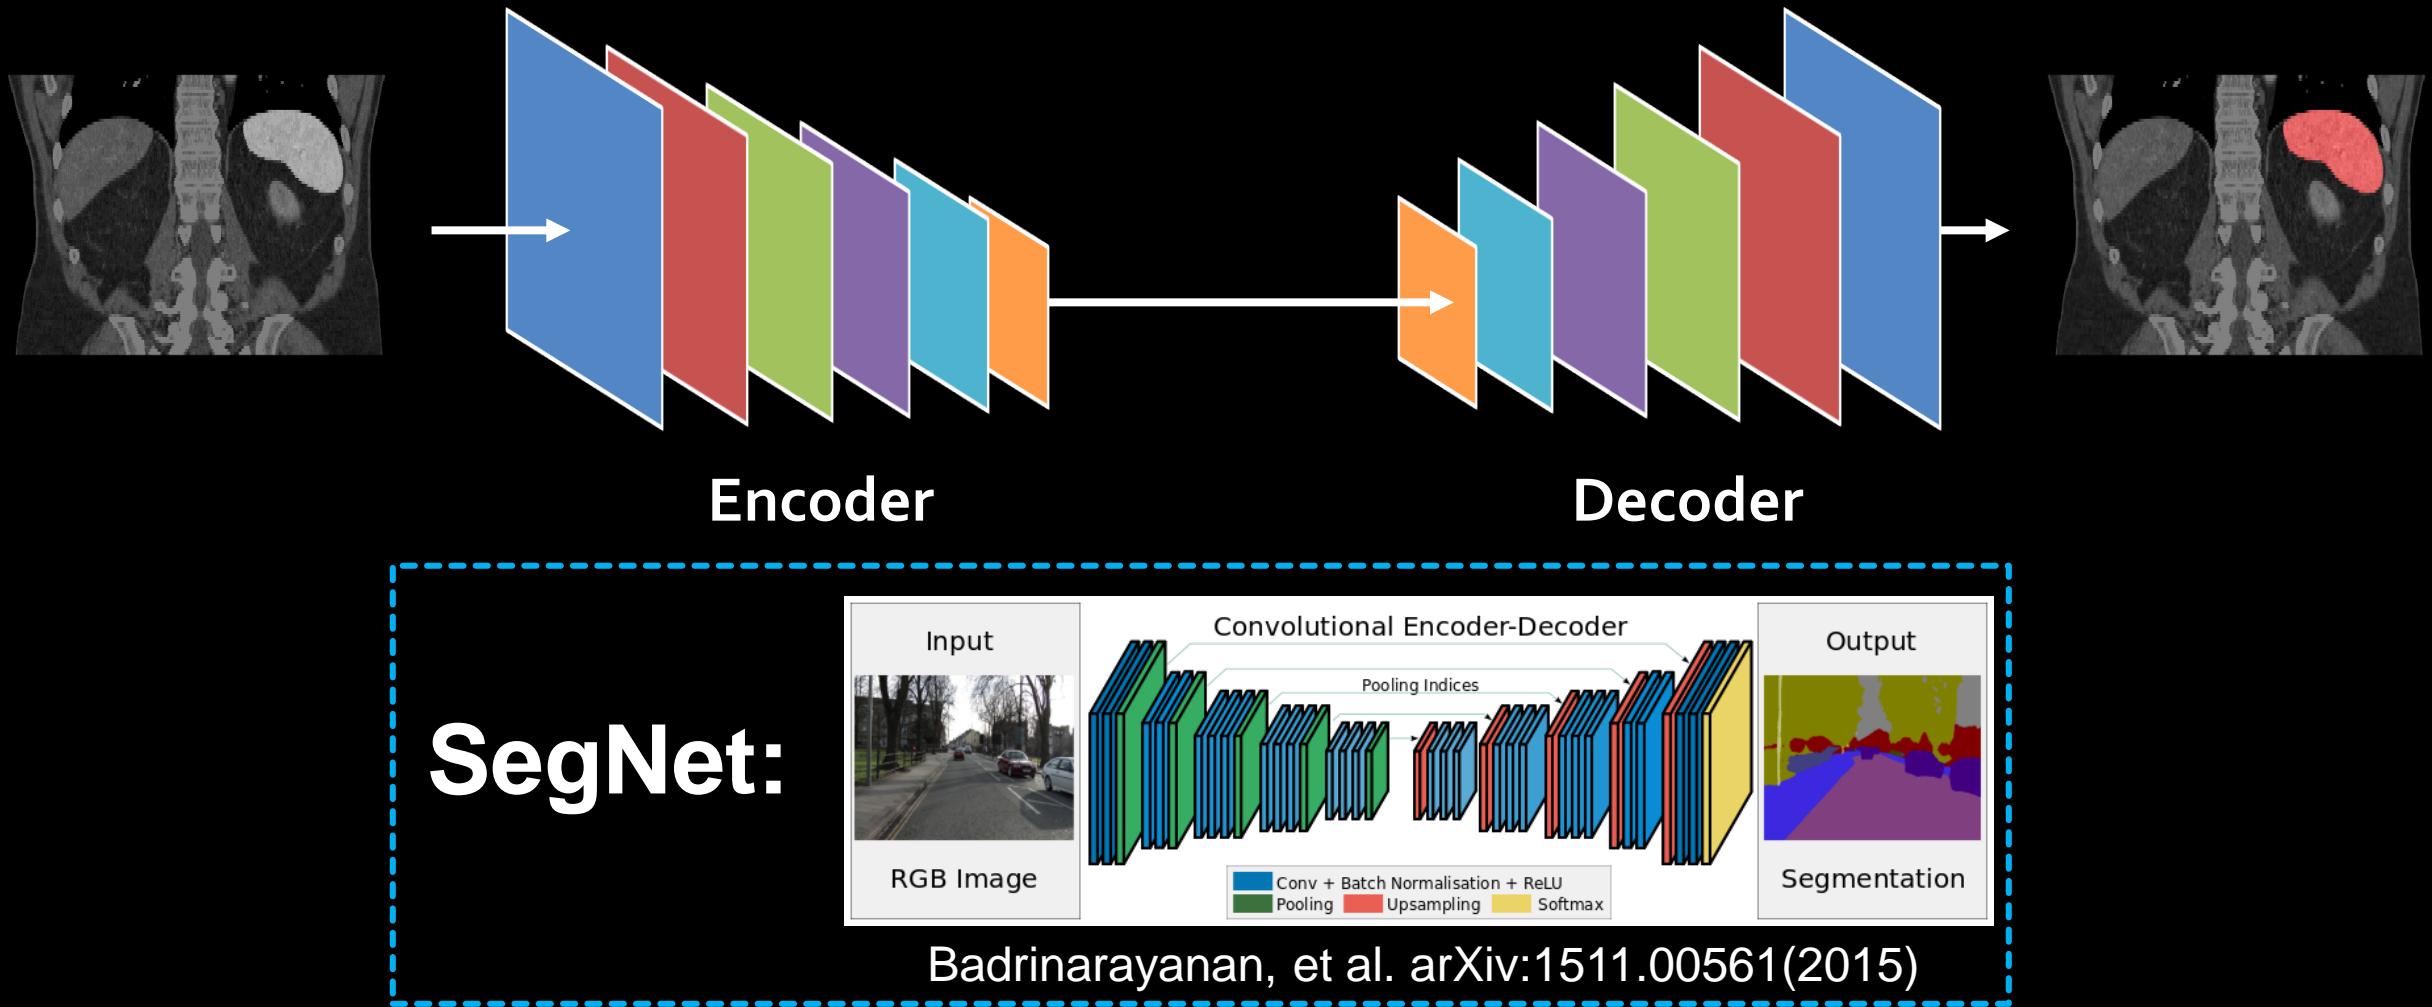
Segmentation

“Fully Conv Net”

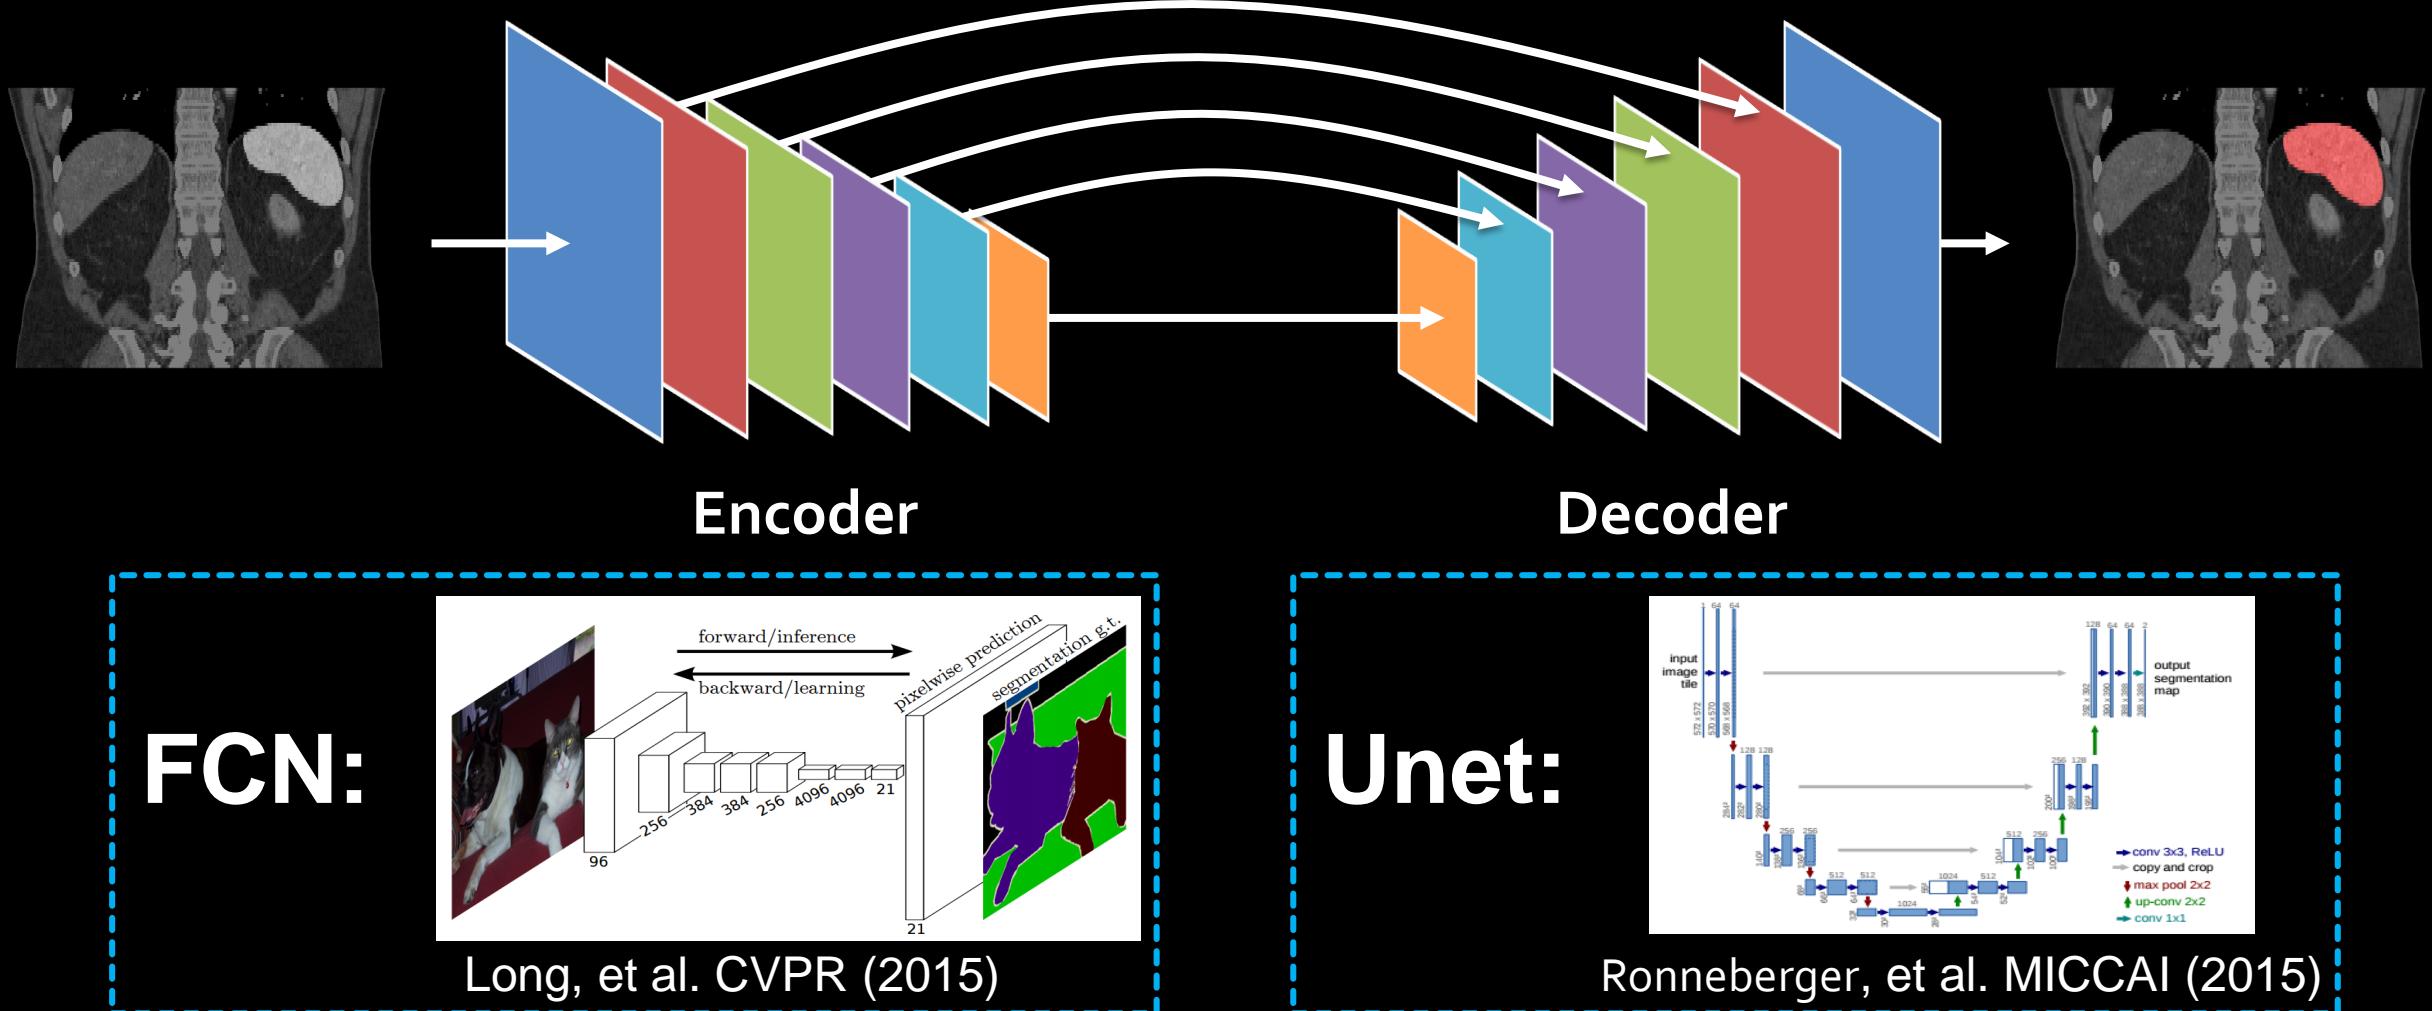


http://slazebni.cs.illinois.edu/spring17/lec07_detection.pdf
https://blog.csdn.net/sum_nap/article/details/80388110
http://cs231n.stanford.edu/slides/2016/winter1516_lecture8.pdf

Encoder – Decoder Strategy



Fully Convolutional Network (FCN)



All Code for U-Net

Encoder

```
layer1 = F.relu(self.conv1_input(x))
layer1 = F.relu(self.conv1(layer1))

layer2 = F.max_pool2d(layer1, 2)
layer2 = F.relu(self.conv2_input(layer2))
layer2 = F.relu(self.conv2(layer2))

layer3 = F.max_pool2d(layer2, 2)
layer3 = F.relu(self.conv3_input(layer3))
layer3 = F.relu(self.conv3(layer3))

layer4 = F.max_pool2d(layer3, 2)
layer4 = F.relu(self.conv4_input(layer4))
layer4 = F.relu(self.conv4(layer4))

layer5 = F.max_pool2d(layer4, 2)
layer5 = F.relu(self.conv5_input(layer5))
layer5 = F.relu(self.conv5(layer5))
```

Decoder

```
layer6 = F.relu(self.conv6_up(layer5))
layer6 = torch.cat((layer4, layer6), 1)
layer6 = F.relu(self.conv6_input(layer6))
layer6 = F.relu(self.conv6(layer6))

layer7 = F.relu(self.conv7_up(layer6))
layer7 = torch.cat((layer3, layer7), 1)
layer7 = F.relu(self.conv7_input(layer7))
layer7 = F.relu(self.conv7(layer7))

layer8 = F.relu(self.conv8_up(layer7))
layer8 = torch.cat((layer2, layer8), 1)
layer8 = F.relu(self.conv8_input(layer8))
layer8 = F.relu(self.conv8(layer8))

layer9 = F.relu(self.conv9_up(layer8))
layer9 = torch.cat((layer1, layer9), 1)
layer9 = F.relu(self.conv9_input(layer9))
layer9 = F.relu(self.conv9(layer9))
layer9 = self.final(self.conv9_output(layer9))

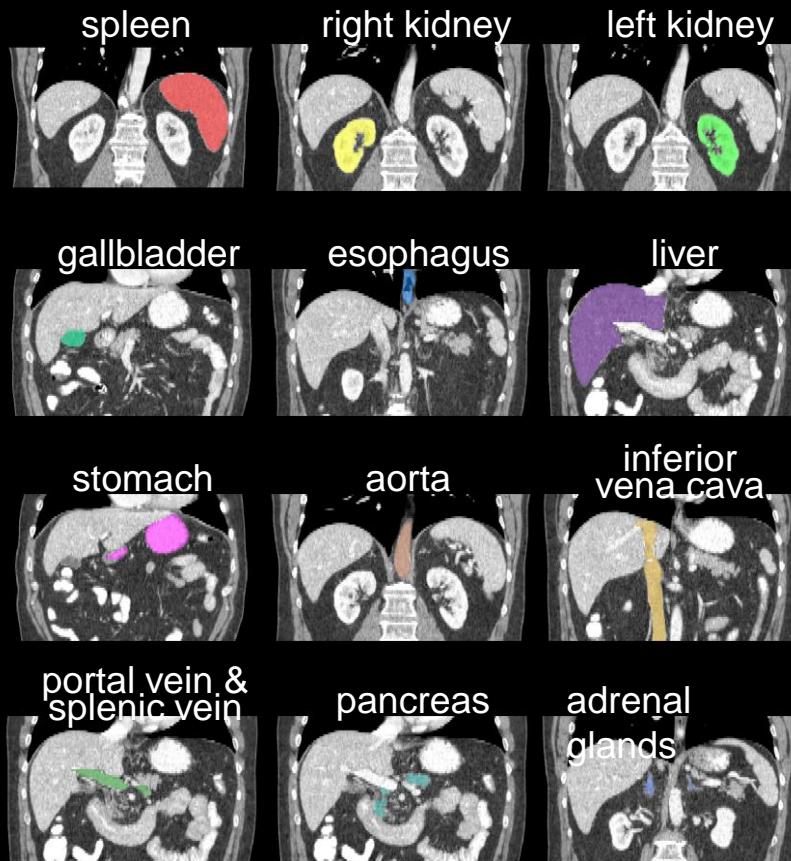
return layer9
```

Parameters

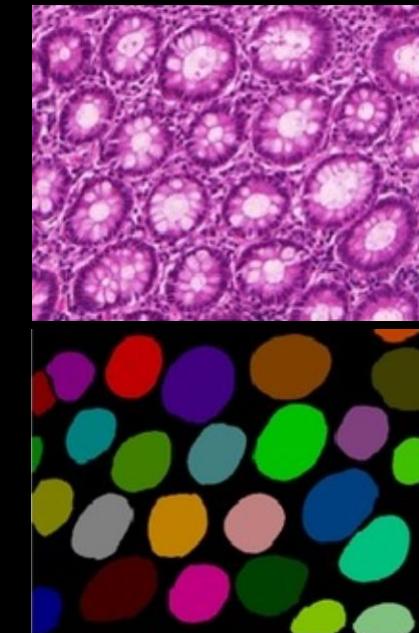
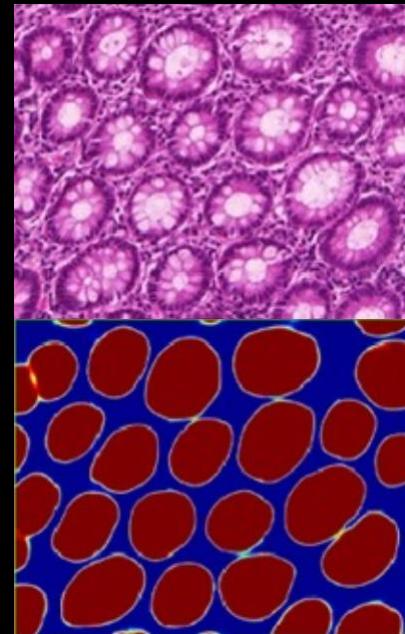
self.conv1_input =	nn.Conv2d(1, 64, 3, padding=1)
self.conv1 =	nn.Conv2d(64, 64, 3, padding=1)
self.conv2_input =	nn.Conv2d(64, 128, 3, padding=1)
self.conv2 =	nn.Conv2d(128, 128, 3, padding=1)
self.conv3_input =	nn.Conv2d(128, 256, 3, padding=1)
self.conv3 =	nn.Conv2d(256, 256, 3, padding=1)
self.conv4_input =	nn.Conv2d(256, 512, 3, padding=1)
self.conv4 =	nn.Conv2d(512, 512, 3, padding=1)
self.conv5_input =	nn.Conv2d(512, 1024, 3, padding=1)
self.conv5 =	nn.Conv2d(1024, 1024, 3, padding=1)
self.conv6_up =	nn.ConvTranspose2d(1024, 512, 2, 2)
self.conv6_input =	nn.Conv2d(1024, 512, 3, padding=1)
self.conv6 =	nn.Conv2d(512, 512, 3, padding=1)
self.conv7_up =	nn.ConvTranspose2d(512, 256, 2, 2)
self.conv7_input =	nn.Conv2d(512, 256, 3, padding=1)
self.conv7 =	nn.Conv2d(256, 256, 3, padding=1)
self.conv8_up =	nn.ConvTranspose2d(256, 128, 2, 2)
self.conv8_input =	nn.Conv2d(256, 128, 3, padding=1)
self.conv8 =	nn.Conv2d(128, 128, 3, padding=1)
self.conv9_up =	nn.ConvTranspose2d(128, 64, 2, 2)
self.conv9_input =	nn.Conv2d(128, 64, 3, padding=1)
self.conv9 =	nn.Conv2d(64, 64, 3, padding=1)
self.conv9_output =	nn.Conv2d(64, 2, 1)

Instance Segmentation

Semantic Segmentation

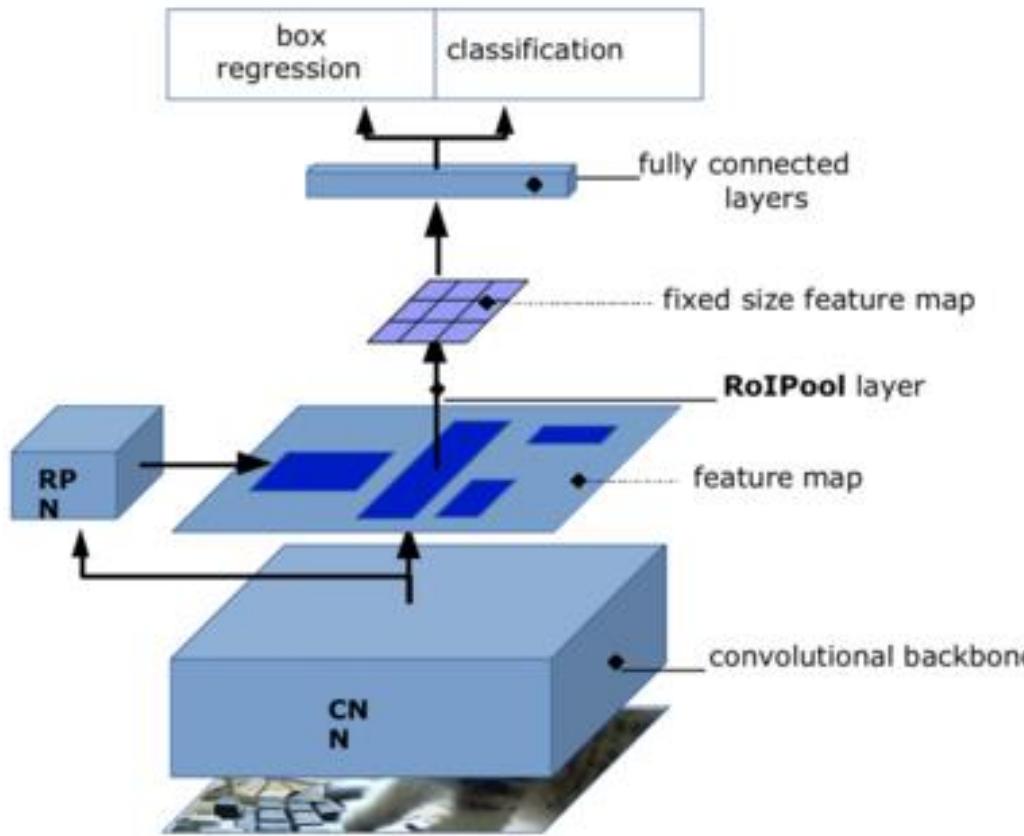


Instance Segmentation



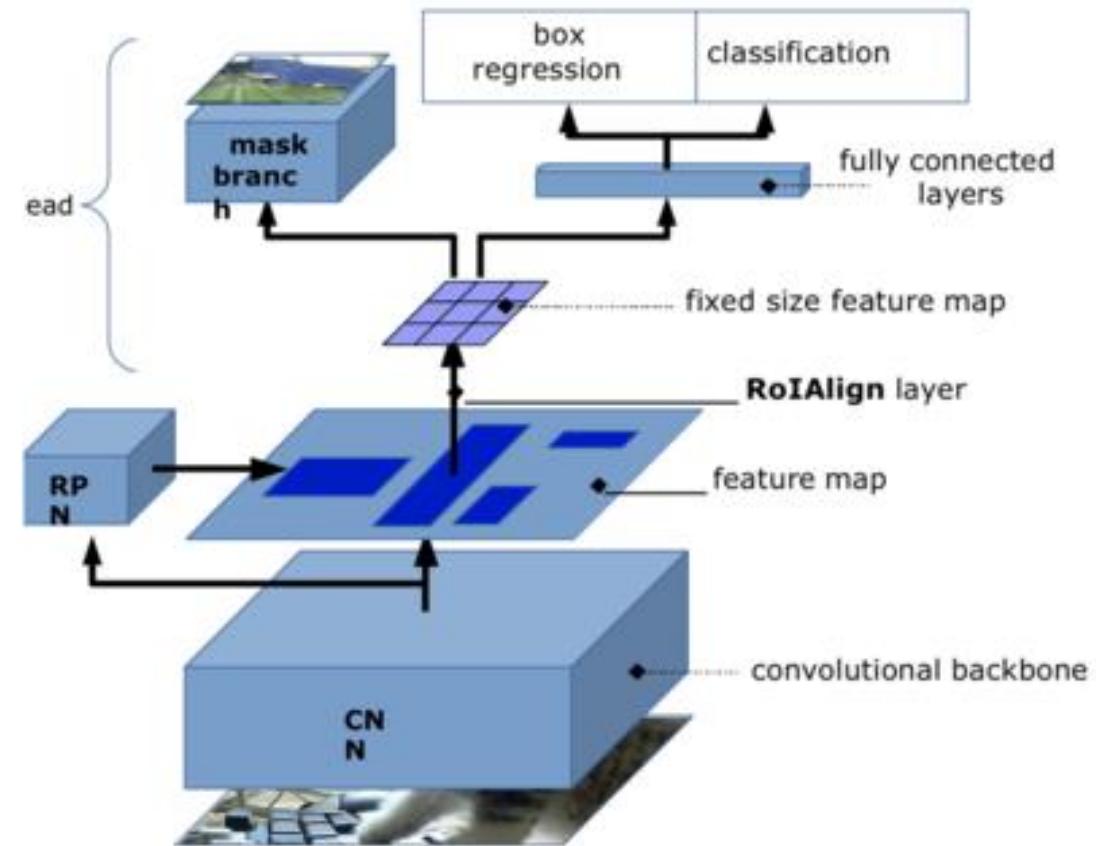
<https://www.sciencedirect.com/science/article/pii/S1361841516302043>

Faster R-CNN vs. Mask R-CNN



a)

https://wiki.ubc.ca/CNNs_in_Image_Segmentation



b)

Image Synthesis



<https://junyanz.github.io/CycleGAN/>

More Examples

Converting Monet into Thomas Kinkade



Portrait to Dollface



Resurrecting Ancient Cities



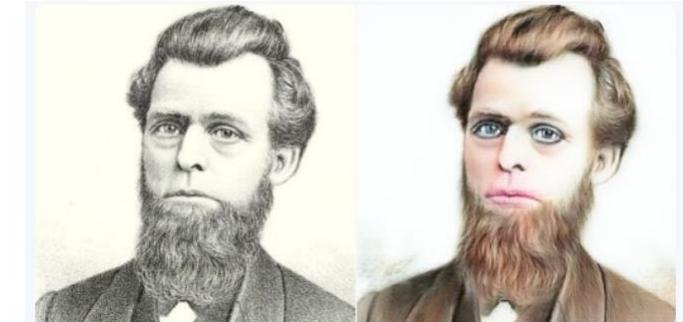
Animal Transfiguration



Face \leftrightarrow Ramen

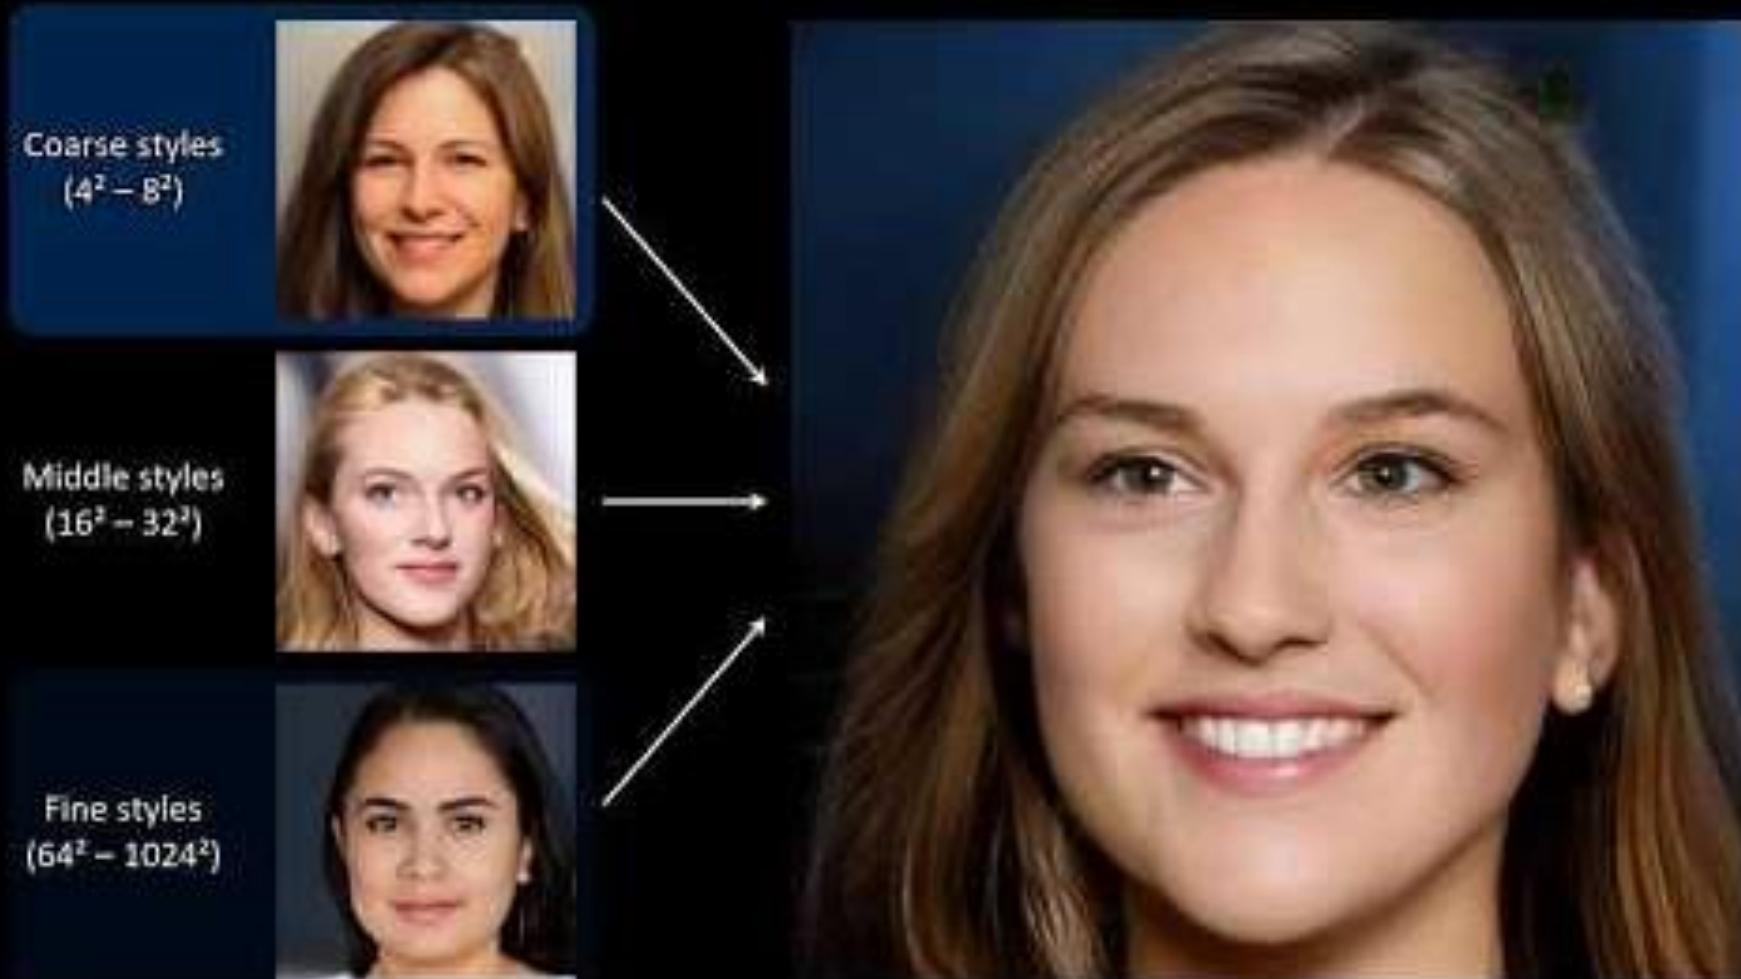


Colorizing legacy photographs



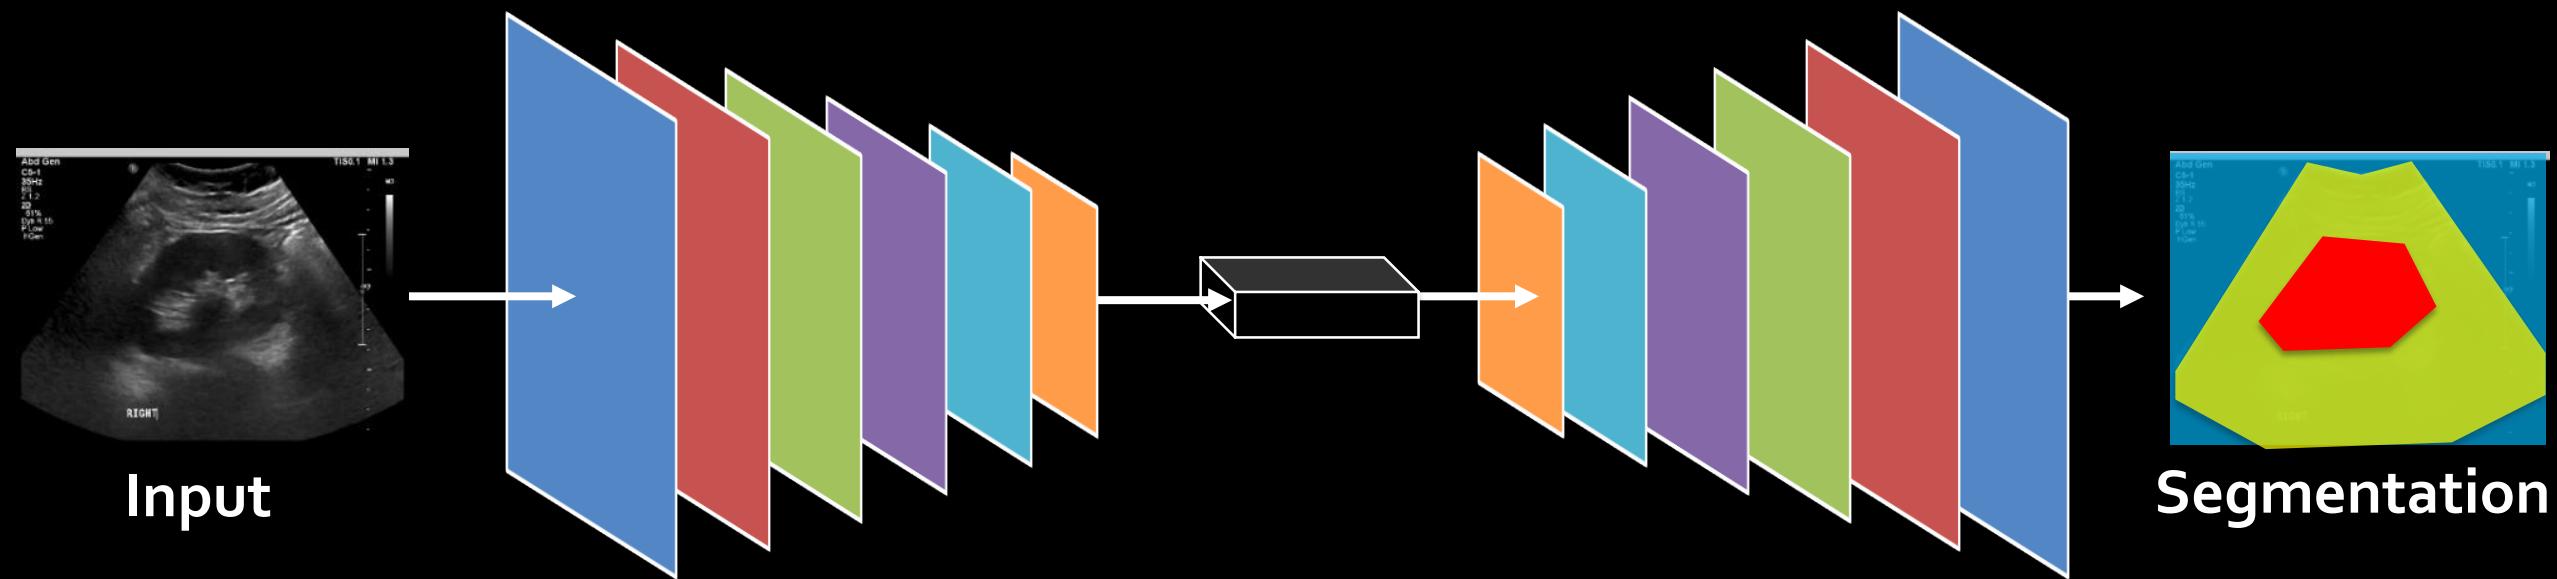
<https://junyanz.github.io/CycleGAN/>

State-of-the-art

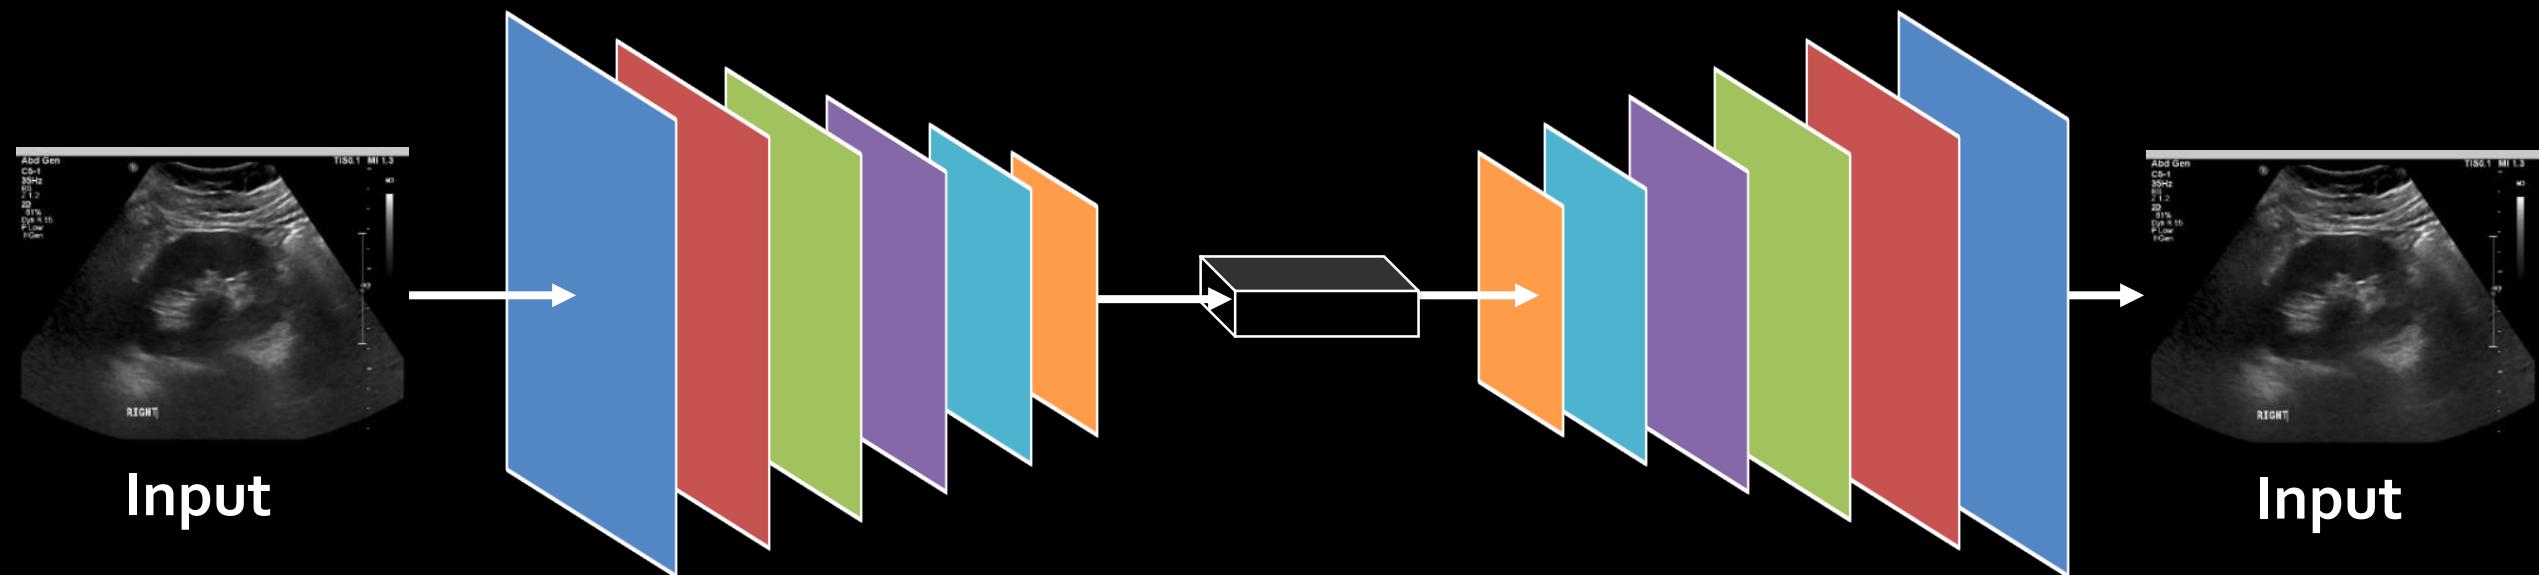


<https://www.youtube.com/watch?v=kSLJriaOumA>

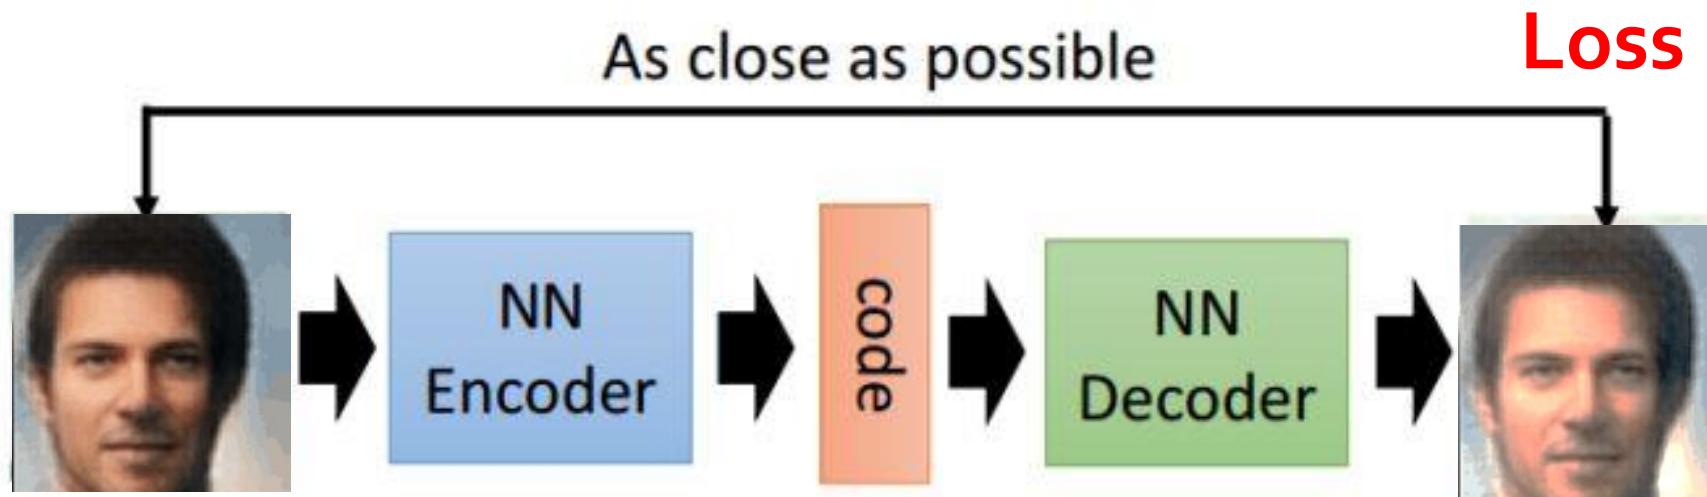
Segmentation (Image to Image)



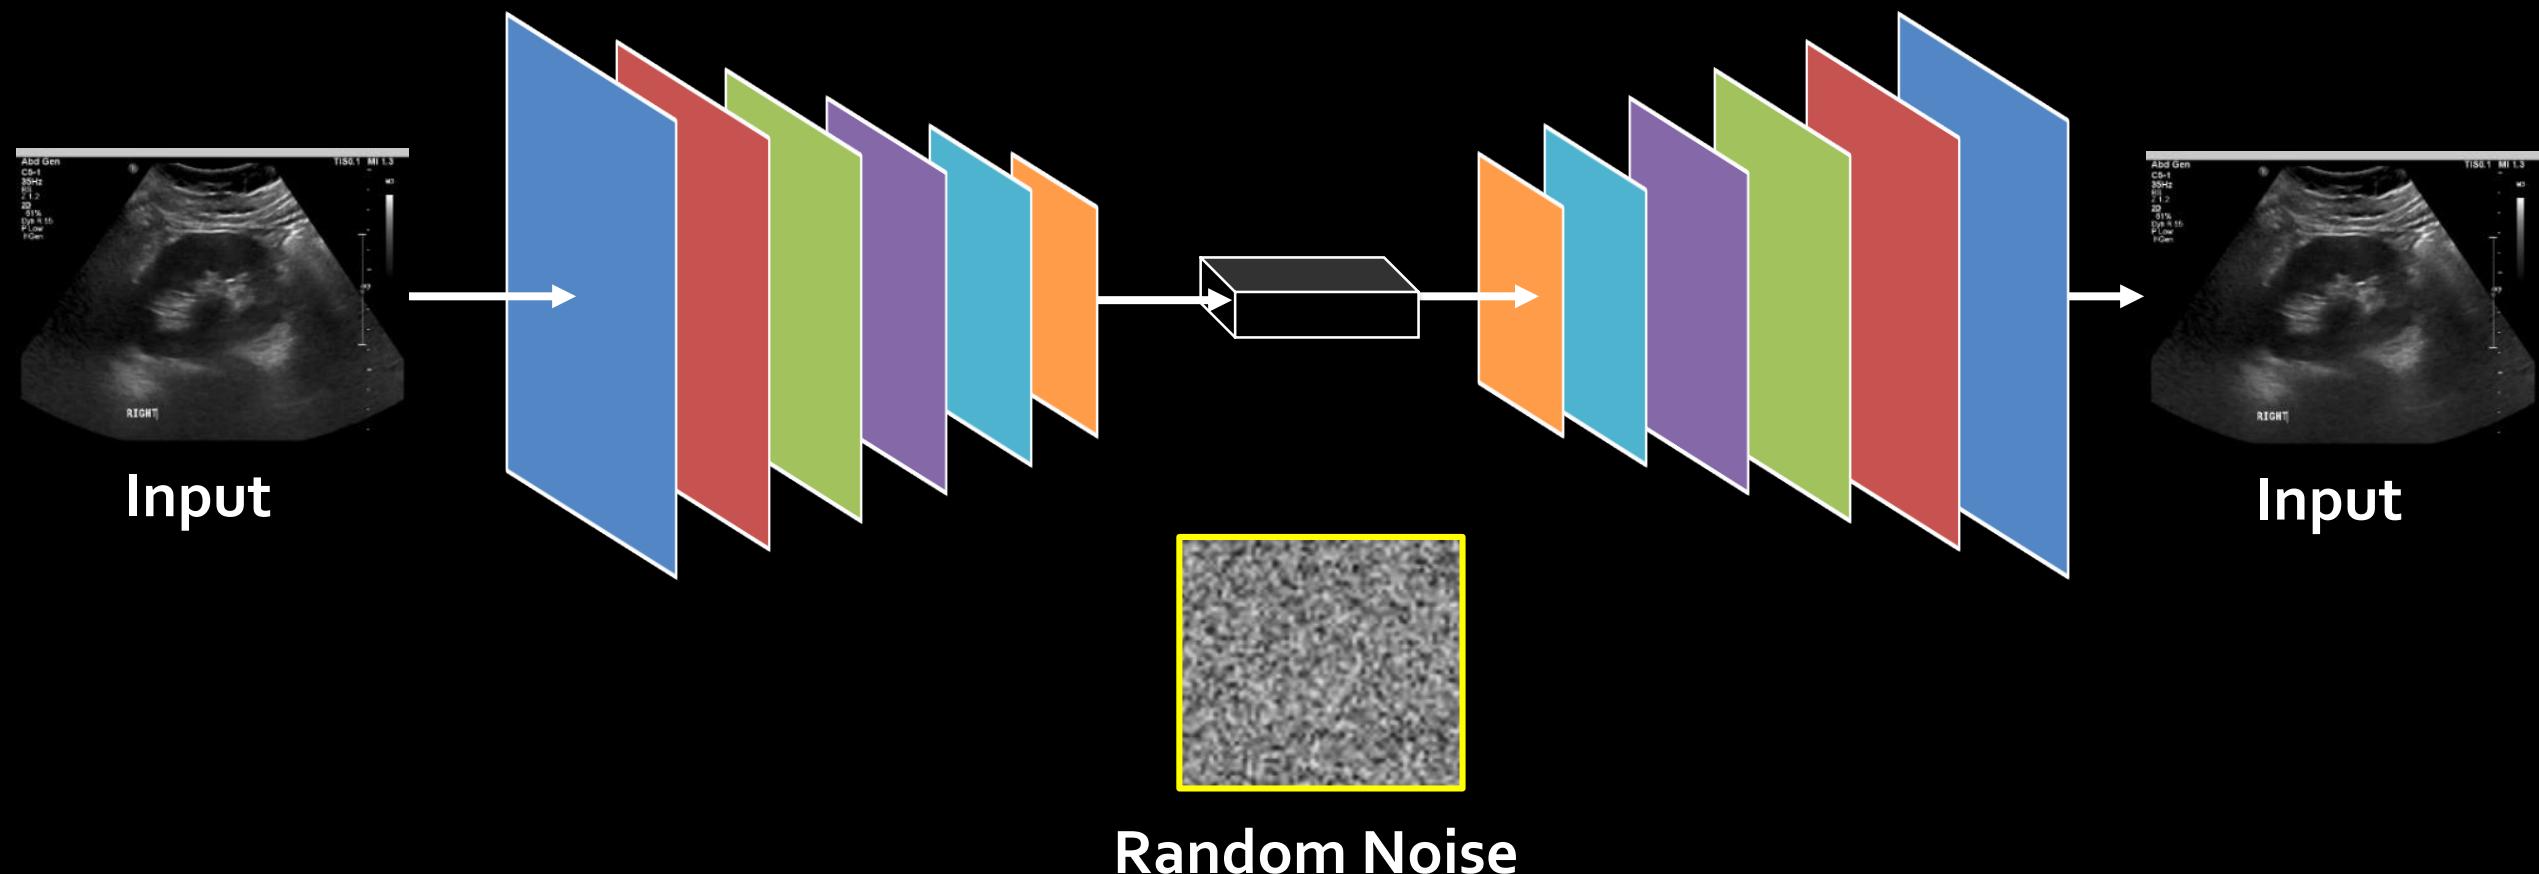
Auto Encoder



Let Deep Network Do It!

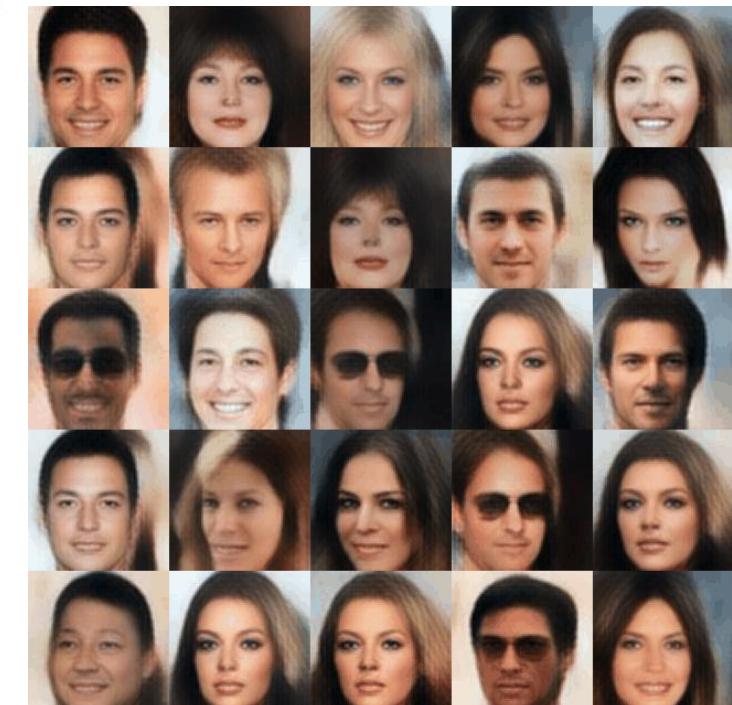
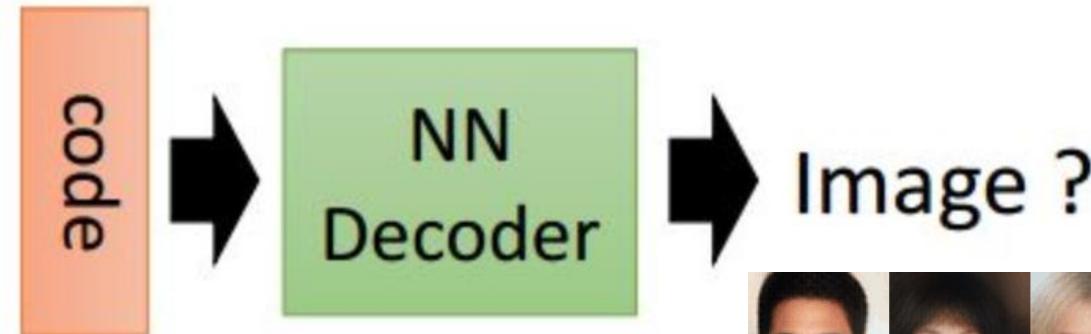


Auto encoder



Auto-encoder

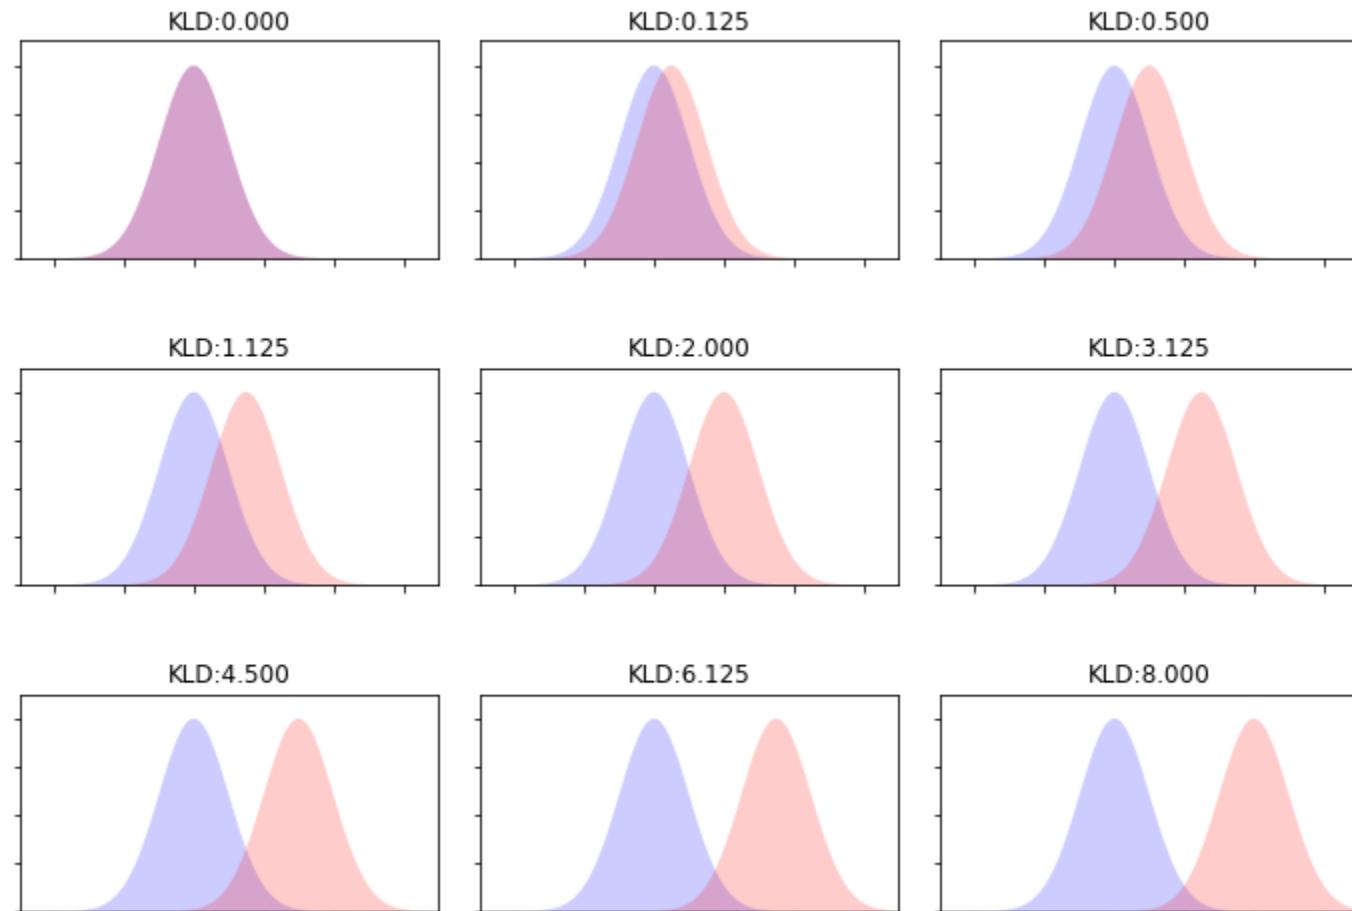
Randomly generate
a vector as code



<https://zhuanlan.zhihu.com/p/27549418>

<https://houxianxu.github.io/assets/project/dfcvae>

KL divergence



GAN

(Generative Adversarial Network)



"Generative Adversarial Networks is the **most interesting idea in the last ten years** in machine learning."

Yann LeCun, Director, Facebook AI

<https://medium.com/@devnag/generative-adversarial-networks-gans-in-50-lines-of-code-pytorch-e81b79659e3f>

The GAN Father

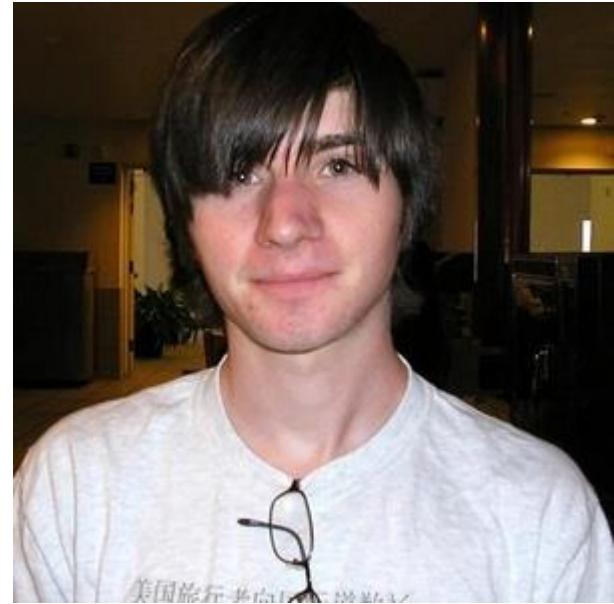


Ian Goodfellow

Goodfellow obtained his B.S. and M.S. in computer science from Stanford University and his Ph.D. in machine learning from the Université de Montréal, under the supervision of Yoshua Bengio and Aaron Courville. After graduation, Goodfellow joined Google as part of the Google Brain research team. He then left Google to join the newly founded OpenAI institute. He returned to Google Research in March 2017.

“I invented generative adversarial networks”

Story of How the GAN was Invented



https://www.peekyou.com/ian_goodfellow



<http://bigbangtheory.wikia.com/wiki/File:Pos8.jpg>



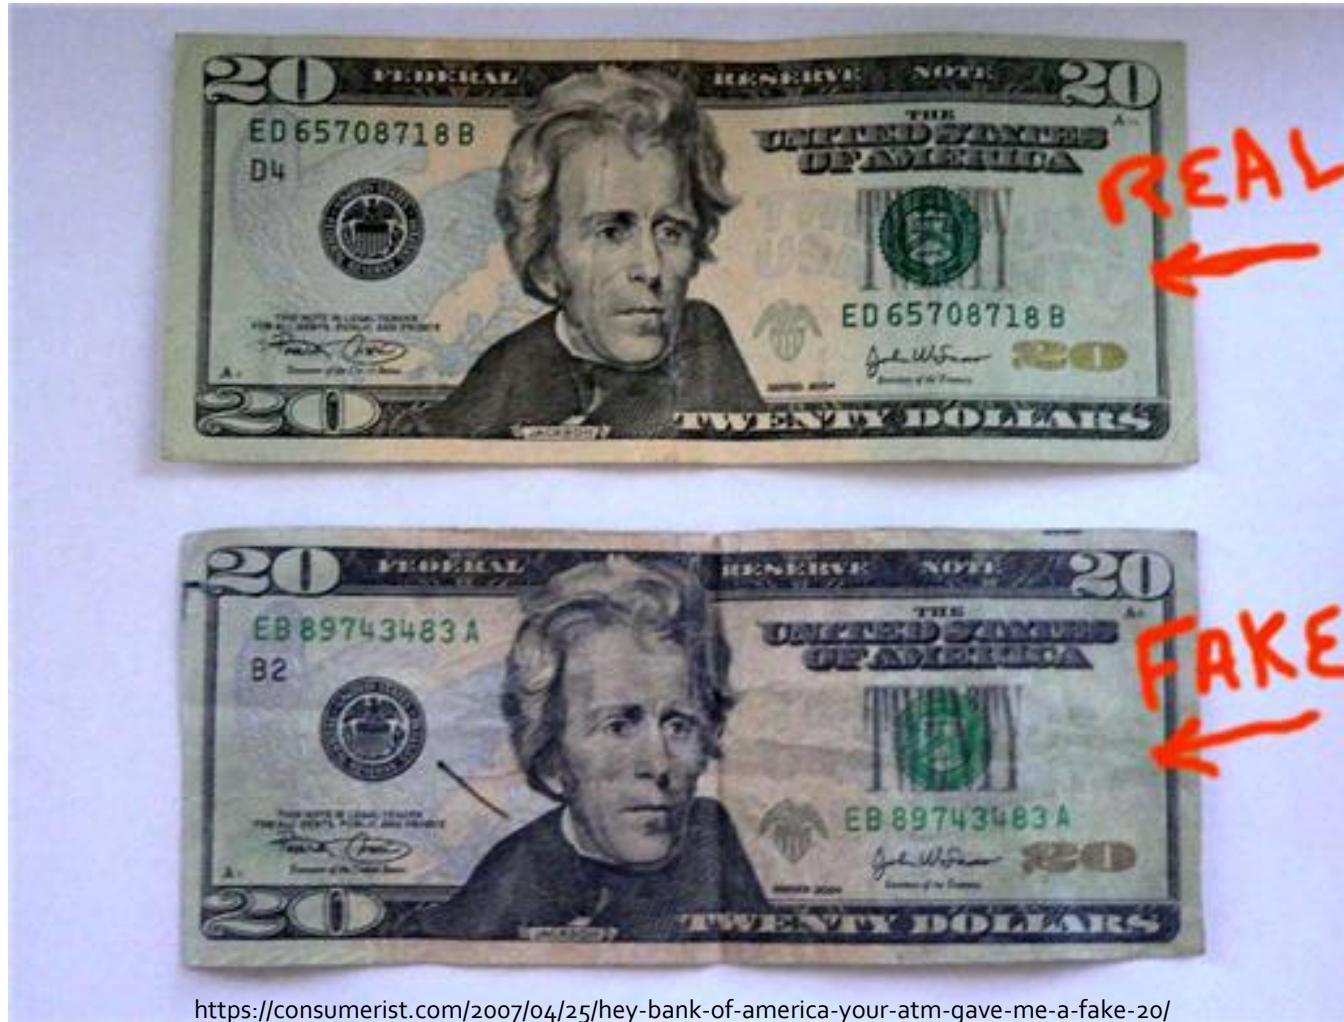
<http://tellmenothing.com/2017/05/13/couple-sleep-memes-night-owls/>

Example



<https://consumerist.com/2007/04/25/hey-bank-of-america-your-atm-gave-me-a-fake-20/>

Example



Example



<https://www.wisebread.com/how-to-spot-counterfeit-money>

Next Iteration



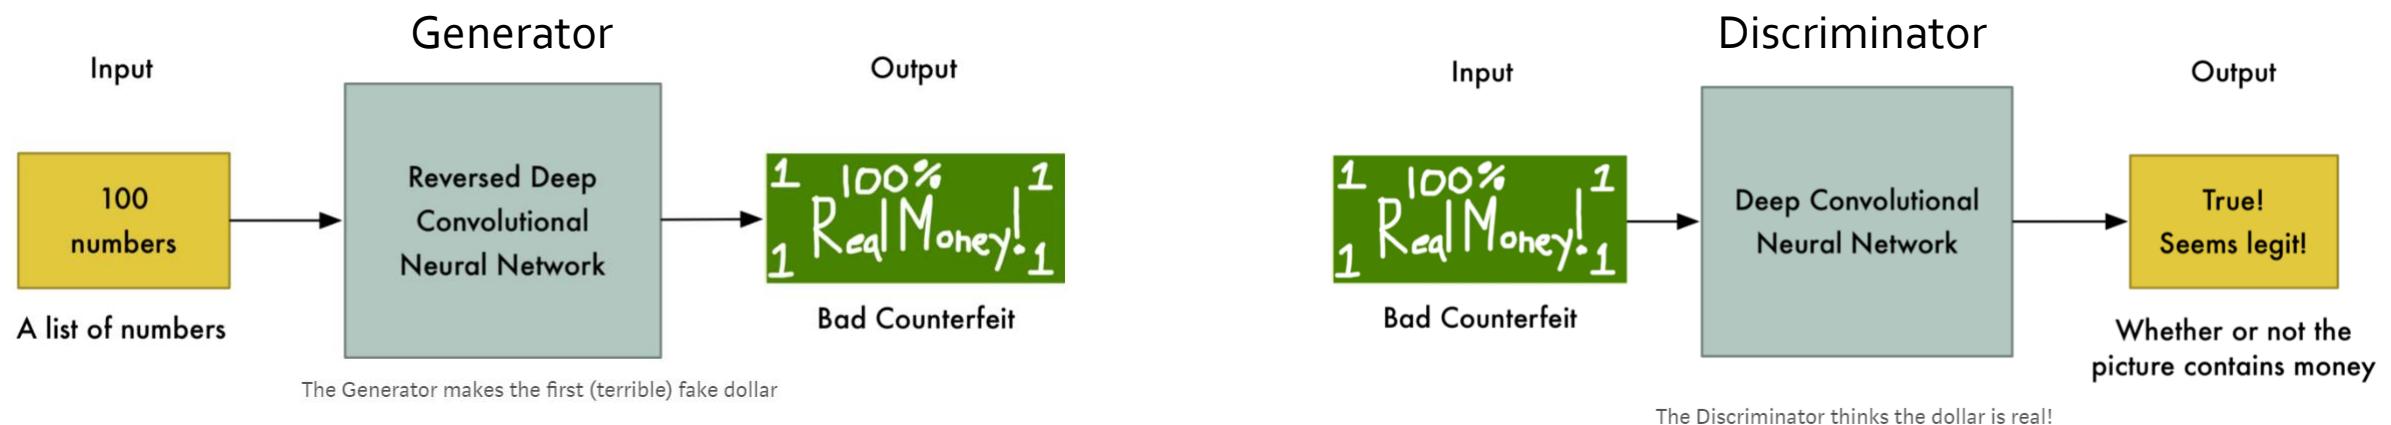
<https://www.youtube.com/watch?v=oOxo6RbJ6is>

Next Iteration



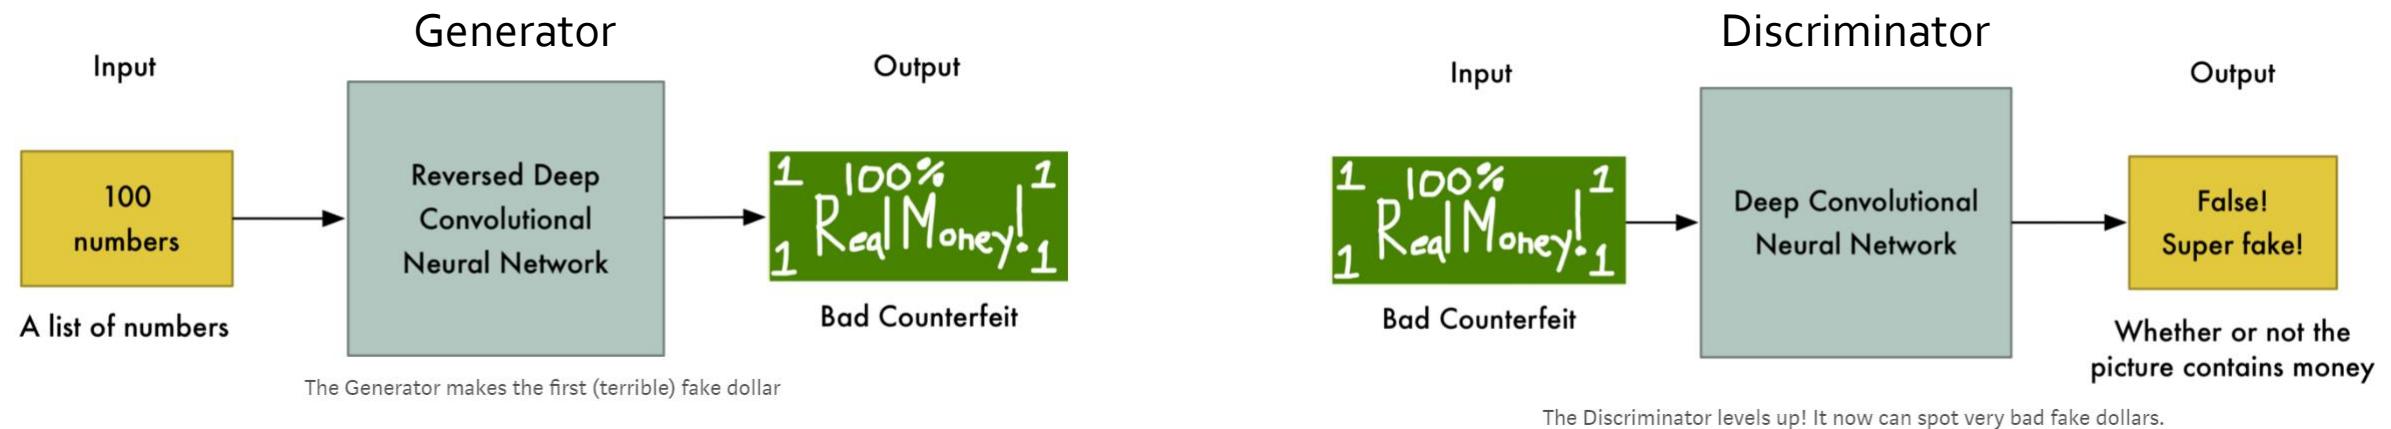
<https://www.ebay.com/bhp/counterfeit-money>

Example



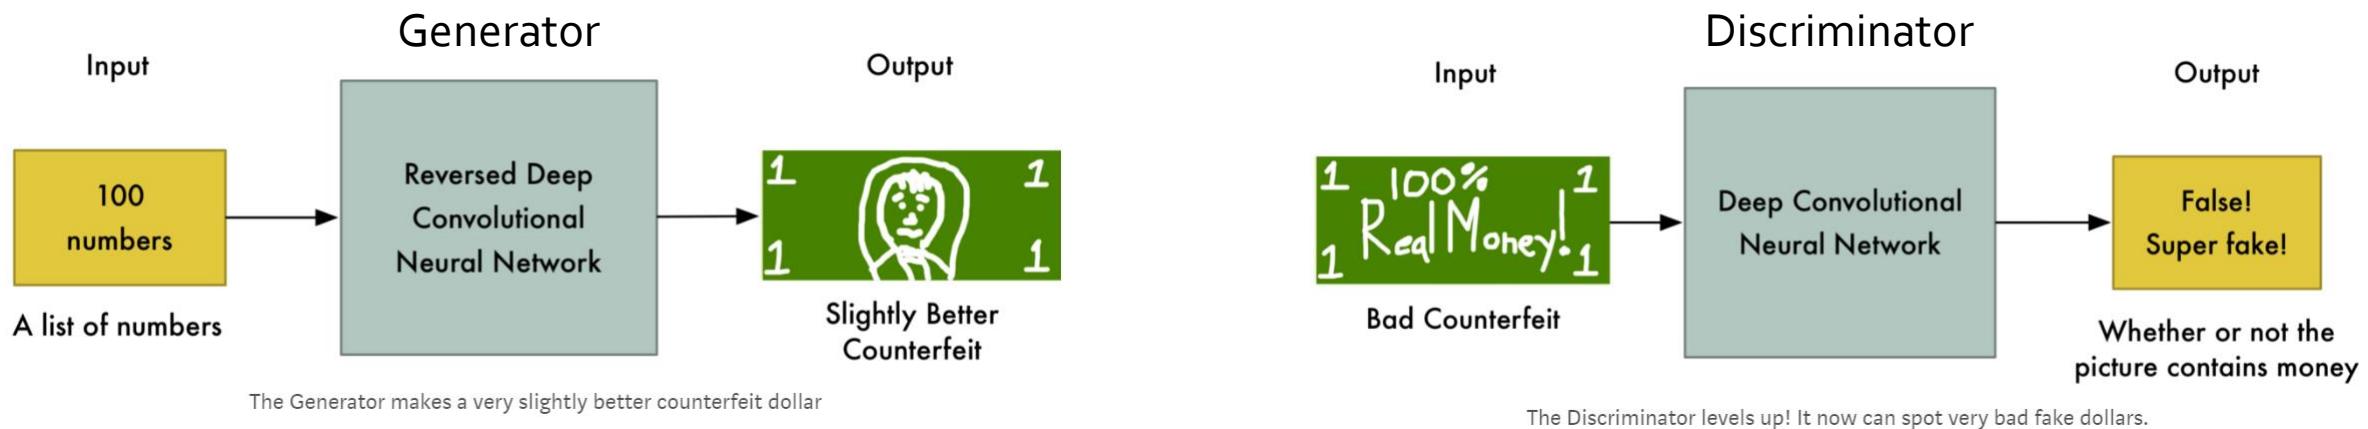
We tell Discriminator that this one is fake, real money has a picture of a person on it and the fake money doesn't.

Update Discriminator



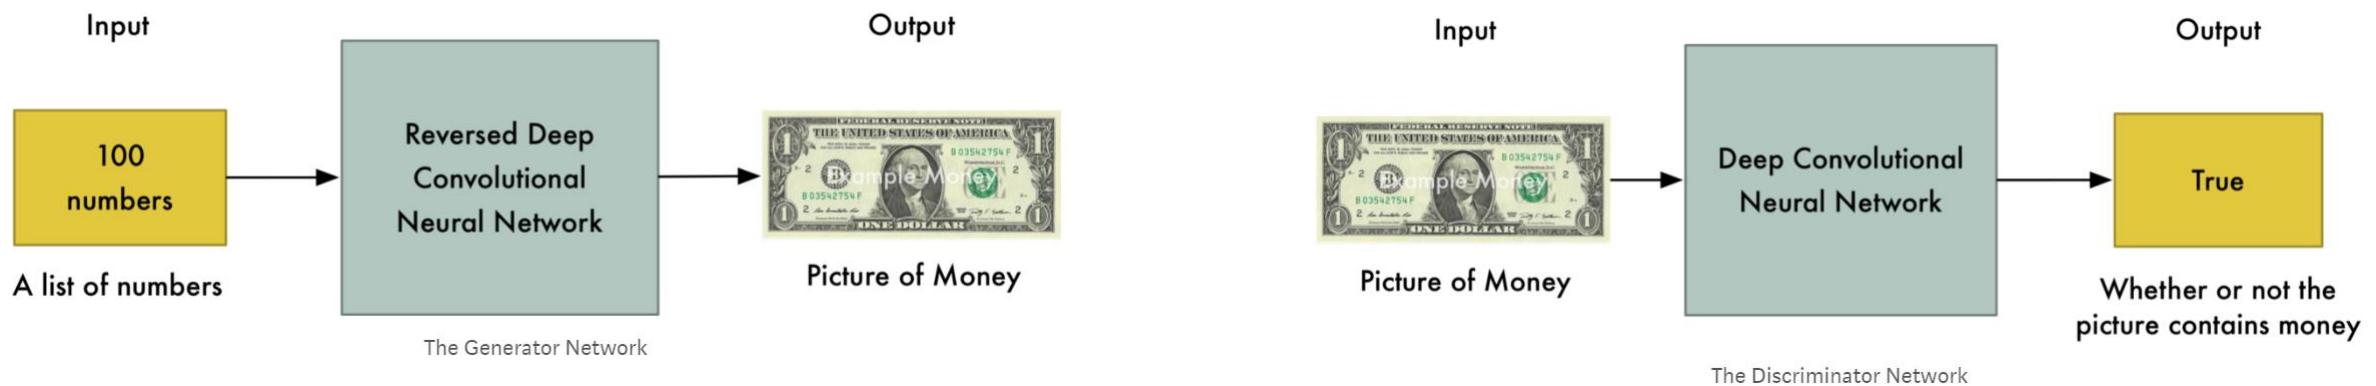
We tell the Generator that it's money images are suddenly getting rejected as fake so it needs to step up it's game. We also tell it that the Discriminator is now looking for faces

Update Generator



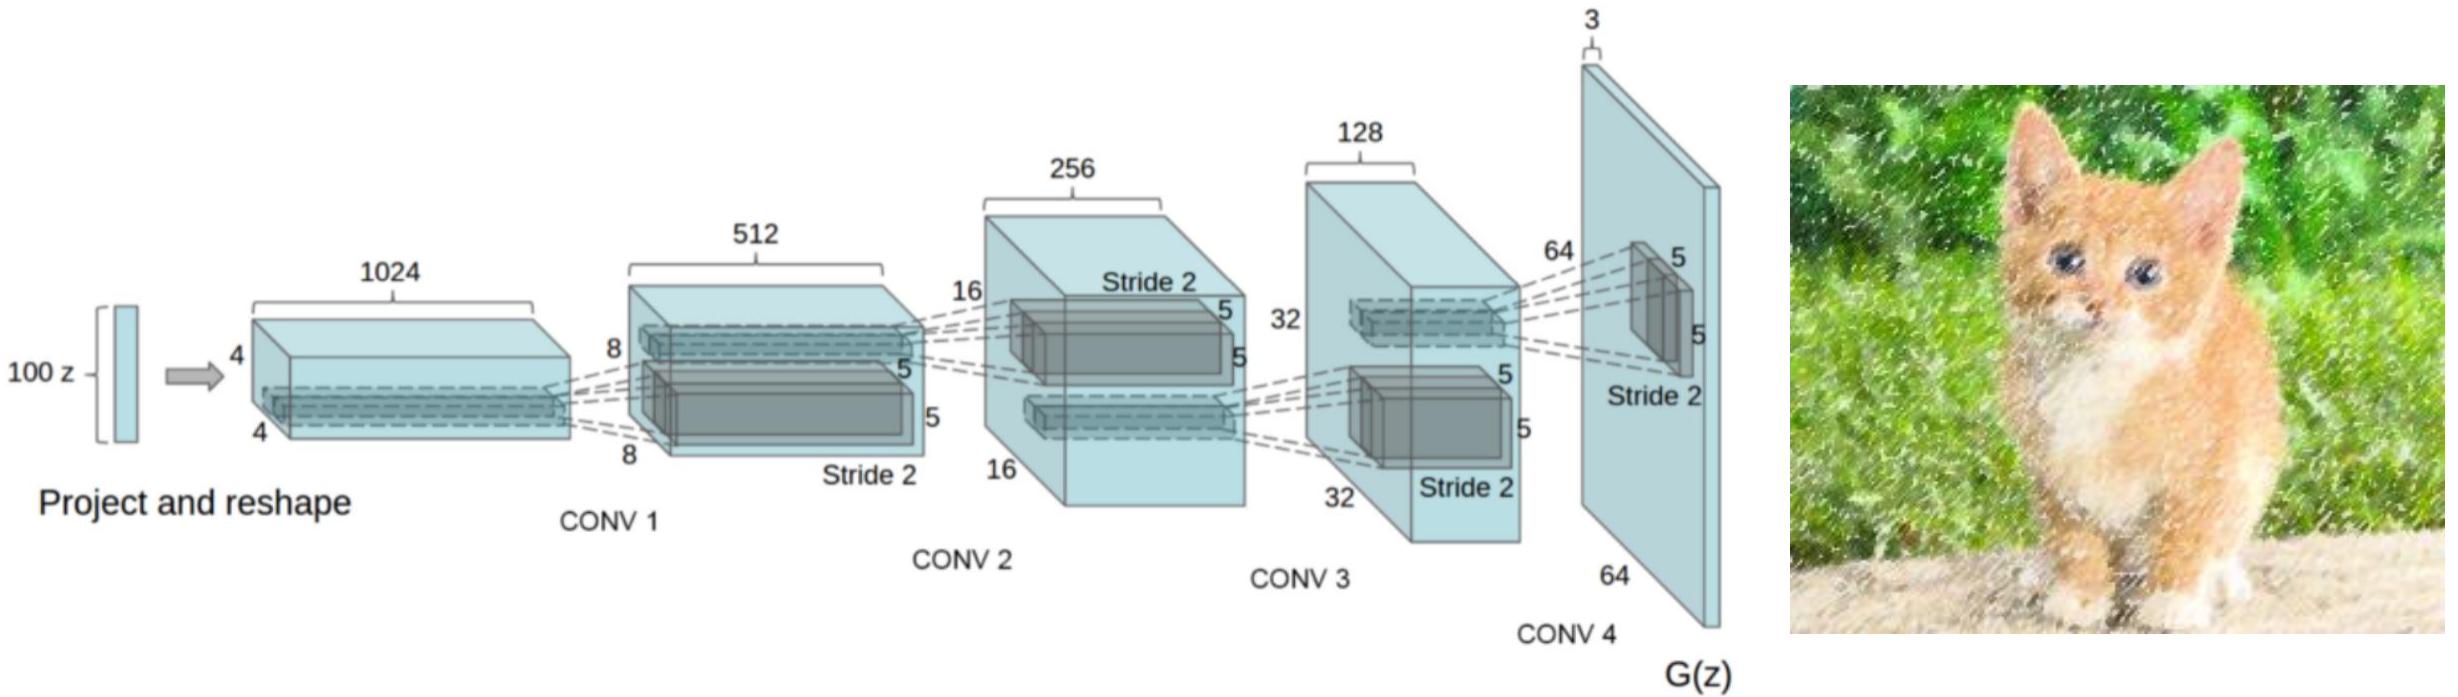
And the fake bills are being accepted as valid again! So now the Discriminator has to look again at the real dollar and find a new way to tell it apart from the fake one.

Final Purpose



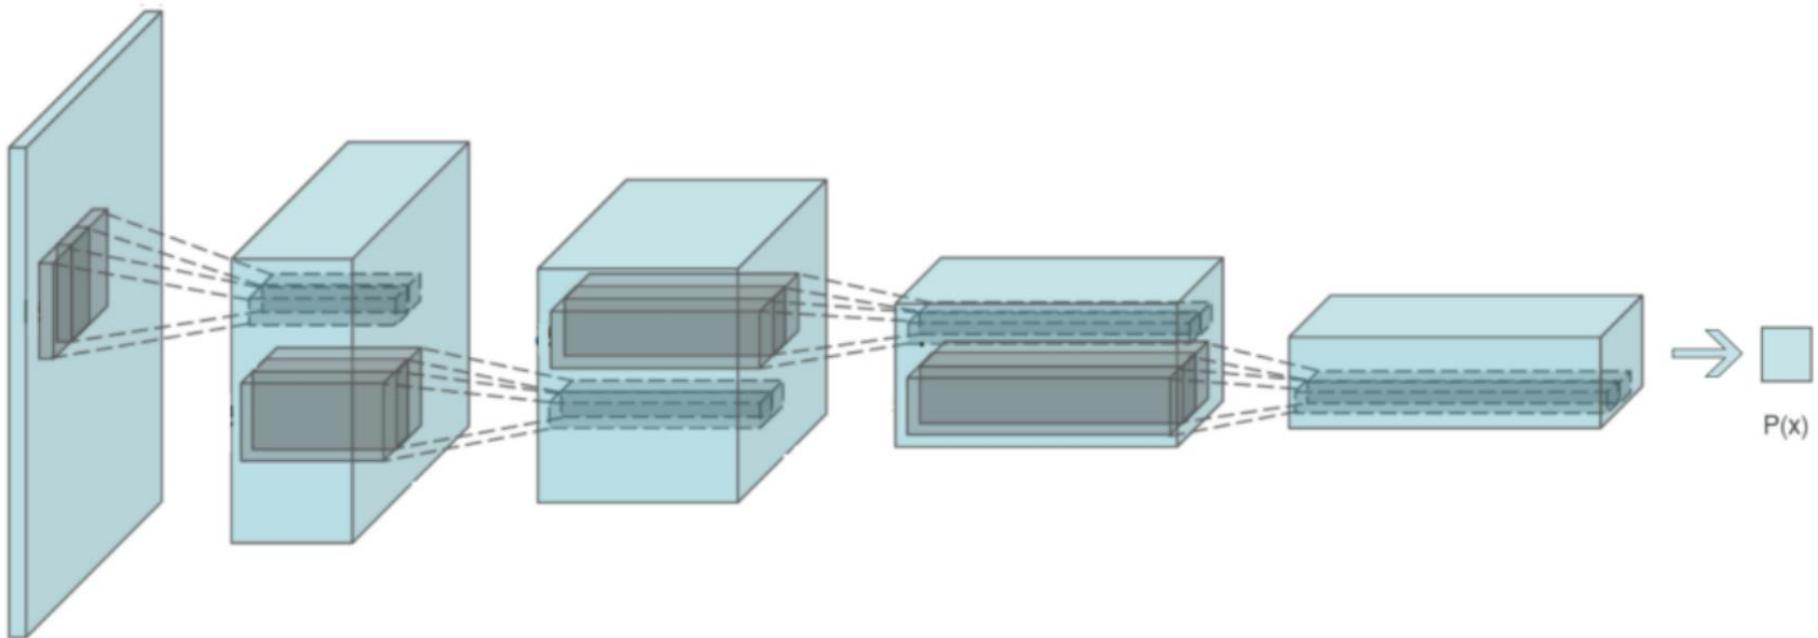
<https://medium.com/@ageitgey/abusing-generative-adversarial-networks-to-make-8-bit-pixel-art-e45d9b96cee7>

Generator



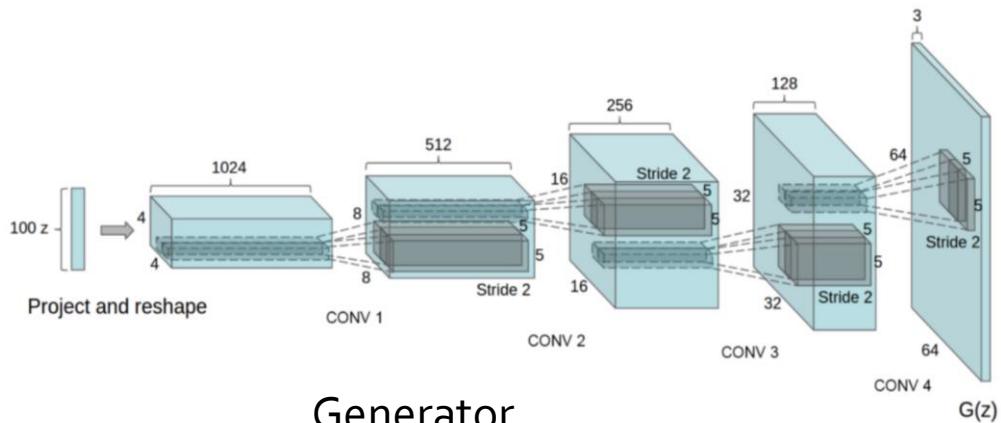
<https://arxiv.org/abs/1511.06434>

Discriminator

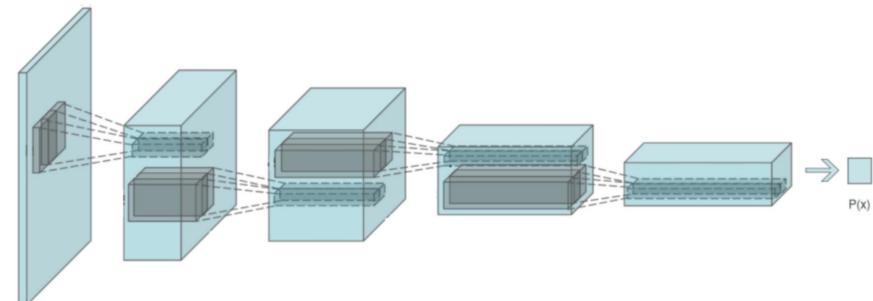


<https://arxiv.org/abs/1511.06434>

GAN



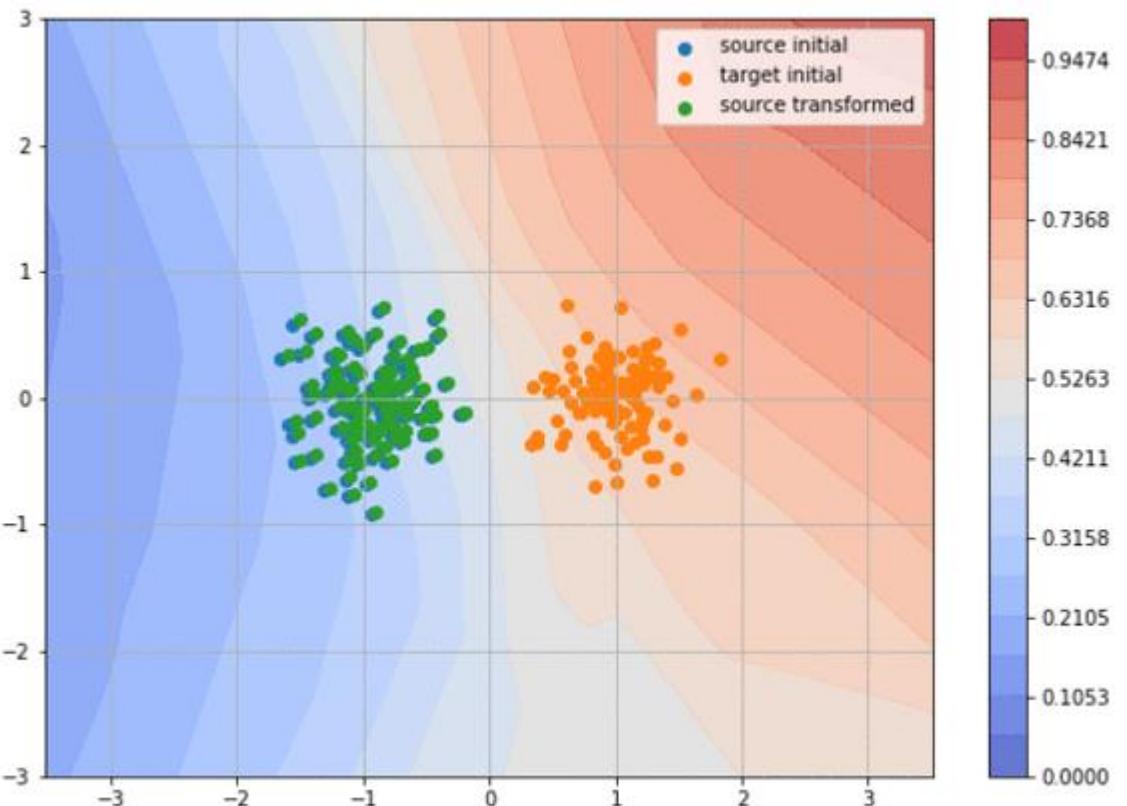
Generator



Discriminator

<https://arxiv.org/abs/1511.06434>

Domain Adaptation



GAN-based adversarial domain adaptation for two Gaussian domains. The discriminator (background) tries to separate the green distribution from the orange distribution, and the generator modifies the green distribution to fool the discriminator.

VAE vs. GAN

Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

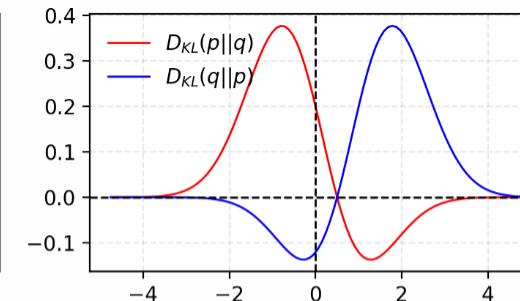
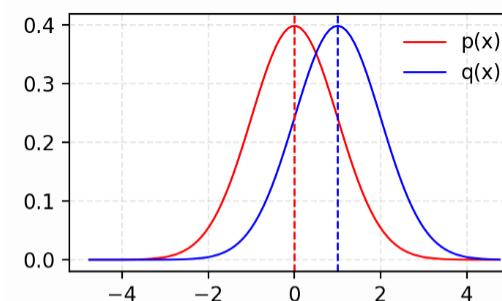
Cross-Entropy:
 $H(p, q) = -\sum_i p_i \log_2(q_i)$

KL Divergence:
 $D_{KL}(p || q) = H(p, q) - H(p)$

Cross-Entropy = Entropy + KL Divergence

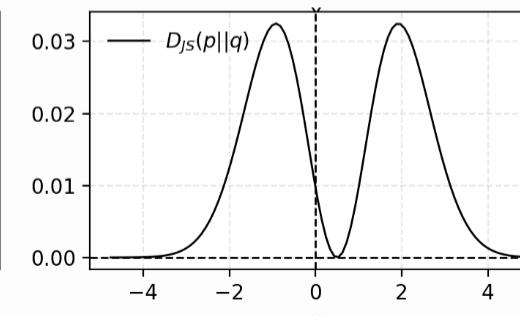
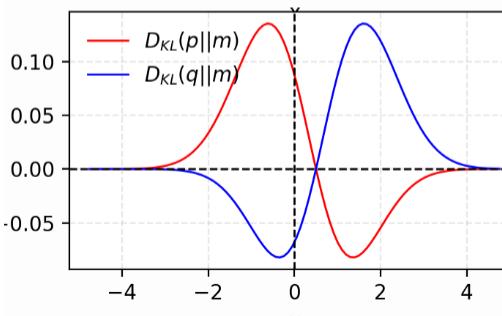
VAE:

$$D_{KL}(p||q) = \int_x p(x) \log \frac{p(x)}{q(x)} dx$$



GAN:

$$D_{JS}(p||q) = \frac{1}{2}D_{KL}(p\|\frac{p+q}{2}) + \frac{1}{2}D_{KL}(q\|\frac{p+q}{2})$$



<https://lilianweng.github.io/lil-log/2017/08/20/from-GAN-to-WGAN.html>
<https://www.youtube.com/watch?v=ErfnhcEV1O8>

GAN paper

Generative Adversarial Nets

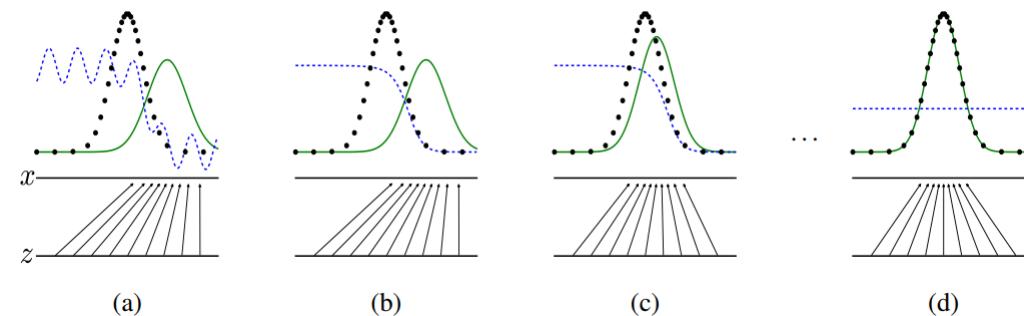
Ian J. Goodfellow, Jean Pouget-Abadie,^{*} Mehdi Mirza, Bing Xu, David Warde-Farley,

Sherjil Ozair,[†] Aaron Courville, Yoshua Bengio[‡]

Département d'informatique et de recherche opérationnelle

Université de Montréal

Montréal, QC H3C 3J7



In other words, D and G play the following two-player minimax game with value function $V(G, D)$:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]. \quad (1)$$

JS Divergence

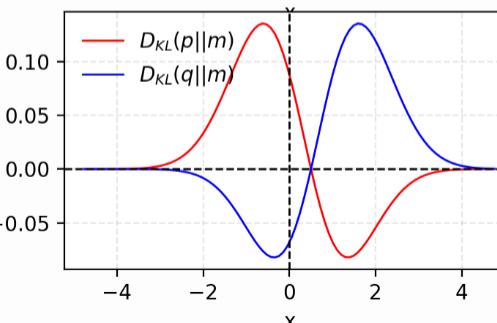
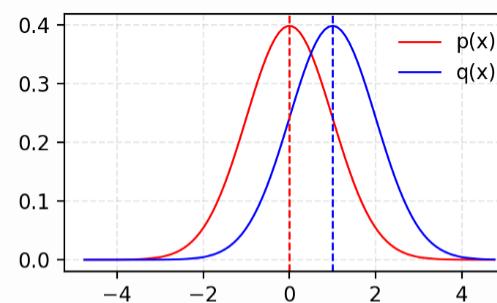
Entropy:
 $H(p) = -\sum_i p_i \log_2(p_i)$

Cross-Entropy:
 $H(p, q) = -\sum_i p_i \log_2(q_i)$

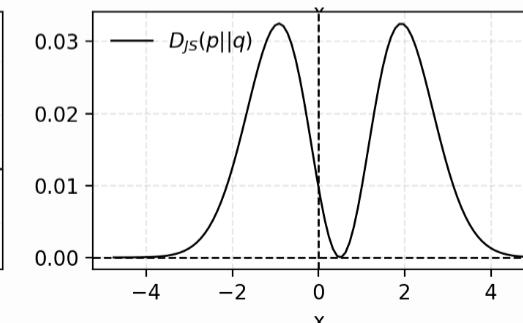
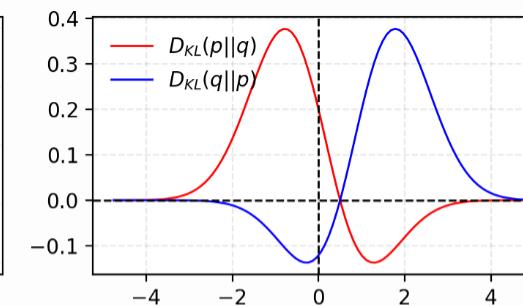
KL Divergence:
 $D_{KL}(p \parallel q) = H(p, q) - H(p)$

Cross-Entropy = Entropy + KL Divergence

$$D_{KL}(p \parallel q) = \int_x p(x) \log \frac{p(x)}{q(x)} dx$$



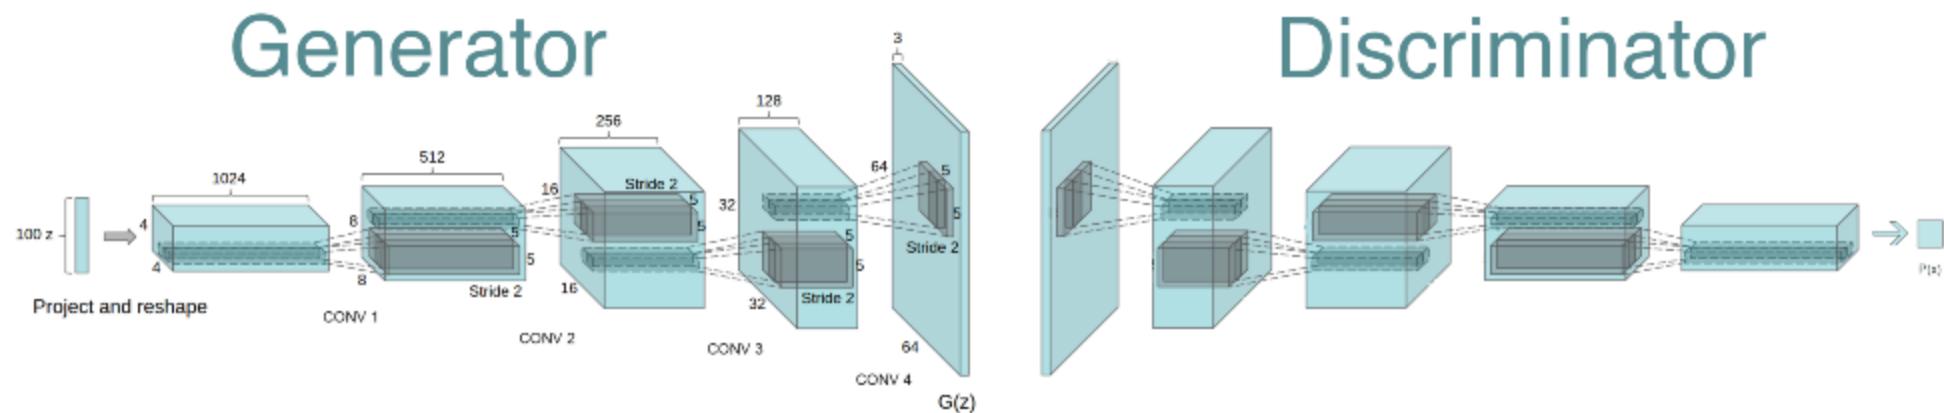
$$D_{JS}(p \parallel q) = \frac{1}{2} D_{KL}(p \parallel \frac{p+q}{2}) + \frac{1}{2} D_{KL}(q \parallel \frac{p+q}{2})$$



<https://lilianweng.github.io/lil-log/2017/08/20/from-GAN-to-WGAN.html>
<https://www.youtube.com/watch?v=ErfnhcEV1O8>

DCGAN

rings convolution operation into GAN



It gets rid of max-pooling which destroys spatial information and hurts the image quality. Its main design includes: Replace all max pooling with convolutional stride, Use transposed convolution for upsampling, Eliminate fully connected layers, and Use Batch normalization **BN**

<https://arxiv.org/abs/1511.06434>

WGAN

Instead of adding noise, Wasserstein GAN (WGAN) proposes a new cost function using Wasserstein distance that has a smoother gradient everywhere. WGAN learns no matter the generator is performing or not.

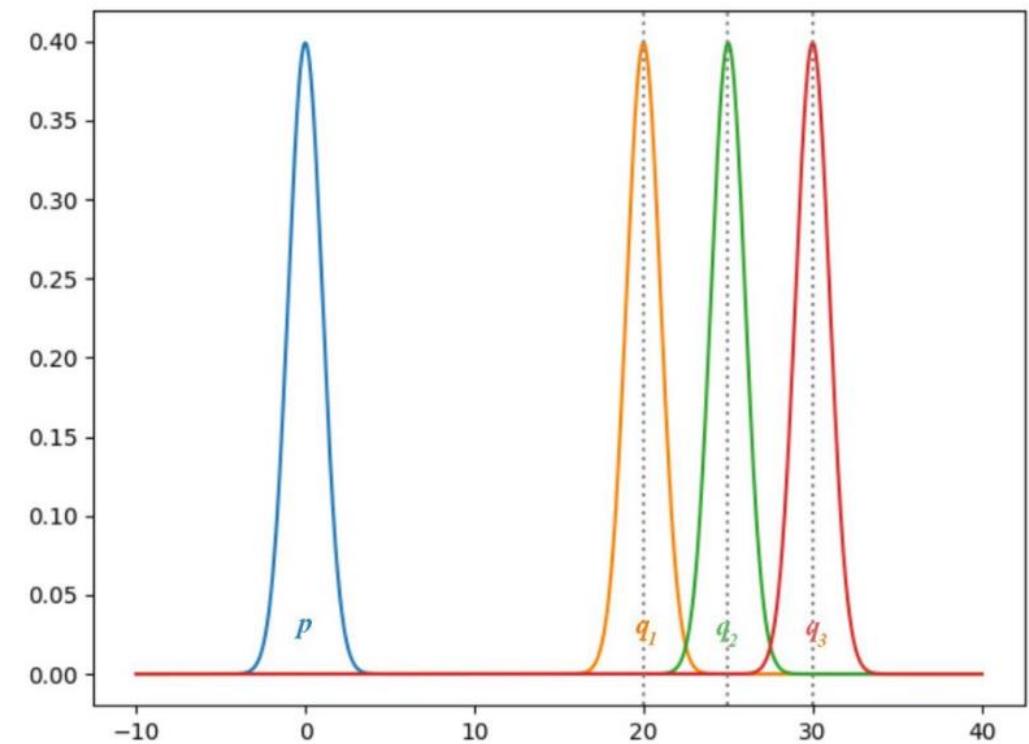
M. Arjovsky, S. Chintala, L. Bottou "Wasserstein GAN" 2016

GAN: minimize Jensen-Shannon divergence between p_X and $p_{G(Z)}$

$$JS(p_X || p_{G(Z)}) = KL(p_X || \frac{p_X + p_{G(Z)}}{2}) + KL(p_{G(Z)} || \frac{p_X + p_{G(Z)}}{2})$$

WGAN: minimize earth mover distance between p_X and $p_{G(Z)}$

$$EM(p_X, p_{G(Z)}) = \inf_{\gamma \in \Pi(p_X, p_{G(Z)})} E_{(x,y) \sim \gamma} [||x - y||]$$



https://medium.com/@jonathan_hui/gan-wasserstein-gan-wgan-gp-6a1a2aa1b490

LSGAN

More Stable than GAN, converge faster than WGAN

Still use a classifier but replace cross-entropy loss with Euclidean loss.

Discriminator

$$\text{GAN} \quad \min_D E_{x \sim p_X} [-\log D(x)] + E_{z \sim p_Z} [-\log(1 - D(G(z)))]$$



$$\text{LSGAN} \quad \min_D E_{x \sim p_X} [(D(x) - 1)^2] + E_{z \sim p_Z} [D(G(z))^2]$$

Generator

$$\text{GAN} \quad \min_G E_{z \sim p_Z} [-\log D(G(z))]$$



$$\text{LSGAN} \quad \min_G E_{z \sim p_Z} [(D(G(z)) - 1)^2]$$

61

EBGAN

The diagram of an EBGAN with an auto-encoder discriminator is depicted in figure 1. The generator (G) takes random vector z as input, and transform it into a sample $G(z)$, for example an image. The discriminator (D), whose output is a scalar energy, takes either real or generated images, and estimates the energy value E accordingly, where $E \in \mathbb{R}$.

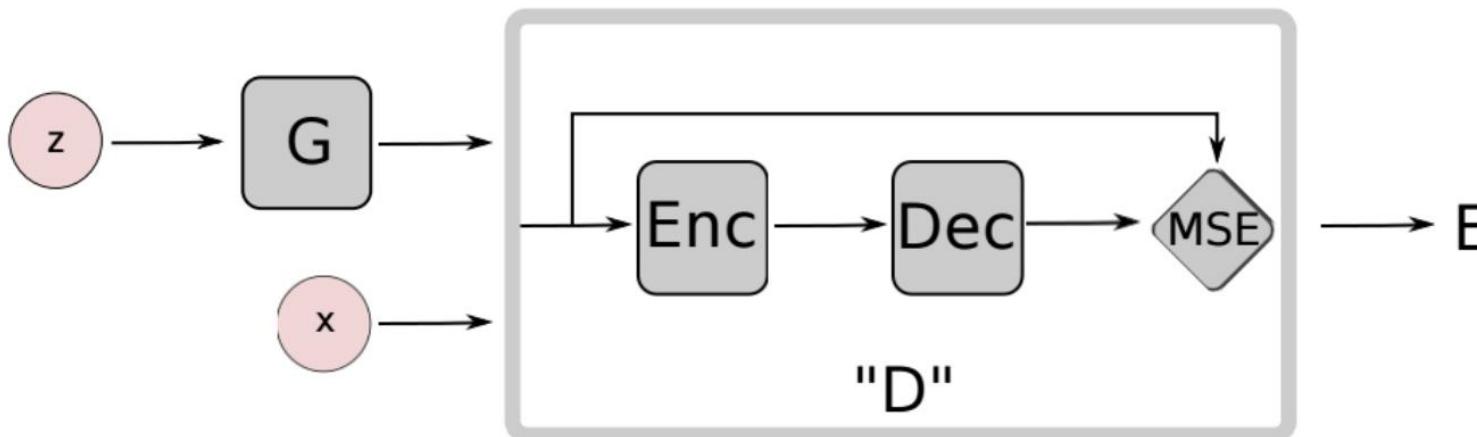
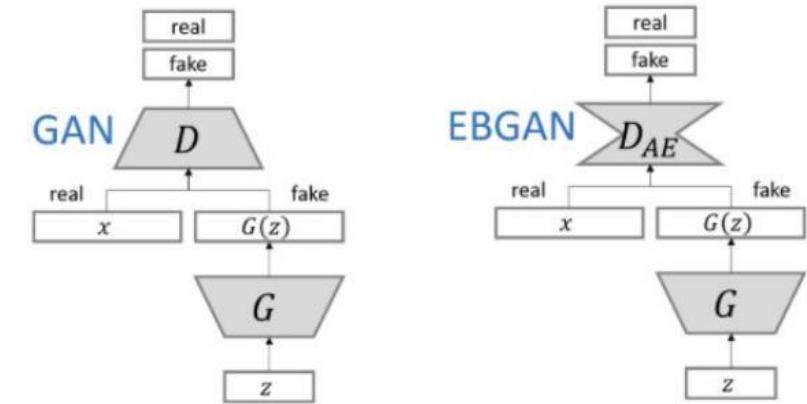
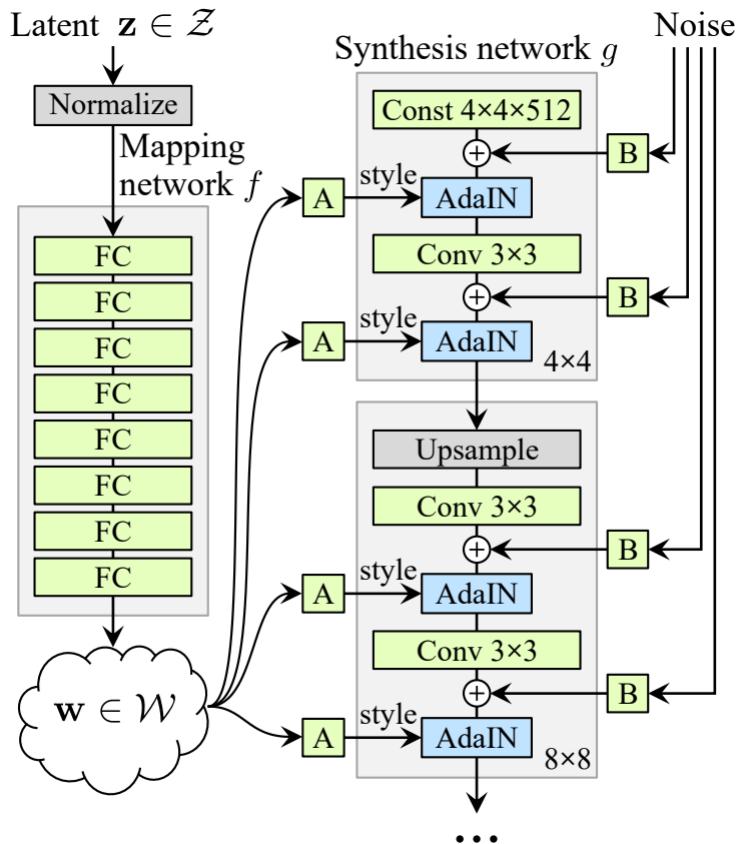


Figure 1: EBGAN architecture.

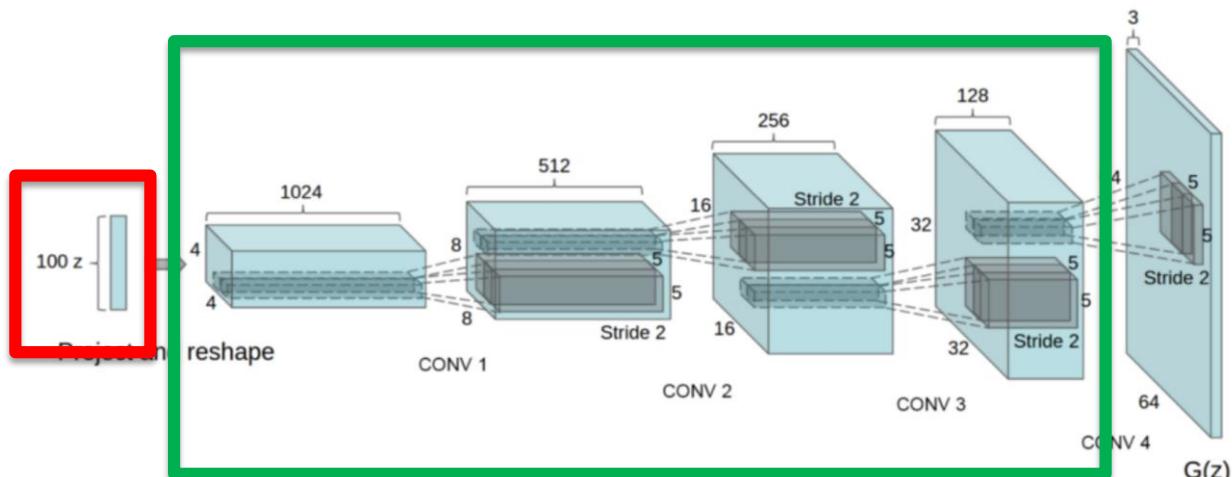
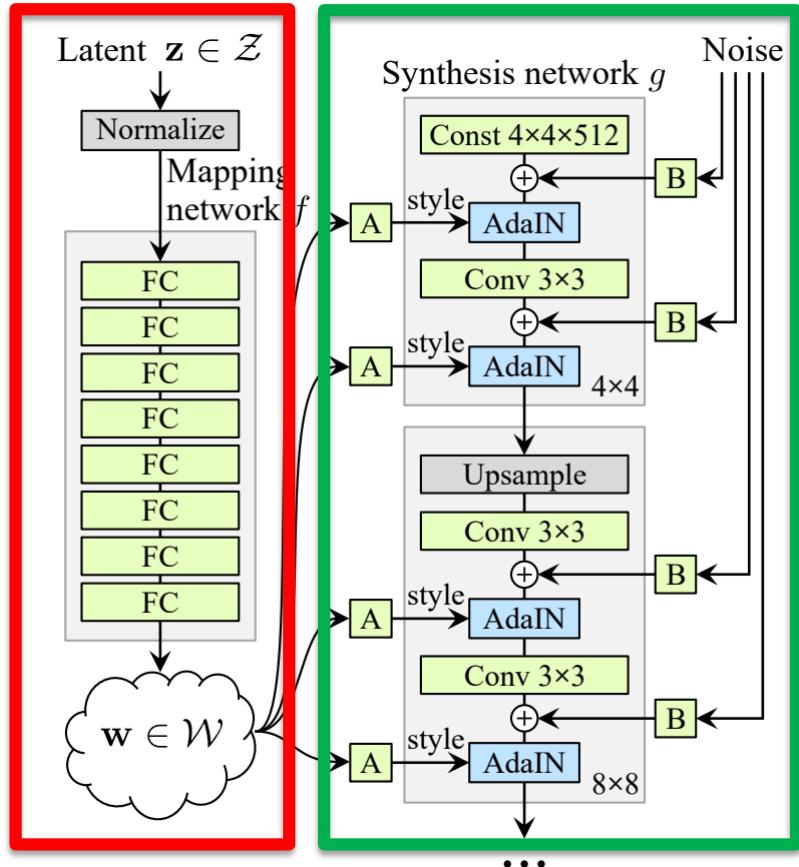


StyleGAN



<https://arxiv.org/abs/1812.04948>

StyleGAN



<https://arxiv.org/abs/1812.04948>

Noise to Medical Image

Learning Implicit Brain MRI Manifolds with Deep Learning

Camilo Bermudez^{*a}, Andrew J. Plassard^b, Larry T. Davis^c, Allen T. Newton^c, Susan M Resnick^b, Bennett A. Landman^a

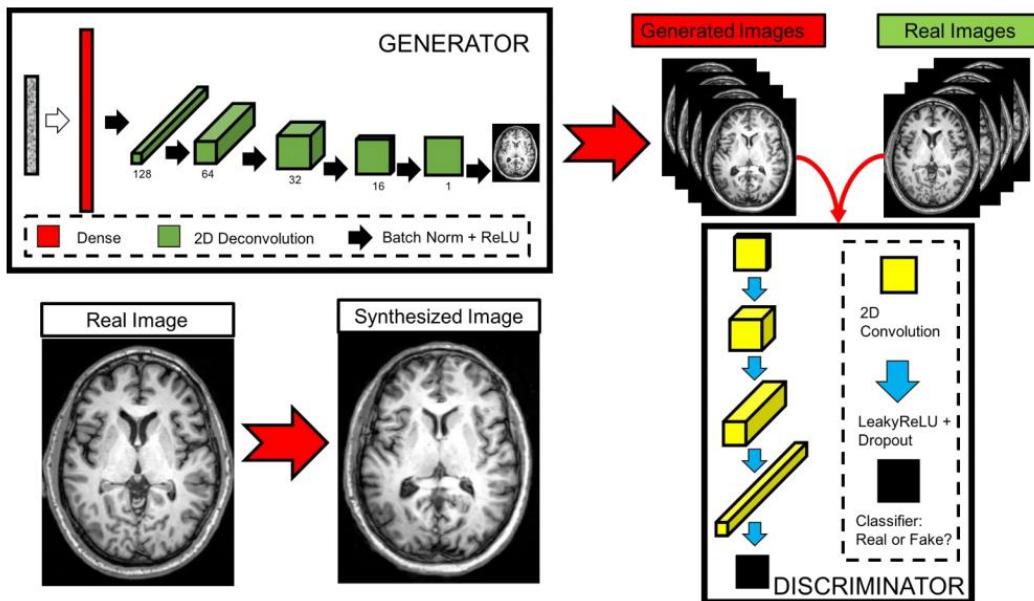


Figure 1. Pipeline and network architecture for Generative Adversarial Network used for image synthesis. The generator is a 2D deconvolutional neural network that takes noise as input and generates a 2D image of the brain. The discriminator is a convolutional neural network that takes real and synthetic images and learns to classify them.

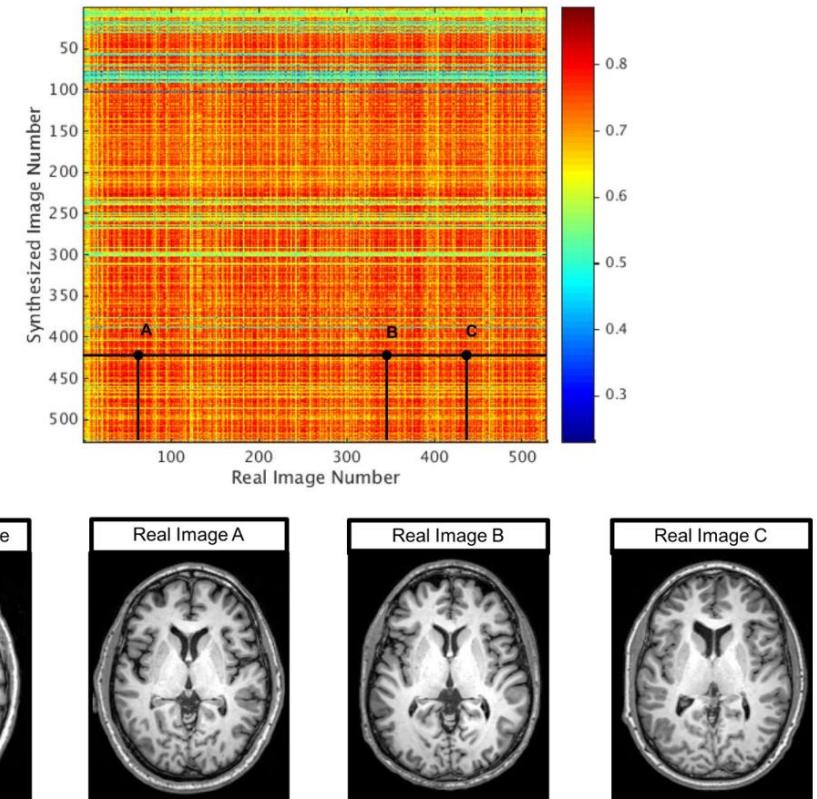
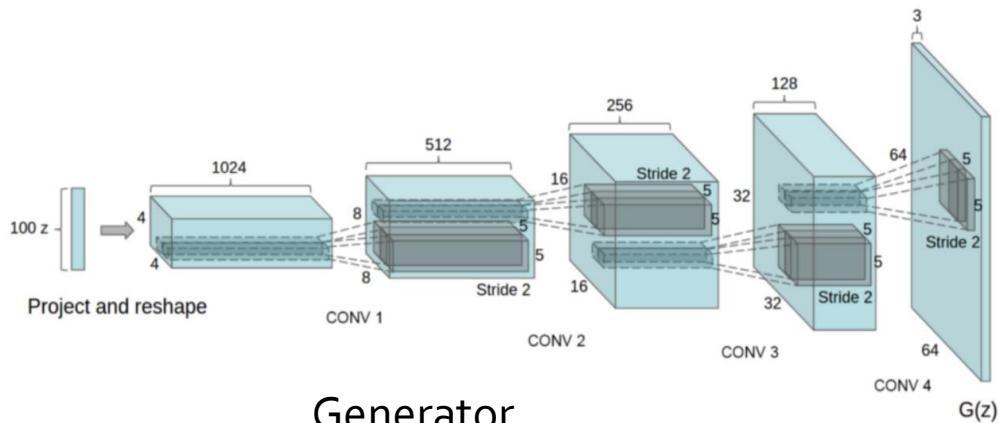


Figure 3. Correlation map between synthesized images and real images. Representative examples a synthetic image (top left) and three real images with the highest correlation values: Real Image A ($\rho = 0.76$), Real Image B ($\rho = 0.77$), and Real Image C ($\rho = 0.78$).

GAN



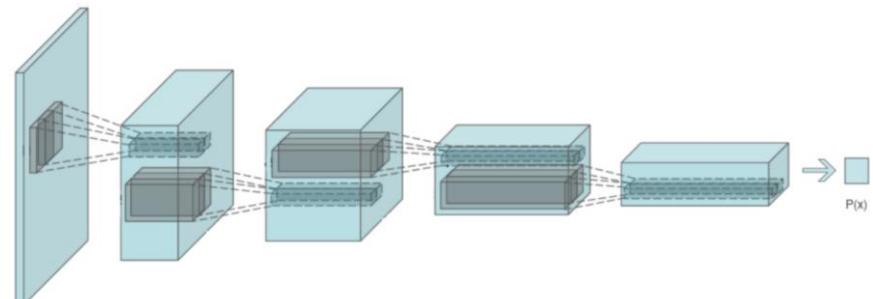
Generator



Fake



True



Discriminator

<https://arxiv.org/abs/1511.06434>

How?

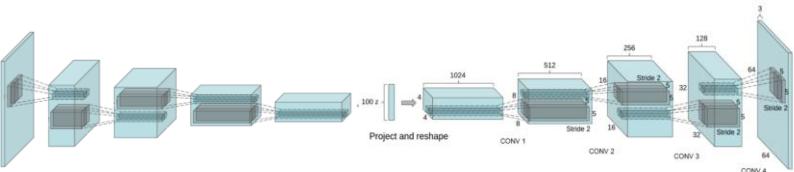


<https://gifer.com/en/SRL8>

How?



Image to Image GAN (Conditional GAN)



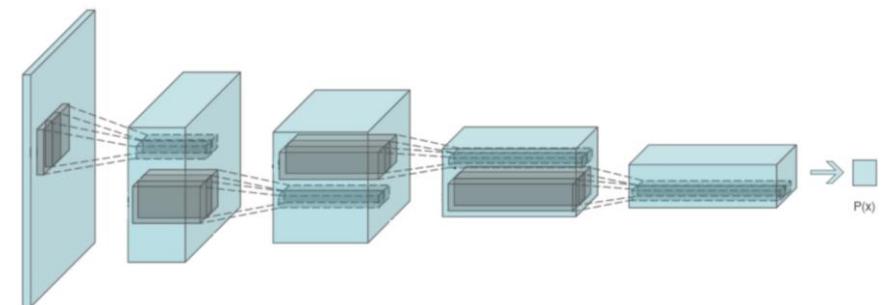
Generator



Fake



True

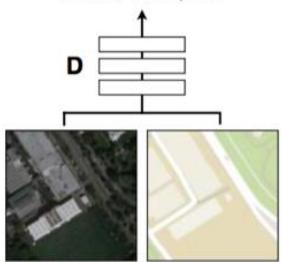


Discriminator

Pix-2-Pix GAN

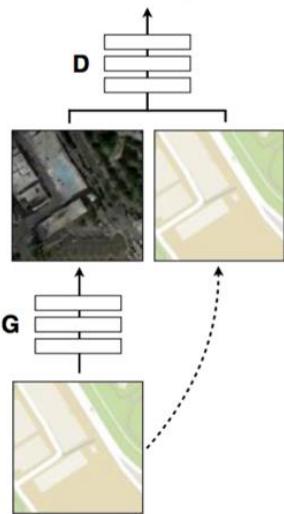
Positive examples

Real or fake pair?



Negative examples

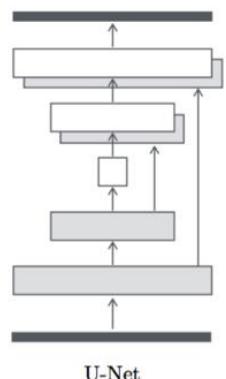
Real or fake pair?



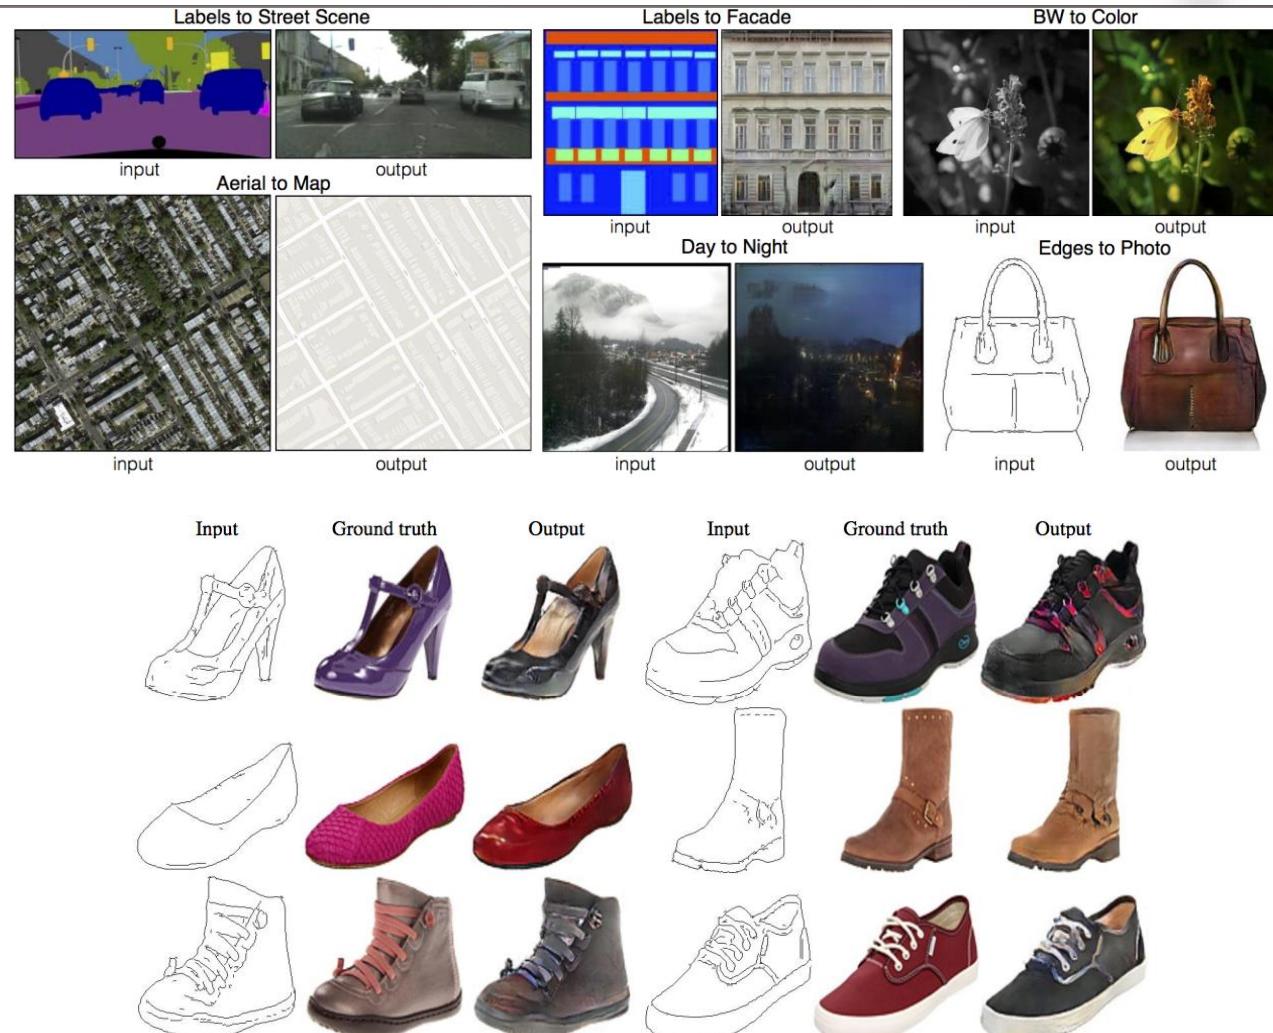
G tries to synthesize fake images that fool **D**

D tries to identify the fakes

Discriminator



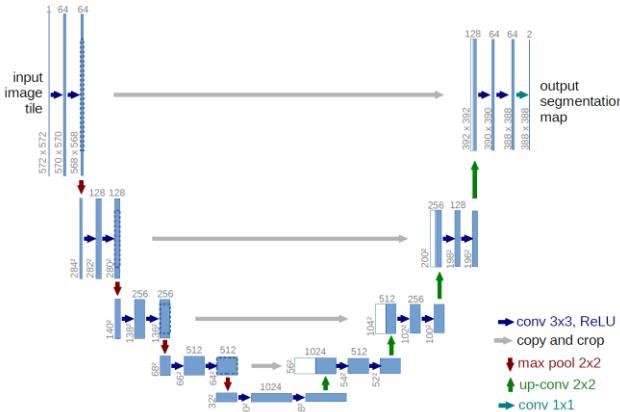
Generator



<https://arxiv.org/abs/1611.07004>

<https://tensorflow.blog/tag/pix2pix/>

How To



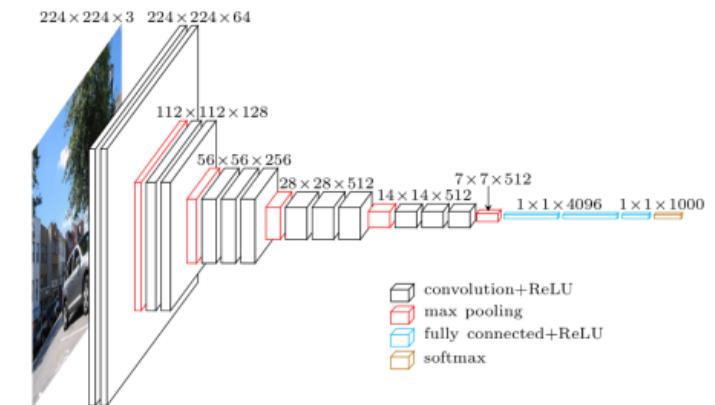
Generator



Fake

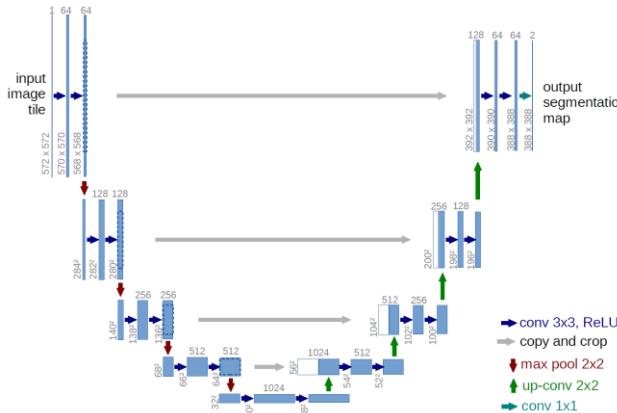


True

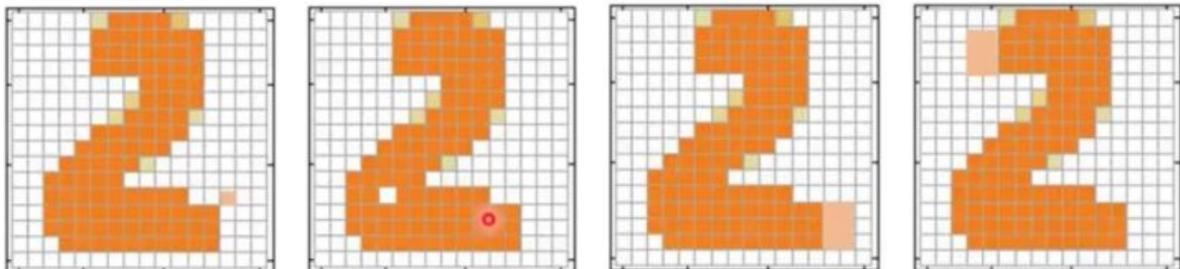


Discriminator

How To



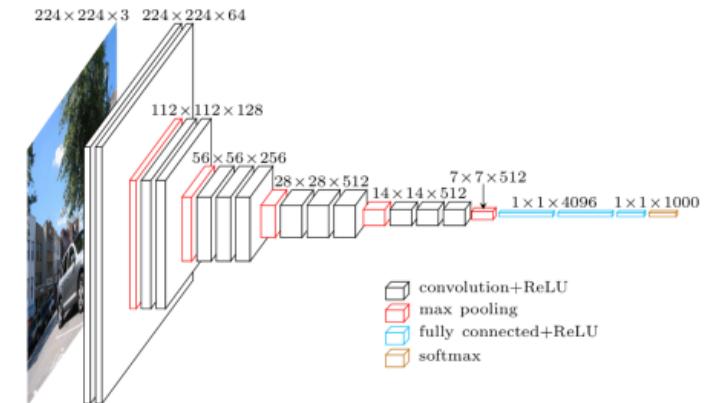
Generator



Fake



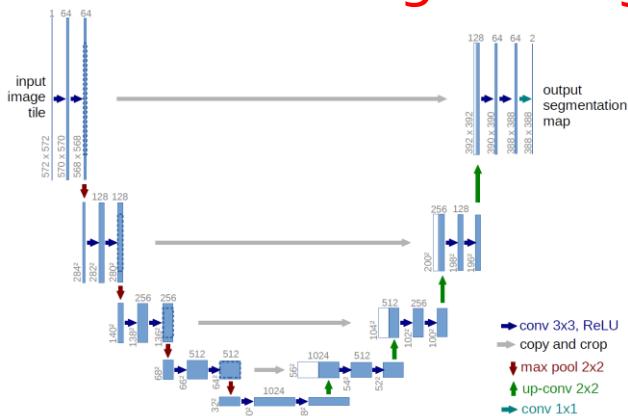
True



Discriminator

How To

Replace Feature Engineering!



Generator

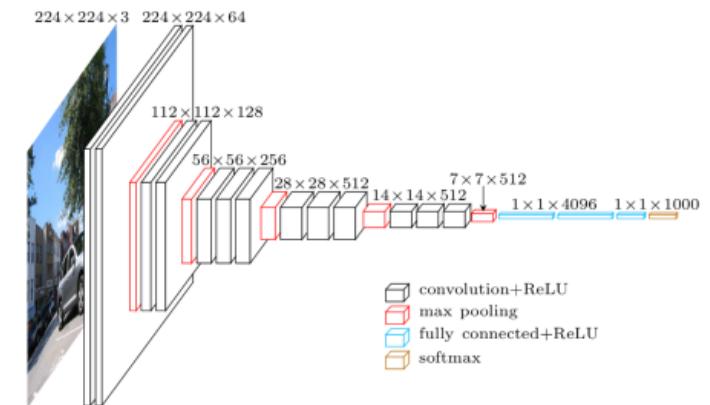


Fake



True

Replace Loss Engineering!

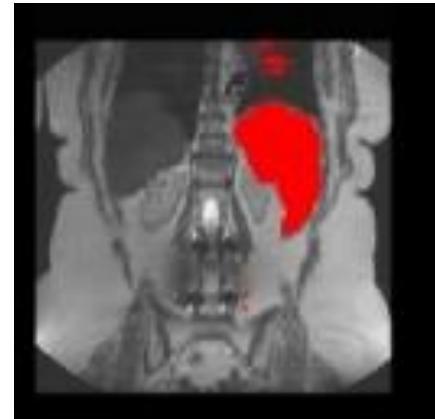
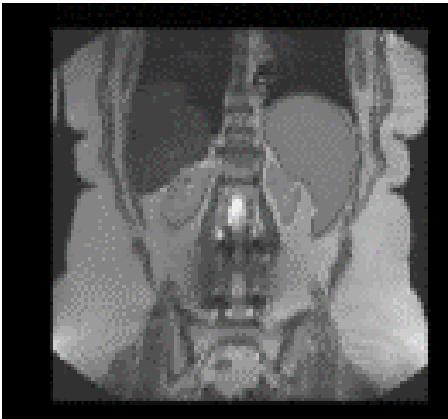


Discriminator

Patch GAN

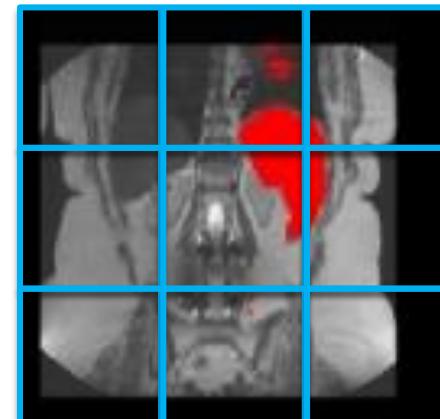
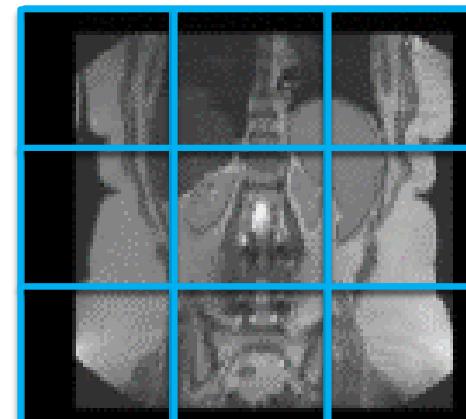
Isola et al, Archive 2017

Image GAN



o/1

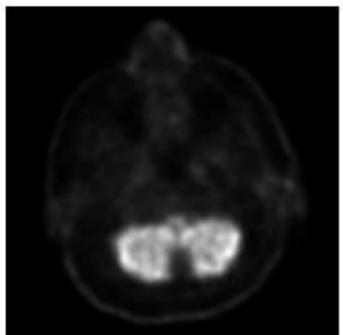
Patch GAN



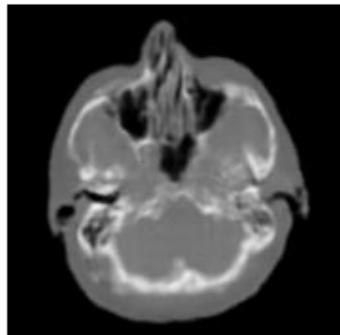
o/1	o/1	o/1
o/1	o/1	o/1
o/1	o/1	o/1

Pix2Pix on Medical Image

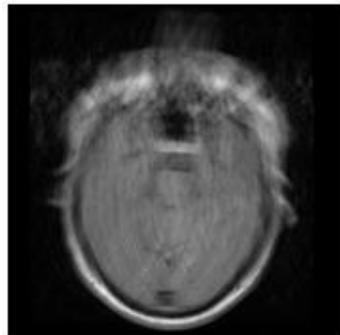
Input



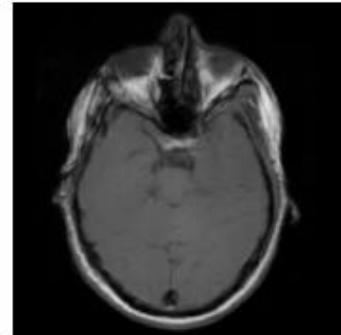
pix2pix



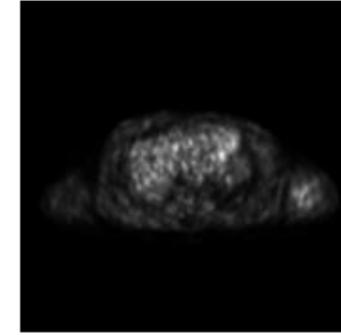
Input



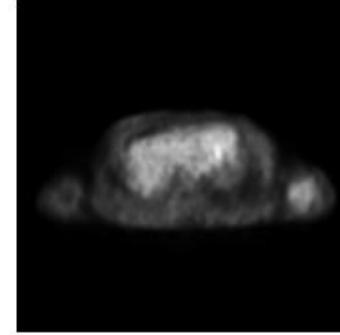
pix2pix



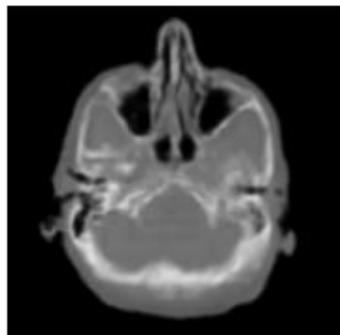
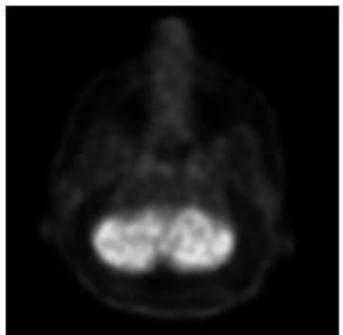
Input



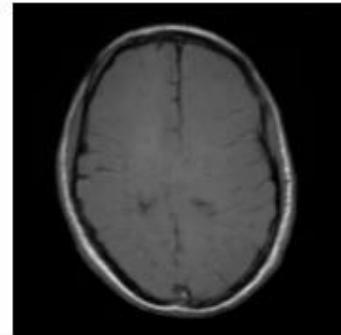
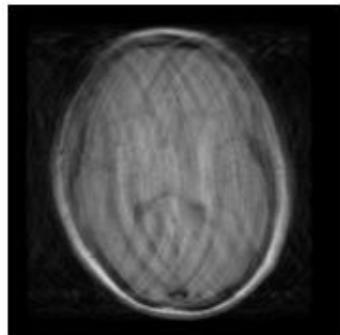
pix2pix



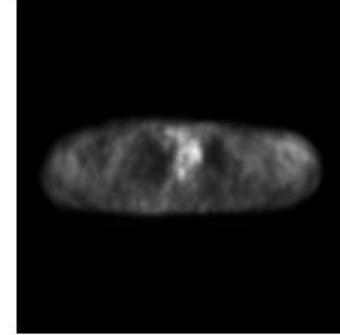
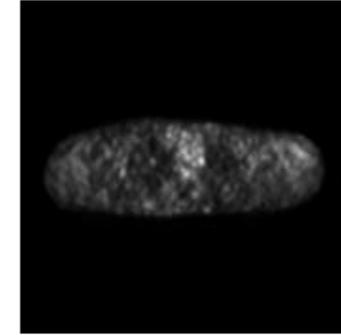
PET-CT translation



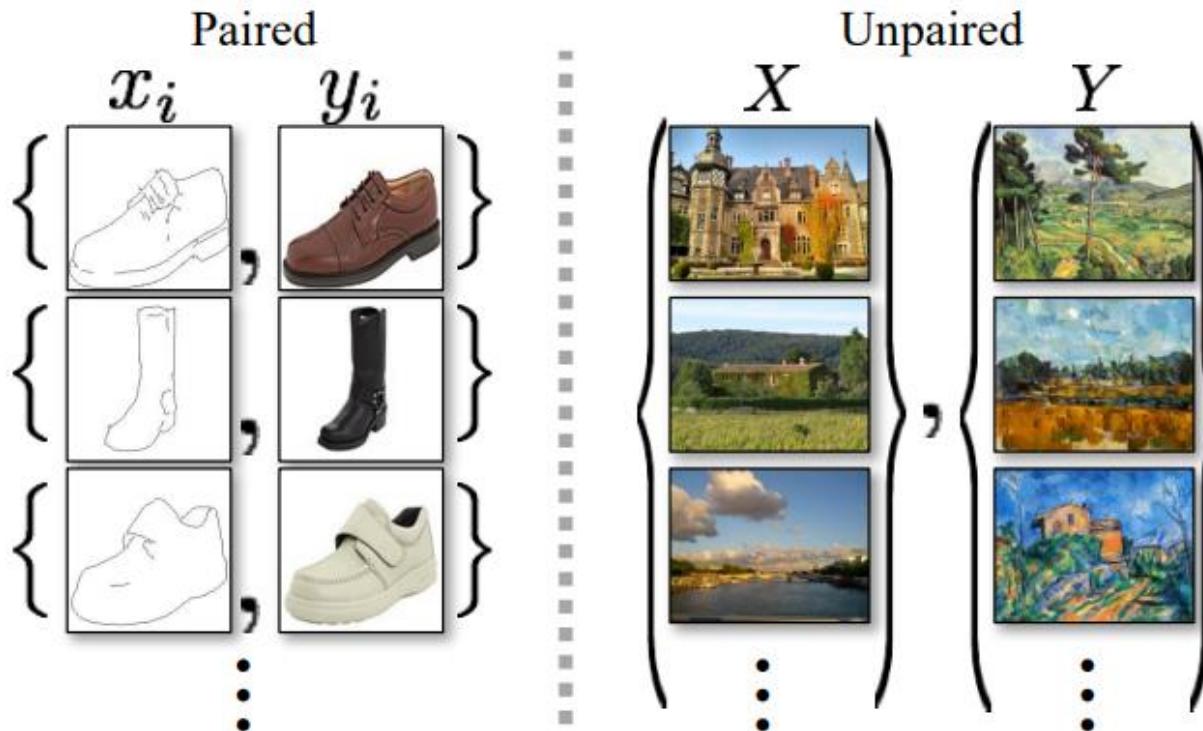
MR motion correction



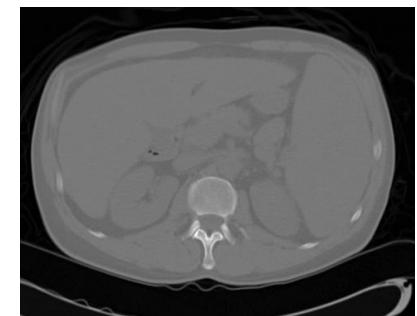
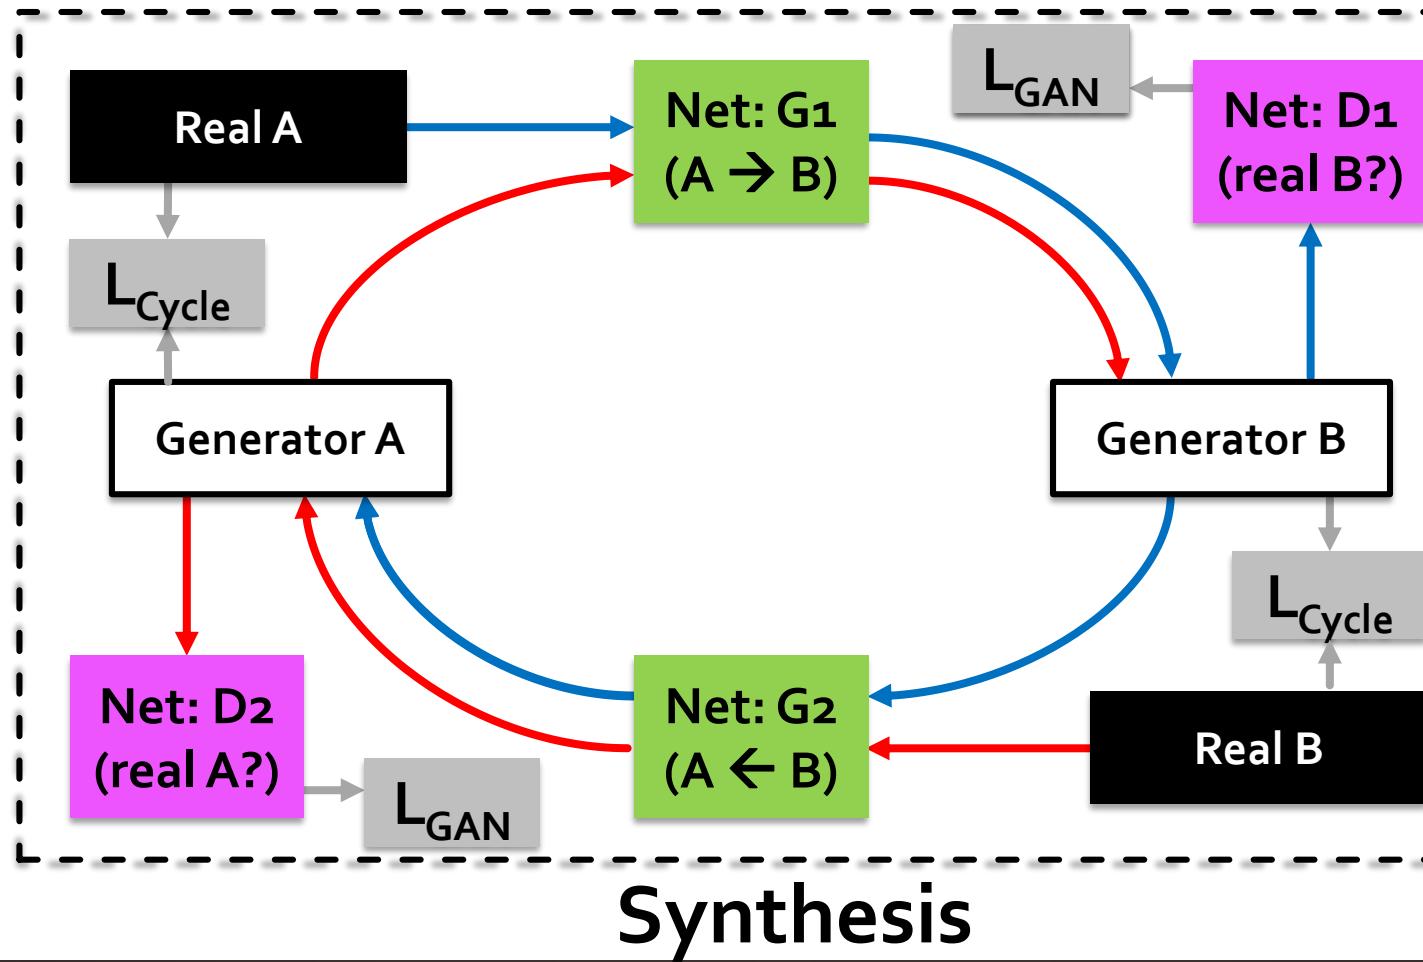
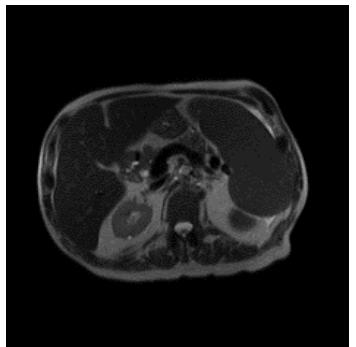
PET denoising



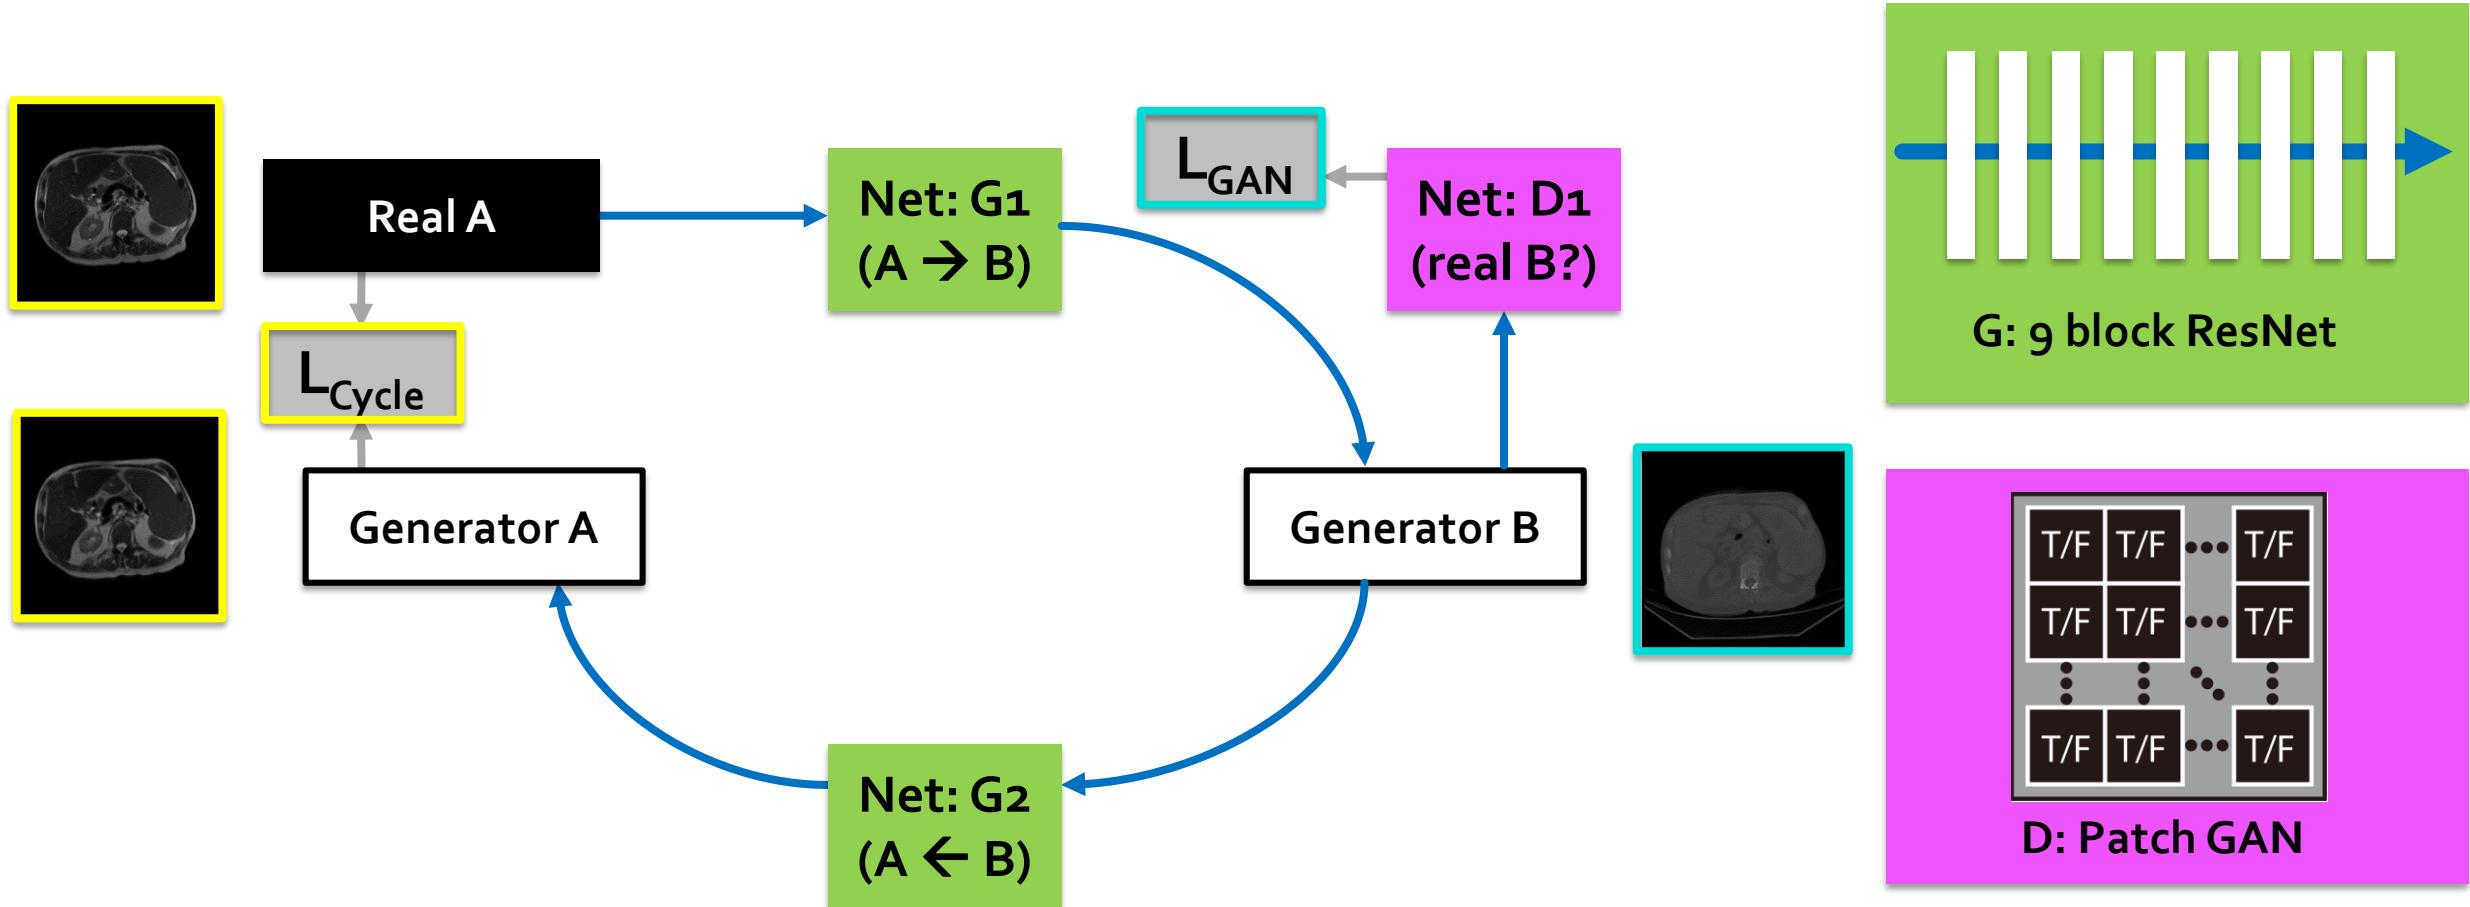
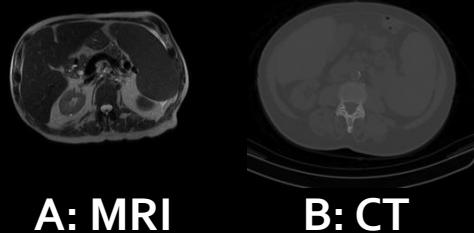
Unpaired Synthesis



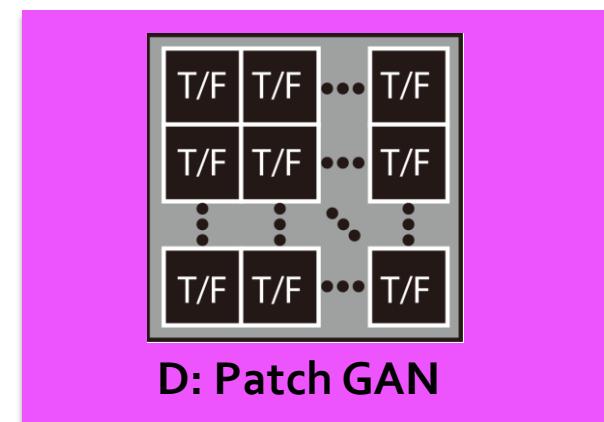
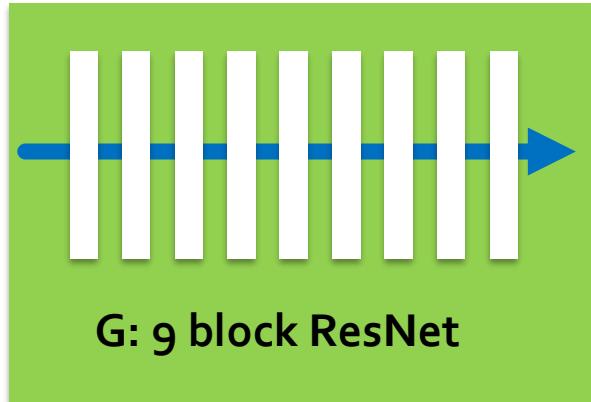
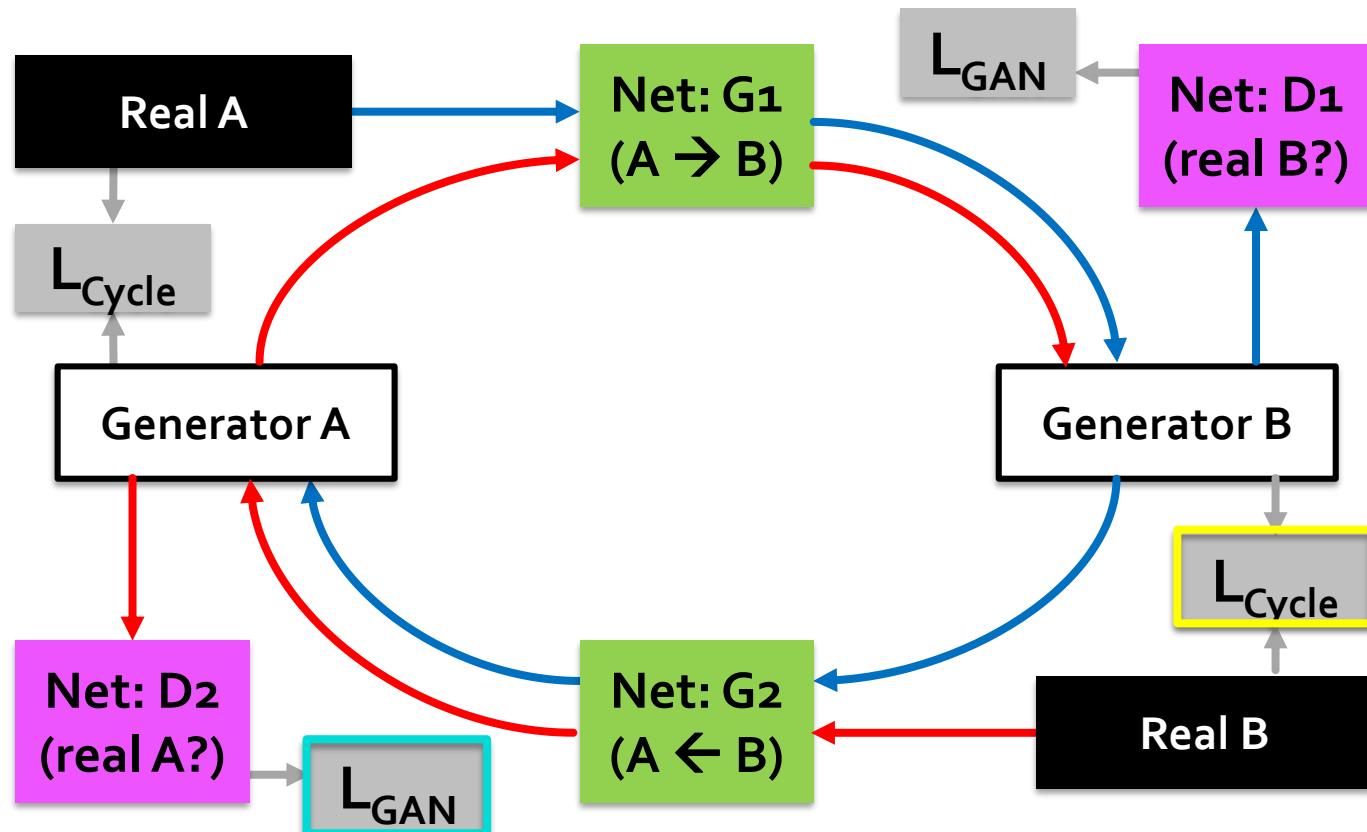
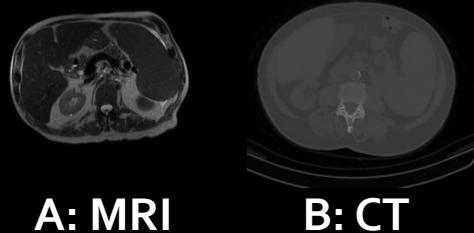
CycleGAN



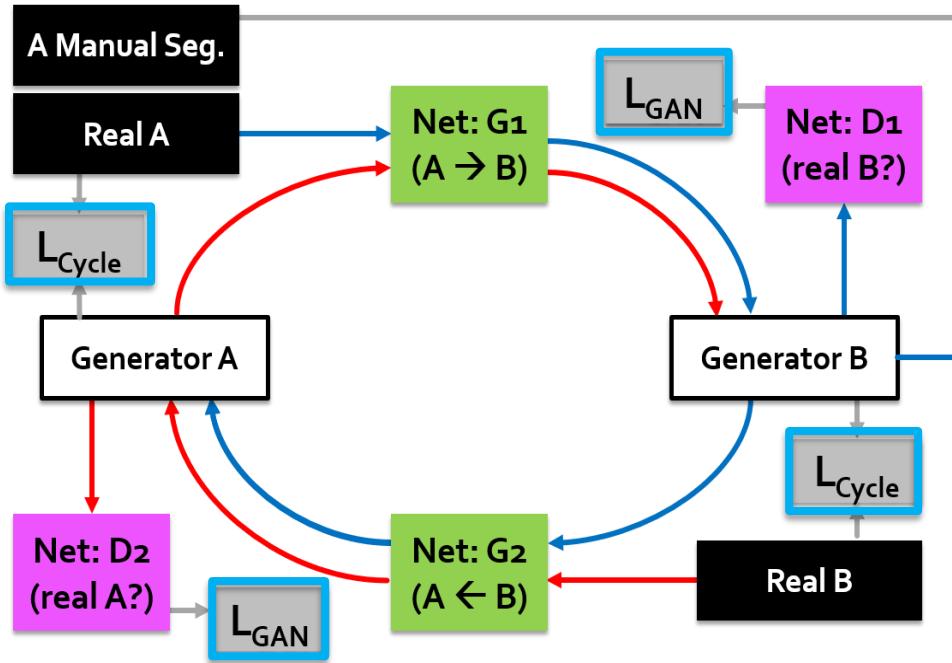
Cycle 1



Cycle 2



Loss Functions



$$\begin{aligned}\mathcal{L}_{GAN}(G_1, D_1, A, B) &= E_{y \sim B}[\log D_1(y)] \\ &\quad + E_{x \sim A}[\log(1 - D_1(G_1(x)))]\end{aligned}$$

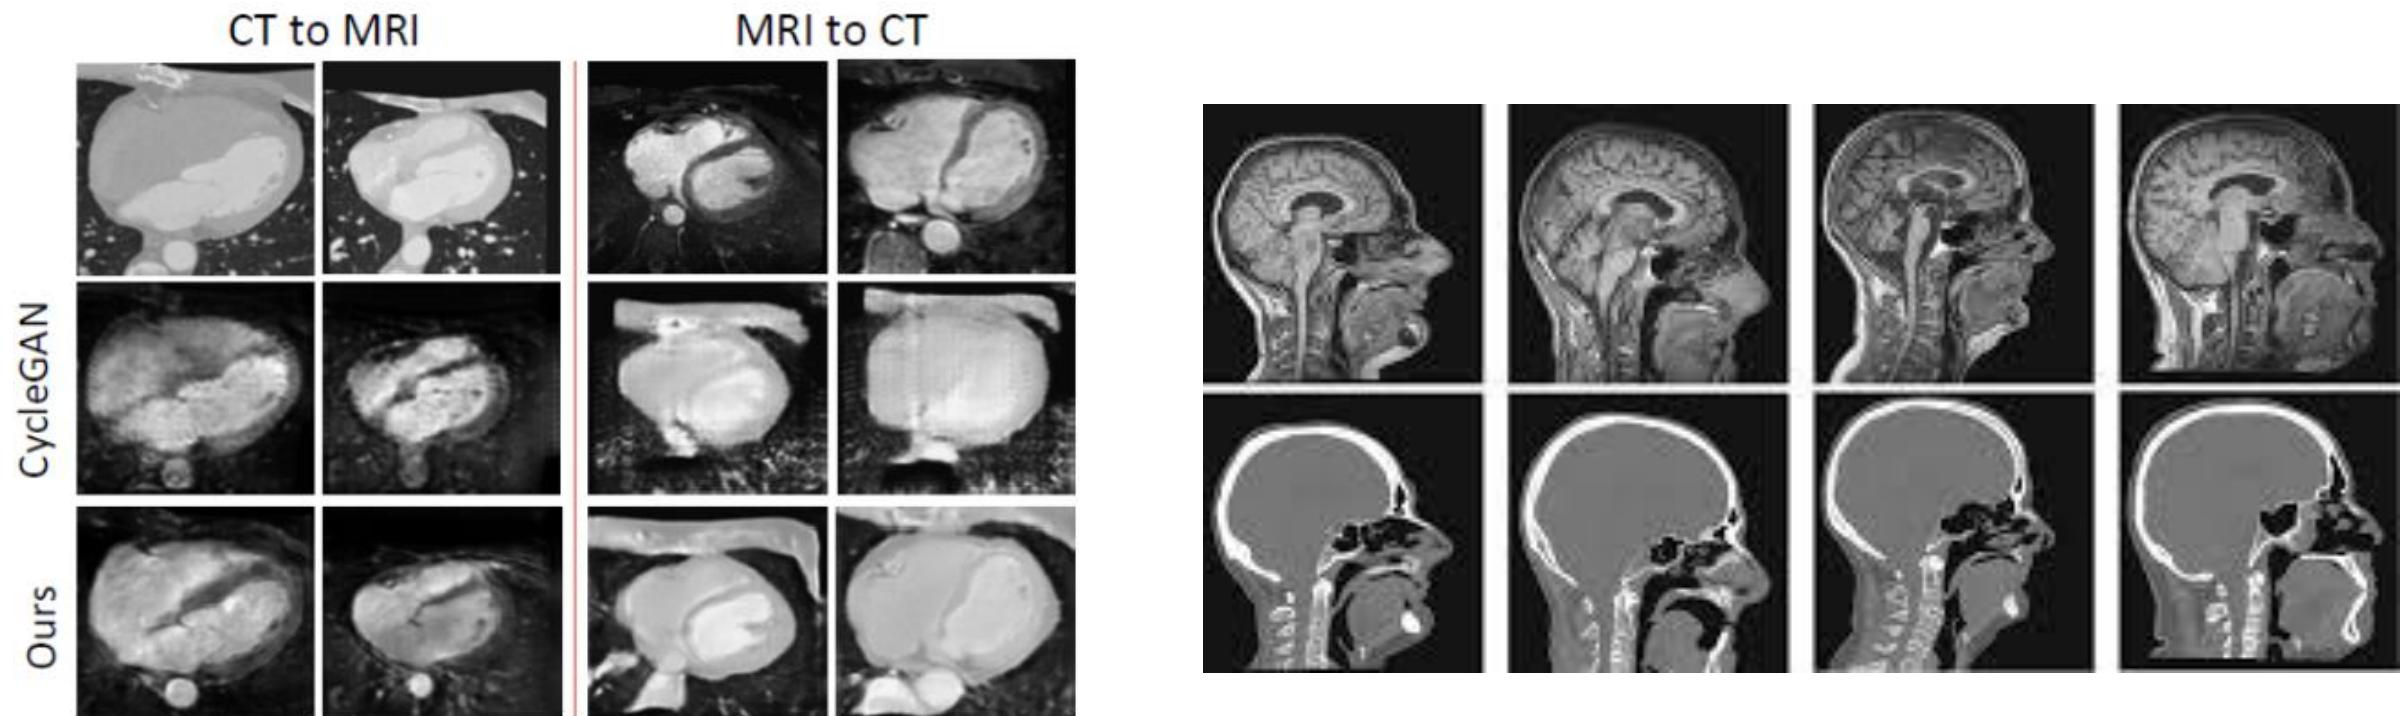
$$\begin{aligned}\mathcal{L}_{GAN}(G_2, D_2, B, A) &= E_{x \sim A}[\log D_2(x)] \\ &\quad + E_{y \sim B}[\log(1 - D_2(G_2(y)))]\end{aligned}$$

$$\mathcal{L}_{cycle}(G_1, G_2, A) = E_{x \sim A}[\|G_2(G_1(x)) - x\|_1]$$

$$\mathcal{L}_{cycle}(G_2, G_1, B) = E_{y \sim B}[\|G_1(G_2(y)) - y\|_1]$$

$$\mathcal{L}_{total} = \lambda_1 \cdot \mathcal{L}_{GAN}(G_1, D_1, A, B) + \lambda_2 \cdot \mathcal{L}_{GAN}(G_2, D_2, B, A) + \lambda_3 \cdot \mathcal{L}_{cycle}(G_1, G_2, A) + \lambda_4 \cdot \mathcal{L}_{cycle}(G_2, G_1, B)$$

On Medical Image



<https://towardsdatascience.com/cvpr-2018-paper-summary-translating-and-segmenting-multimodal-medical-volumes-with-cycle-and-e6381b4a2690>

Author

Image-to-Image Translation with Conditional Adversarial Nets

Phillip Isola

Jun-Yan Zhu

Tinghui Zhou

Alexei A. Efros

University of California, Berkeley
In CVPR 2017

Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks

Jun-Yan Zhu* **Taesung Park*** **Phillip Isola** **Alexei A. Efros**

UC Berkeley

In ICCV 2017

Outstanding Doctoral Dissertation Award: Jun-Yan Zhu
SIGGRAPH 2018



Code

train.py

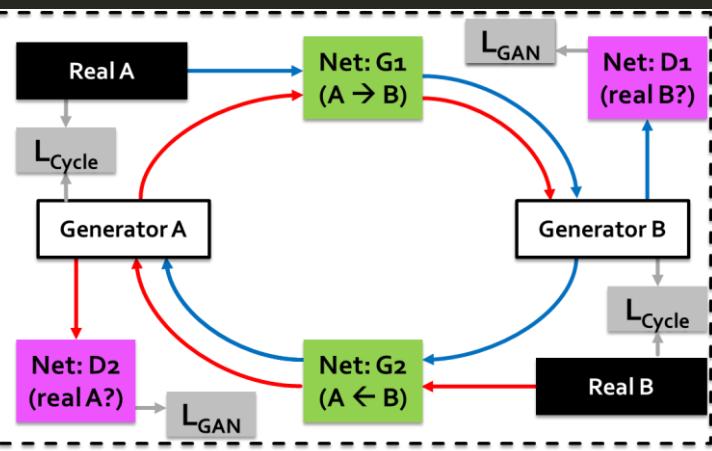
```
model = create_model(opt)
model.setup(opt)
visualizer = Visualizer(opt)
total_steps = 0

for epoch in range(opt.epoch_count, opt.niter + opt.niter_decay + 1):
    epoch_start_time = time.time()
    iter_data_time = time.time()
    epoch_iter = 0

    for i, data in enumerate(dataset):
        iter_start_time = time.time()
        if total_steps % opt.print_freq == 0:
            t_data = iter_start_time - iter_data_time
        visualizer.reset()
        total_steps += opt.batch_size
        epoch_iter += opt.batch_size
        model.set_input(data)
        model.optimize_parameters()
```

Code

cycle_gan_model.py



```
# load/define networks
# The naming conversion is different from those used in the paper
# Code (paper): G_A (G), G_B (F), D_A (D_Y), D_B (D_X)
self.netG_A = networks.define_G(opt.input_nc, opt.output_nc, opt.ngf, opt.netG, opt.norm,
                                not opt.no_dropout, opt.init_type, opt.init_gain, self.gpu_ids)
self.netG_B = networks.define_G(opt.output_nc, opt.input_nc, opt.ngf, opt.netG, opt.norm,
                                not opt.no_dropout, opt.init_type, opt.init_gain, self.gpu_ids)

if self.isTrain:
    use_sigmoid = opt.no_lsgan
    self.netD_A = networks.define_D(opt.output_nc, opt.ndf, opt.netD,
                                    opt.n_layers_D, opt.norm, use_sigmoid, opt.init_type, opt.init_gain, self.gpu_ids)
    self.netD_B = networks.define_D(opt.input_nc, opt.ndf, opt.netD,
                                    opt.n_layers_D, opt.norm, use_sigmoid, opt.init_type, opt.init_gain, self.gpu_ids)

if self.isTrain:
    self.fake_A_pool = ImagePool(opt.pool_size)
    self.fake_B_pool = ImagePool(opt.pool_size)
    # define loss functions
    self.criterionGAN = networks.GANLoss(use_Lsgan=not opt.no_lsgan).to(self.device)
    self.criterionCycle = torch.nn.L1Loss()
    self.criterionIdt = torch.nn.L1Loss()
    # initialize optimizers
    self.optimizer_G = torch.optim.Adam(itertools.chain(self.netG_A.parameters(), self.netG_B.parameters()),
                                       lr=opt.lr, betas=(opt.beta1, 0.999))
    self.optimizer_D = torch.optim.Adam(itertools.chain(self.netD_A.parameters(), self.netD_B.parameters()),
                                       lr=opt.lr, betas=(opt.beta1, 0.999))
    self.optimizers = []
    self.optimizers.append(self.optimizer_G)
    self.optimizers.append(self.optimizer_D)
```

Code

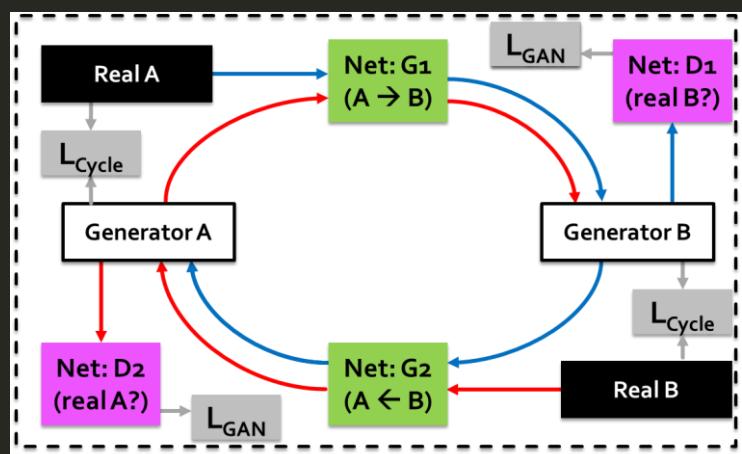
unaligned_dataset.py

```
def __getitem__(self, index):
    A_path = self.A_paths[index % self.A_size]
    if self.opt.serial_batches:
        index_B = index % self.B_size
    else:
        index_B = random.randint(0, self.B_size - 1)
    B_path = self.B_paths[index_B]
    # print('(A, B) = (%d, %d)' % (index_A, index_B))
    A_img = Image.open(A_path).convert('RGB')
    B_img = Image.open(B_path).convert('RGB')

    A = self.transform(A_img)
    B = self.transform(B_img)
    if self.opt.direction == 'BtoA':
        input_nc = self.opt.output_nc
        output_nc = self.opt.input_nc
    else:
        input_nc = self.opt.input_nc
        output_nc = self.opt.output_nc
```

Code

cycle_gan_model.py



```
def optimize_parameters(self):  
    # forward  
    self.forward()  
    # G_A and G_B  
    self.set_requires_grad([self.netD_A, self.netD_B], False)  
    self.optimizer_G.zero_grad()  
    self.backward_G()  
    self.optimizer_G.step()  
    # D_A and D_B  
    self.set_requires_grad([self.netD_A, self.netD_B], True)  
    self.optimizer_D.zero_grad()  
    self.backward_D_A()  
    self.backward_D_B()  
    self.optimizer_D.step()  
  
def forward(self):  
    self.fake_B = self.netG_A(self.real_A)  
    self.rec_A = self.netG_B(self.fake_B)  
  
    self.fake_A = self.netG_B(self.real_B)  
    self.rec_B = self.netG_A(self.fake_A)
```

Code

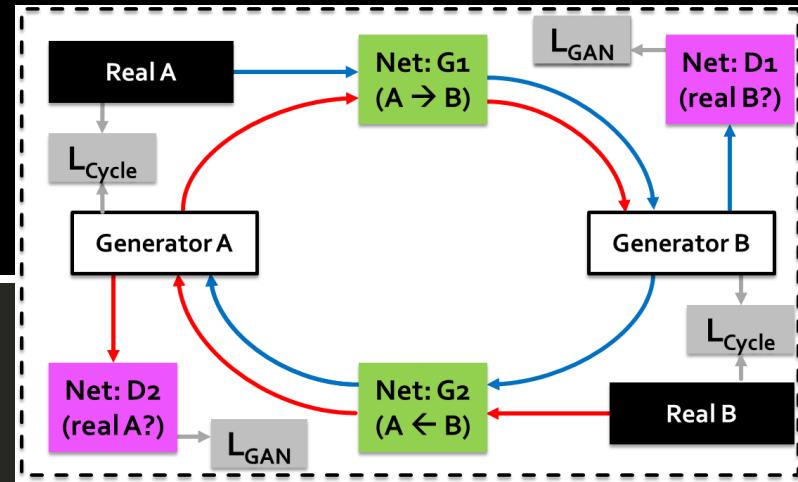
cycle_gan_model.py

```
def optimize_parameters(self):
    # forward
    self.forward()
    # G_A and G_B
    self.set_requires_grad([self.netD_A, self.netD_B], False)
    self.optimizer_G.zero_grad()
    self.backward_G()
    self.optimizer_G.step()
    # D_A and D_B
    self.set_requires_grad([self.netD_A, self.netD_B], True)
    self.optimizer_D.zero_grad()
    self.backward_D_A()
    self.backward_D_B()
    self.optimizer_D.step()
```

```
def backward_G(self):
    lambda_idt = self.opt.lambda_identity
    lambda_A = self.opt.lambda_A
    lambda_B = self.opt.lambda_B

    # GAN loss D_A(G_A(A))
    self.loss_G_A = self.criterionGAN(self.netD_A(self.fake_B), True)
    # GAN loss D_B(G_B(B))
    self.loss_G_B = self.criterionGAN(self.netD_B(self.fake_A), True)
    # Forward cycle loss
    self.loss_cycle_A = self.criterionCycle(self.rec_A, self.real_A) * lambda_A
    # Backward cycle loss
    self.loss_cycle_B = self.criterionCycle(self.rec_B, self.real_B) * lambda_B
    # combined loss
    self.loss_G = self.loss_G_A + self.loss_G_B + self.loss_cycle_A +
                  self.loss_cycle_B + self.loss_idt_A + self.loss_idt_B

    self.loss_G.backward()
```



Code

cycle_gan_model.py

```
def optimize_parameters(self):
    # forward
    self.forward()
    # G_A and G_B
    self.set_requires_grad([self.netD_A, self.netD_B], False)
    self.optimizer_G.zero_grad()
    self.backward_G()
    self.optimizer_G.step()
    # D_A and D_B
    self.set_requires_grad([self.netD_A, self.netD_B], True)
    self.optimizer_D.zero_grad()
    self.backward_D_A()
    self.backward_D_B()
    self.optimizer_D.step()
```

```
def backward_D_basic(self, netD, real, fake):
    # Real
    pred_real = netD(real)
    loss_D_real = self.criterionGAN(pred_real, True)
    # Fake
    pred_fake = netD(fake.detach())
    loss_D_fake = self.criterionGAN(pred_fake, False)
    # Combined loss
    loss_D = (loss_D_real + loss_D_fake) * 0.5
    # backward
    loss_D.backward()
    return loss_D

def backward_D_A(self):
    fake_B = self.fake_B_pool.query(self.fake_B)
    self.loss_D_A = self.backward_D_basic(self.netD_A, self.real_B, fake_B)

def backward_D_B(self):
    fake_A = self.fake_A_pool.query(self.fake_A)
    self.loss_D_B = self.backward_D_basic(self.netD_B, self.real_A, fake_A)
```

